

Lecture 5: Understanding Convnets and Knowledge Transfer

Deep Learning @ UvA

Previous Lecture

- What are the Convolutional Neural Networks?
- Why are they so important for Computer Vision?
- How do they differ from standard Neural Networks?
- How can we train a Convolutional Neural Network?

Lecture Overview

- What do convolutions look like?
- Build on the visual intuition behind Convnets
- Deep Learning Feature maps
- Transfer Learning

Understanding convnets

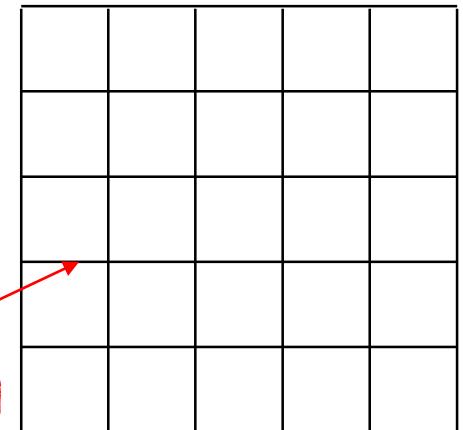
UVA DEEP LEARNING COURSE
EFSTRATIOS GAVVES
UNDERSTANDING CONVNETS AND KNOWLEDGE TRANSFER - 4



How large filters?

- Traditionally, medium sized filters (smaller than 11×11)
- Modern architectures prefer small filter sizes (e.g. 3×3)
- We lose frequency resolution
- Fewer parameters to train
- Deeper networks of cascade filters
 - Still, the same output dimensionalities

$$(2d + 1)^2$$

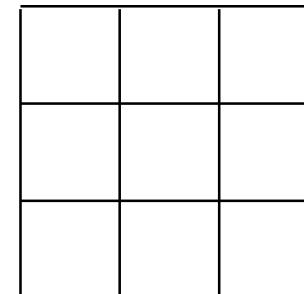


For stride 1 the final
image has dimensionality
 $\frac{H-2d-1}{1} + 1 = H - 2d$

vs.

$$(d + 1)^2$$

$$(d + 1)^2$$



*

$$(Layer l) * (Layer l + 1)$$

For stride 1 the first image has dimensionality
 $H - d$, the second image $\frac{H-d-d-1}{1} + 1 = H - 2d$

Invariance and covariance of filters

- Filters learn how different variances affect appearance
- Different layers and different hierarchies focus on different transformations
- For different objects filters reproduce different behaviors

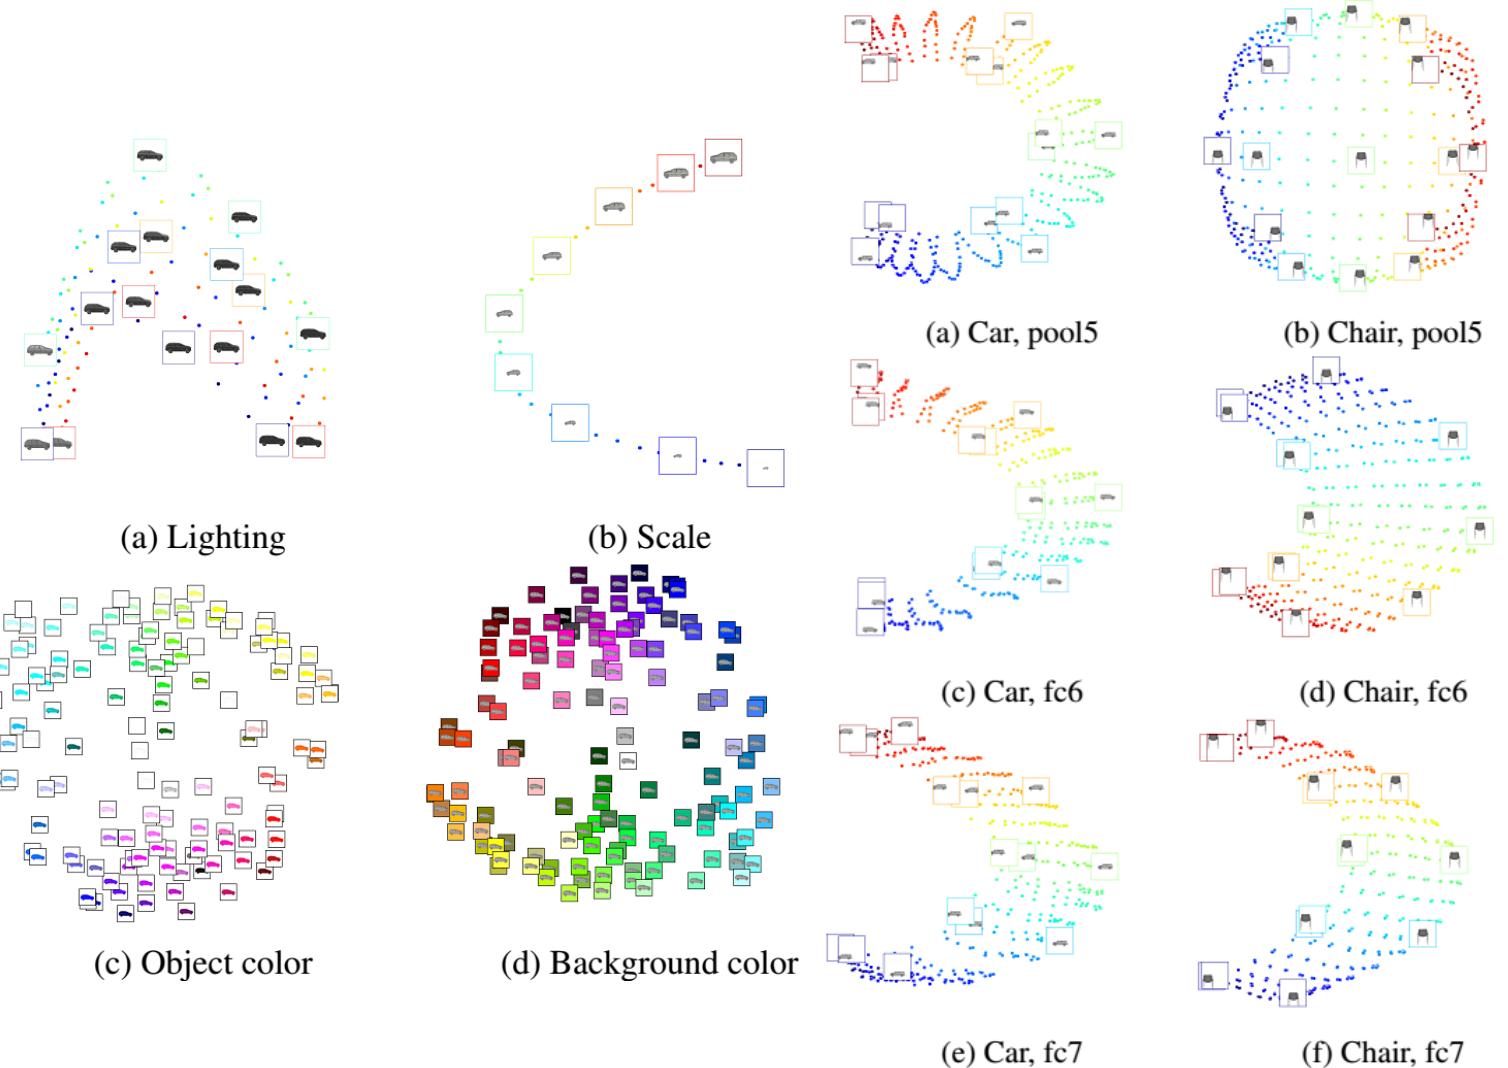


Figure 3: PCA embeddings for 2D position on AlexNet.

Invariance and covariance of filters

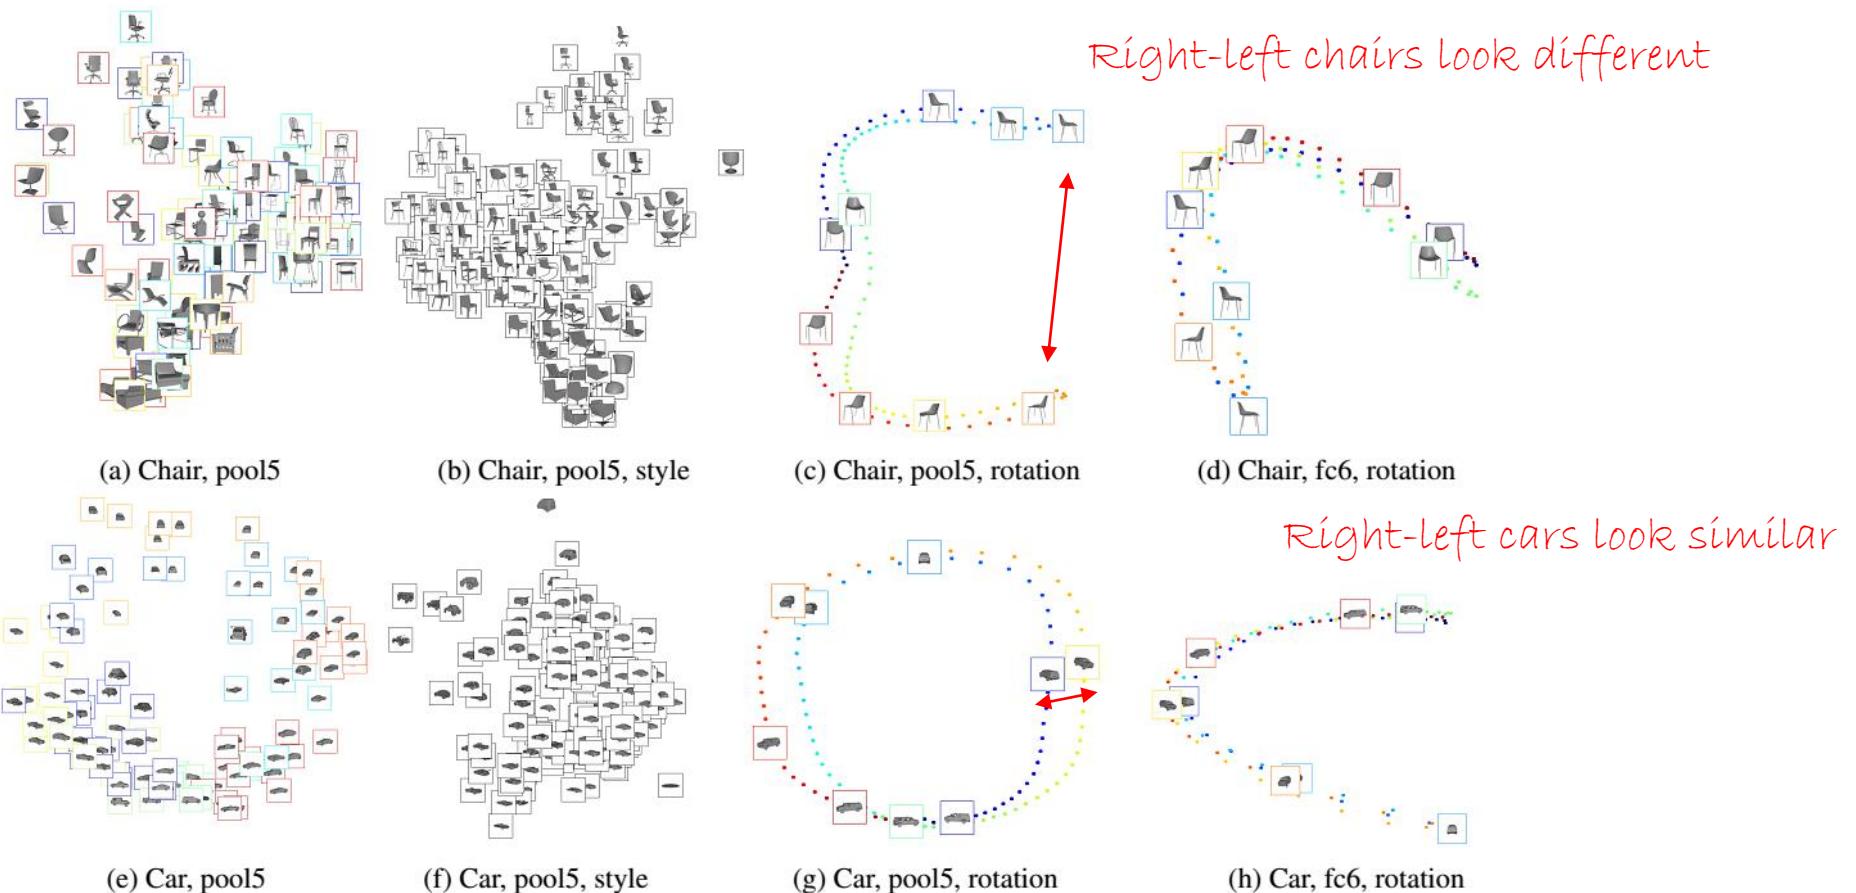


Figure 2: **Best viewed in the electronic version.** PCA embeddings (dims. 1,2) of AlexNet features for “chairs” (first row) and “cars” (second row). Column 1 – Direct embedding of the rendered images without viewpoint-style separation. Columns 2,3 – Embeddings associated with style (for all rotations) and rotation (for all styles). Column 4 – Rotation embedding for fc6, which is qualitatively different than pool5. Colors correspond to orientation and can be interpreted via the example images in columns 3,4. Similar results for other categories and PCA dimensions are available in the supplementary material.

The effect of the architecture

Depth is important

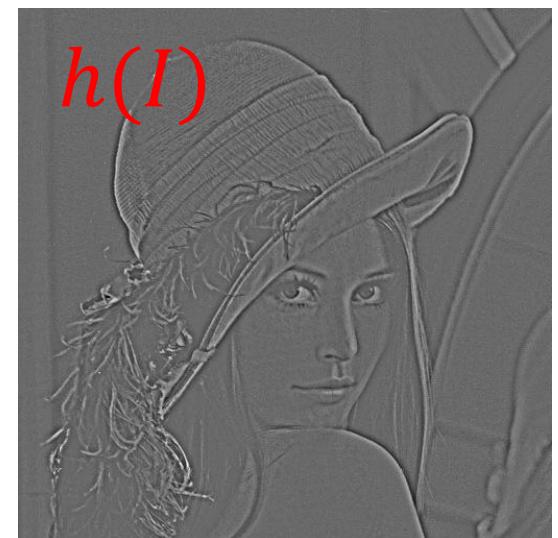
Early signs of overfitting

Overfitting

Error %	Train Top-1	Val Top-1	Val Top-5
Our replication of Krizhevsky <i>et al.</i> [18], 1 convnet	35.1	40.5	18.1
Removed layers 3,4	41.8	45.4	22.1
Removed layer 7	27.4	40.0	18.4
Removed layers 6,7	27.4	44.8	22.4
Removed layer 3,4,6,7	71.1	71.3	50.1
Adjust layers 6,7: 2048 units	40.3	41.7	18.8
Adjust layers 6,7: 8192 units	26.8	40.0	18.1
Our Model (as per Fig. 3)	33.1	38.4	16.5
Adjust layers 6,7: 2048 units	38.2	40.2	17.6
Adjust layers 6,7: 8192 units	22.0	38.8	17.0
Adjust layers 3,4,5: 512,1024,512 maps	18.8	37.5	16.0
Adjust layers 6,7: 8192 units and Layers 3,4,5: 512,1024,512 maps	10.0	38.3	16.9

Convolution activations are “images”

- $$\begin{aligned} h_I(x, y) &= I(x, y) * h \\ &= \sum_{i=-a}^a \sum_{j=-b}^b I(x - i, y - j) \cdot h(i, j) \end{aligned}$$
- For every image pixel we compute a convolved image pixel
- After each convolution we end up with a "new image"



Several interesting questions

- What do the image activations in different layers look like?
- What types of images create the strongest activations?
- What are the activations for the class “ostrich”?
- Do the activations occur in meaningful positions?

Feature maps

- The convolution activations are also called feature maps
- A deep network has several hierarchical layers
 - hence several feature maps

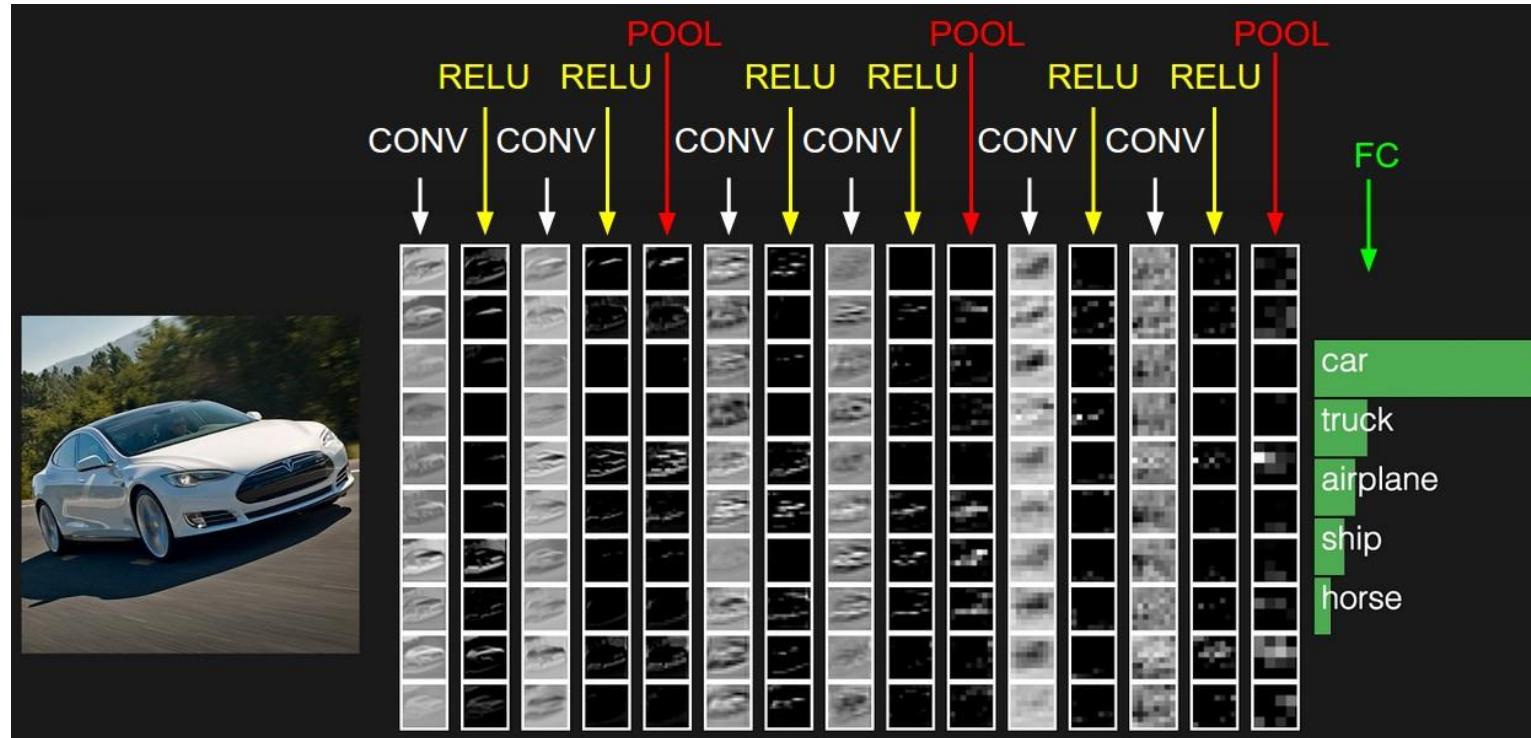


Image borrowed by
A. Karpathy

Feature map strength

- Naturally, given an image some feature maps will be “stronger”
 - Contribute higher amount to the final classification score
- Why? What pixel structure excited these particular feature maps?
 - Generally, what part of the picture excites a certain feature map?

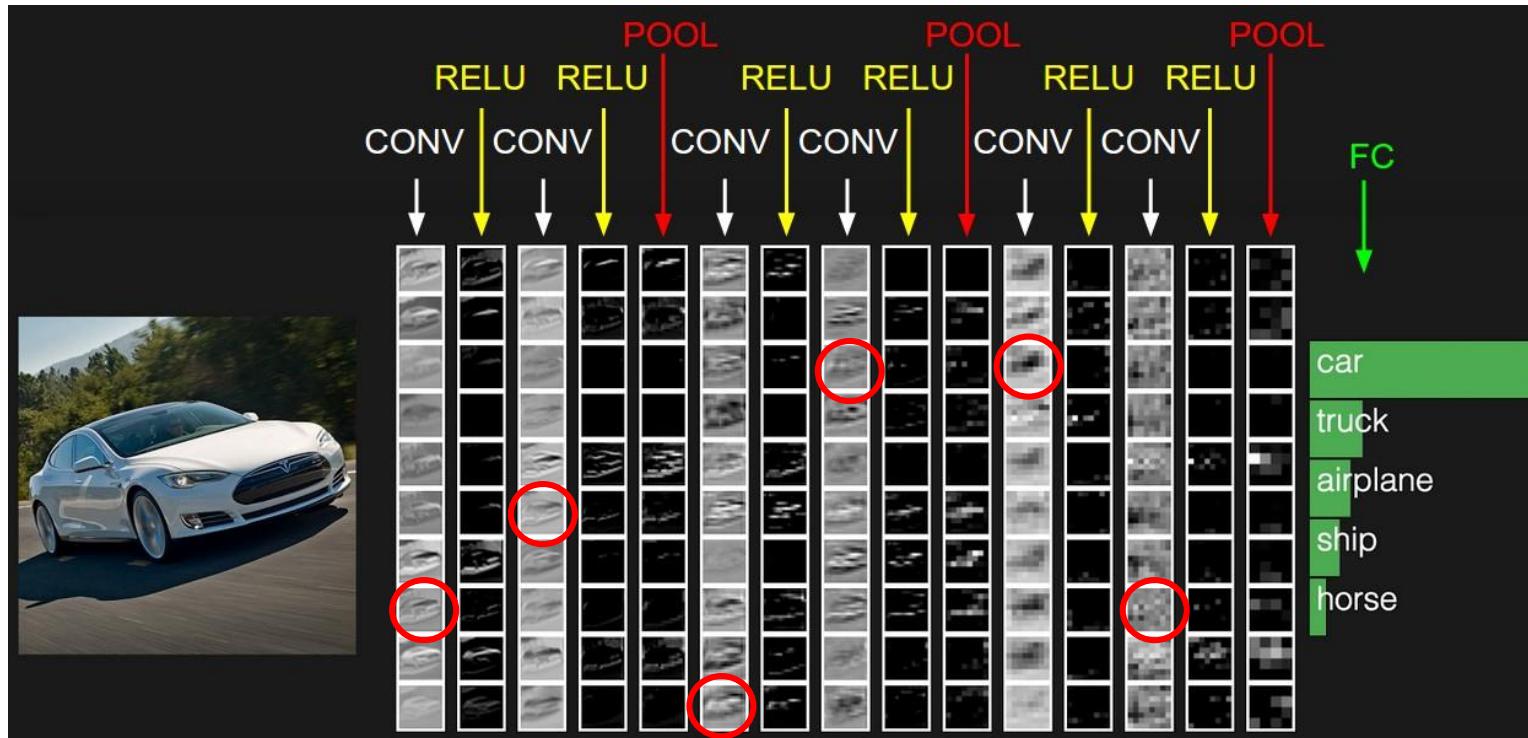
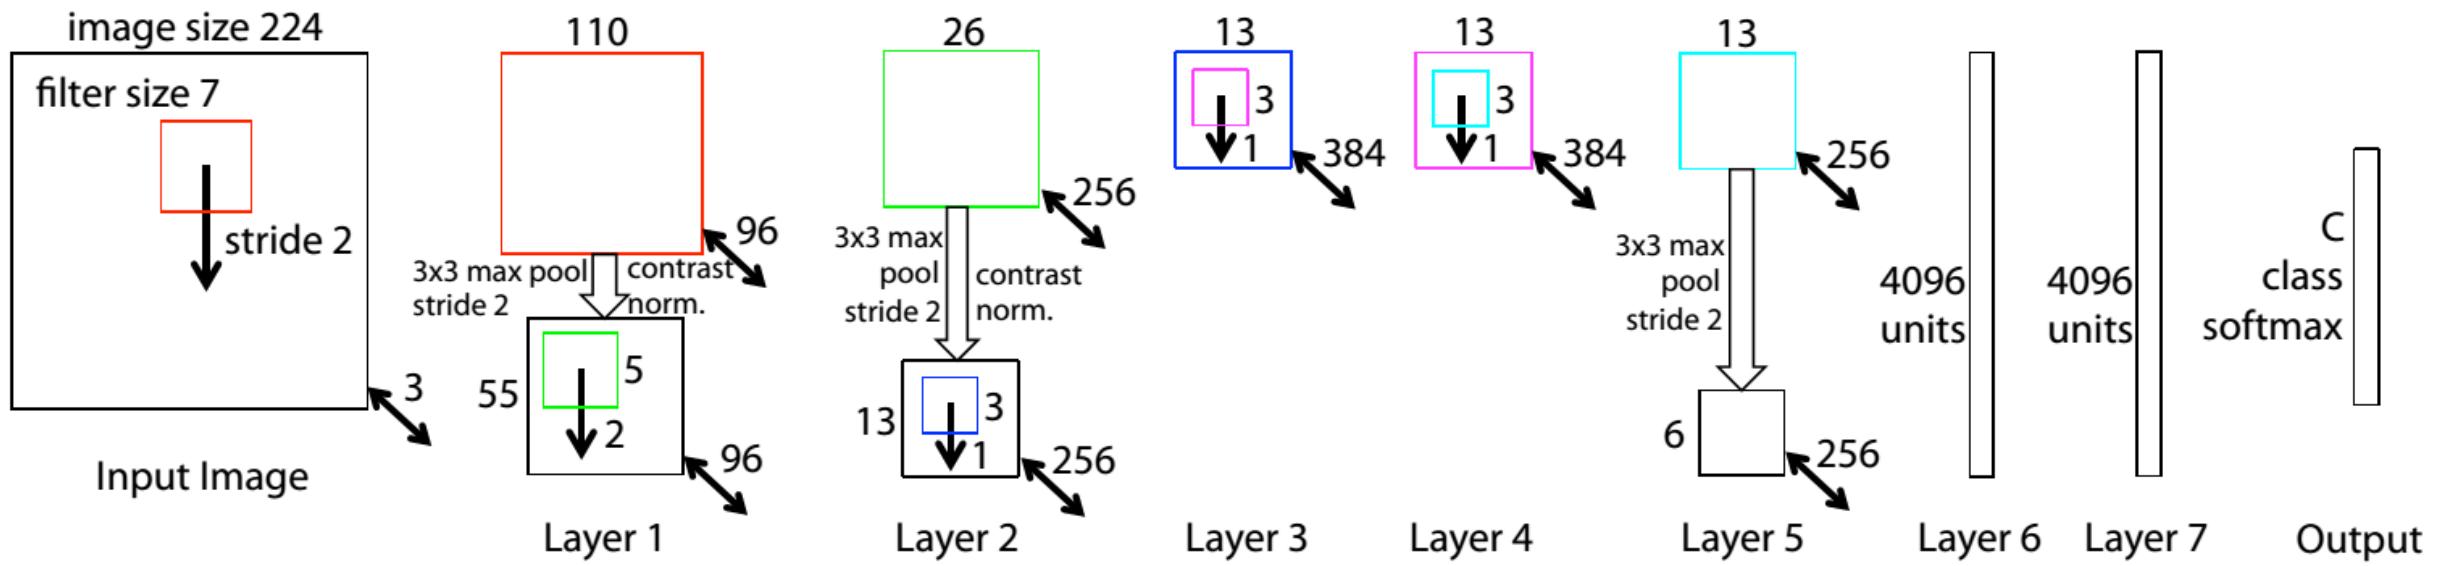


Image borrowed by
A. Karpathy

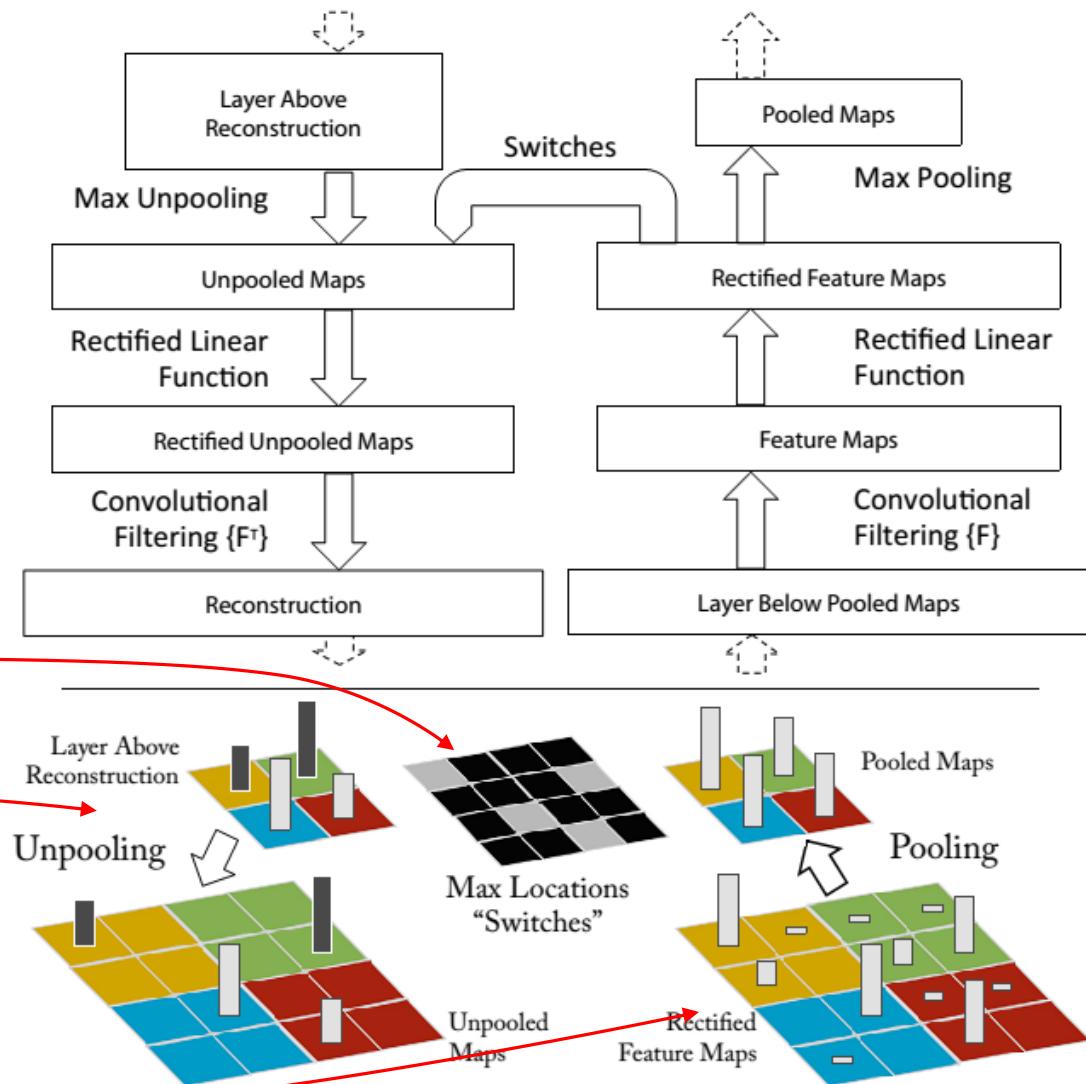
Starting from a convnet [Zeiler2014]



Visualization with a Deconvnet [Zeiler2014]

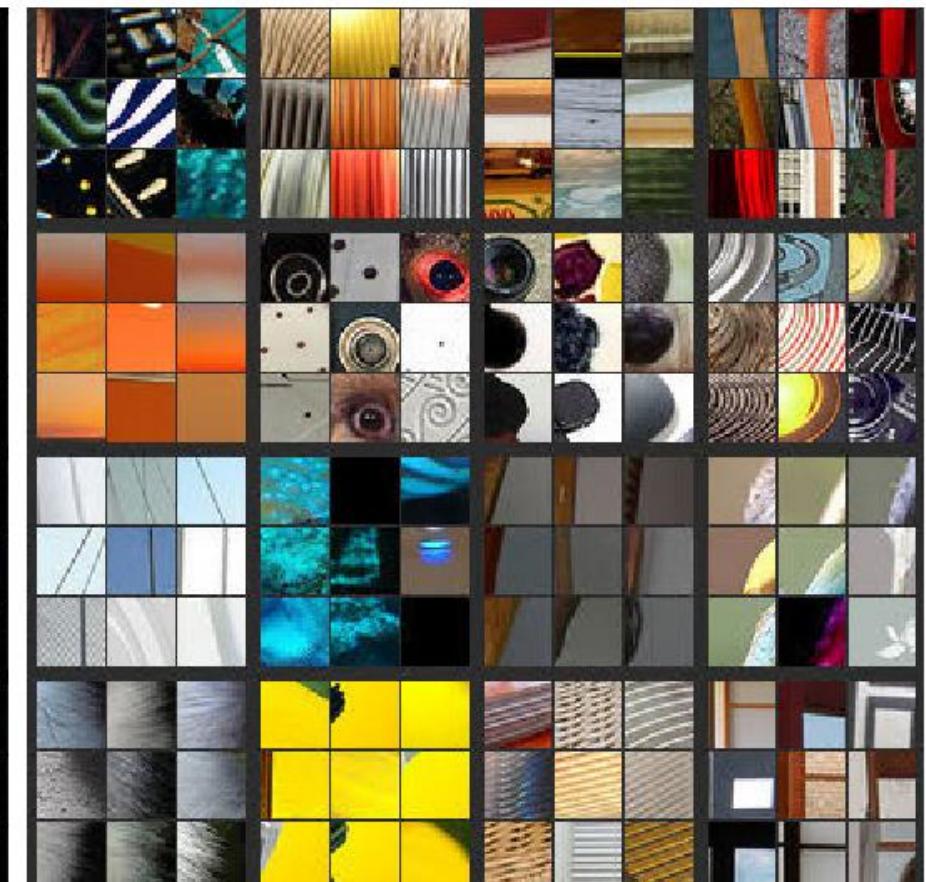
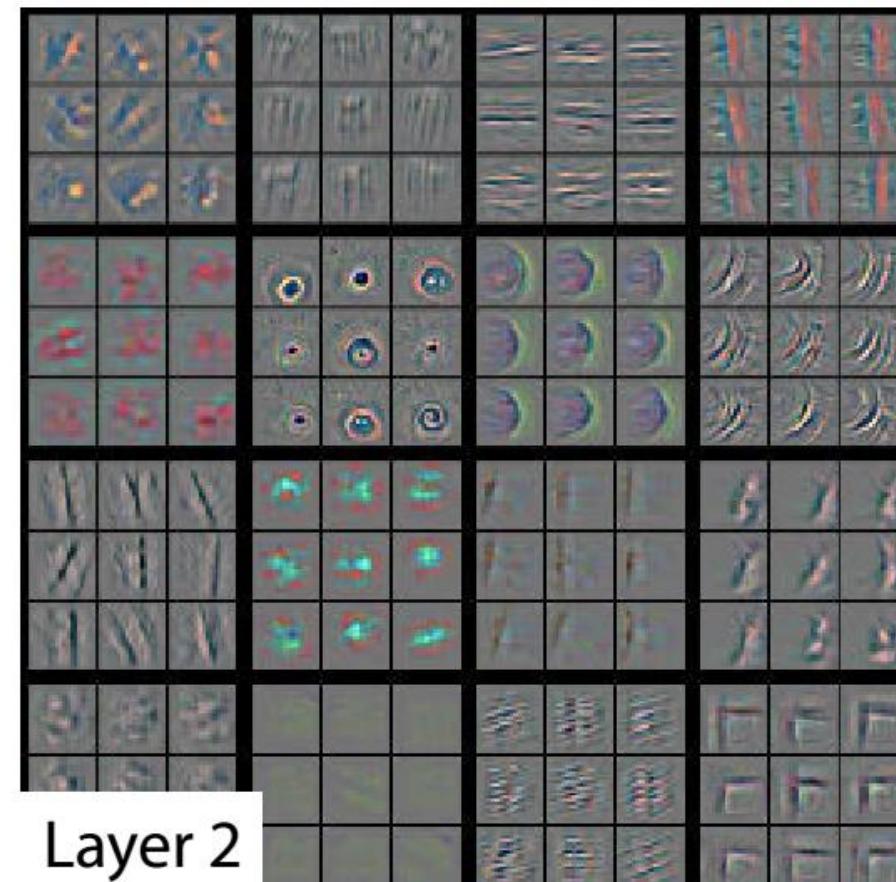
- Same as convnet, but in reverse
 - Instead of mapping pixels to activations, mapping activations to pixels
- Attach a deconvnet layer on every convnet layer
- Unpooling via switches that remember where max pooling was activated
- Deconv filtering equals to “reverse conv filtering”

$$F = \begin{bmatrix} 1 & 5 \\ 6 & 3 \end{bmatrix} \quad F^T = \begin{bmatrix} 3 & 6 \\ 5 & 1 \end{bmatrix}$$



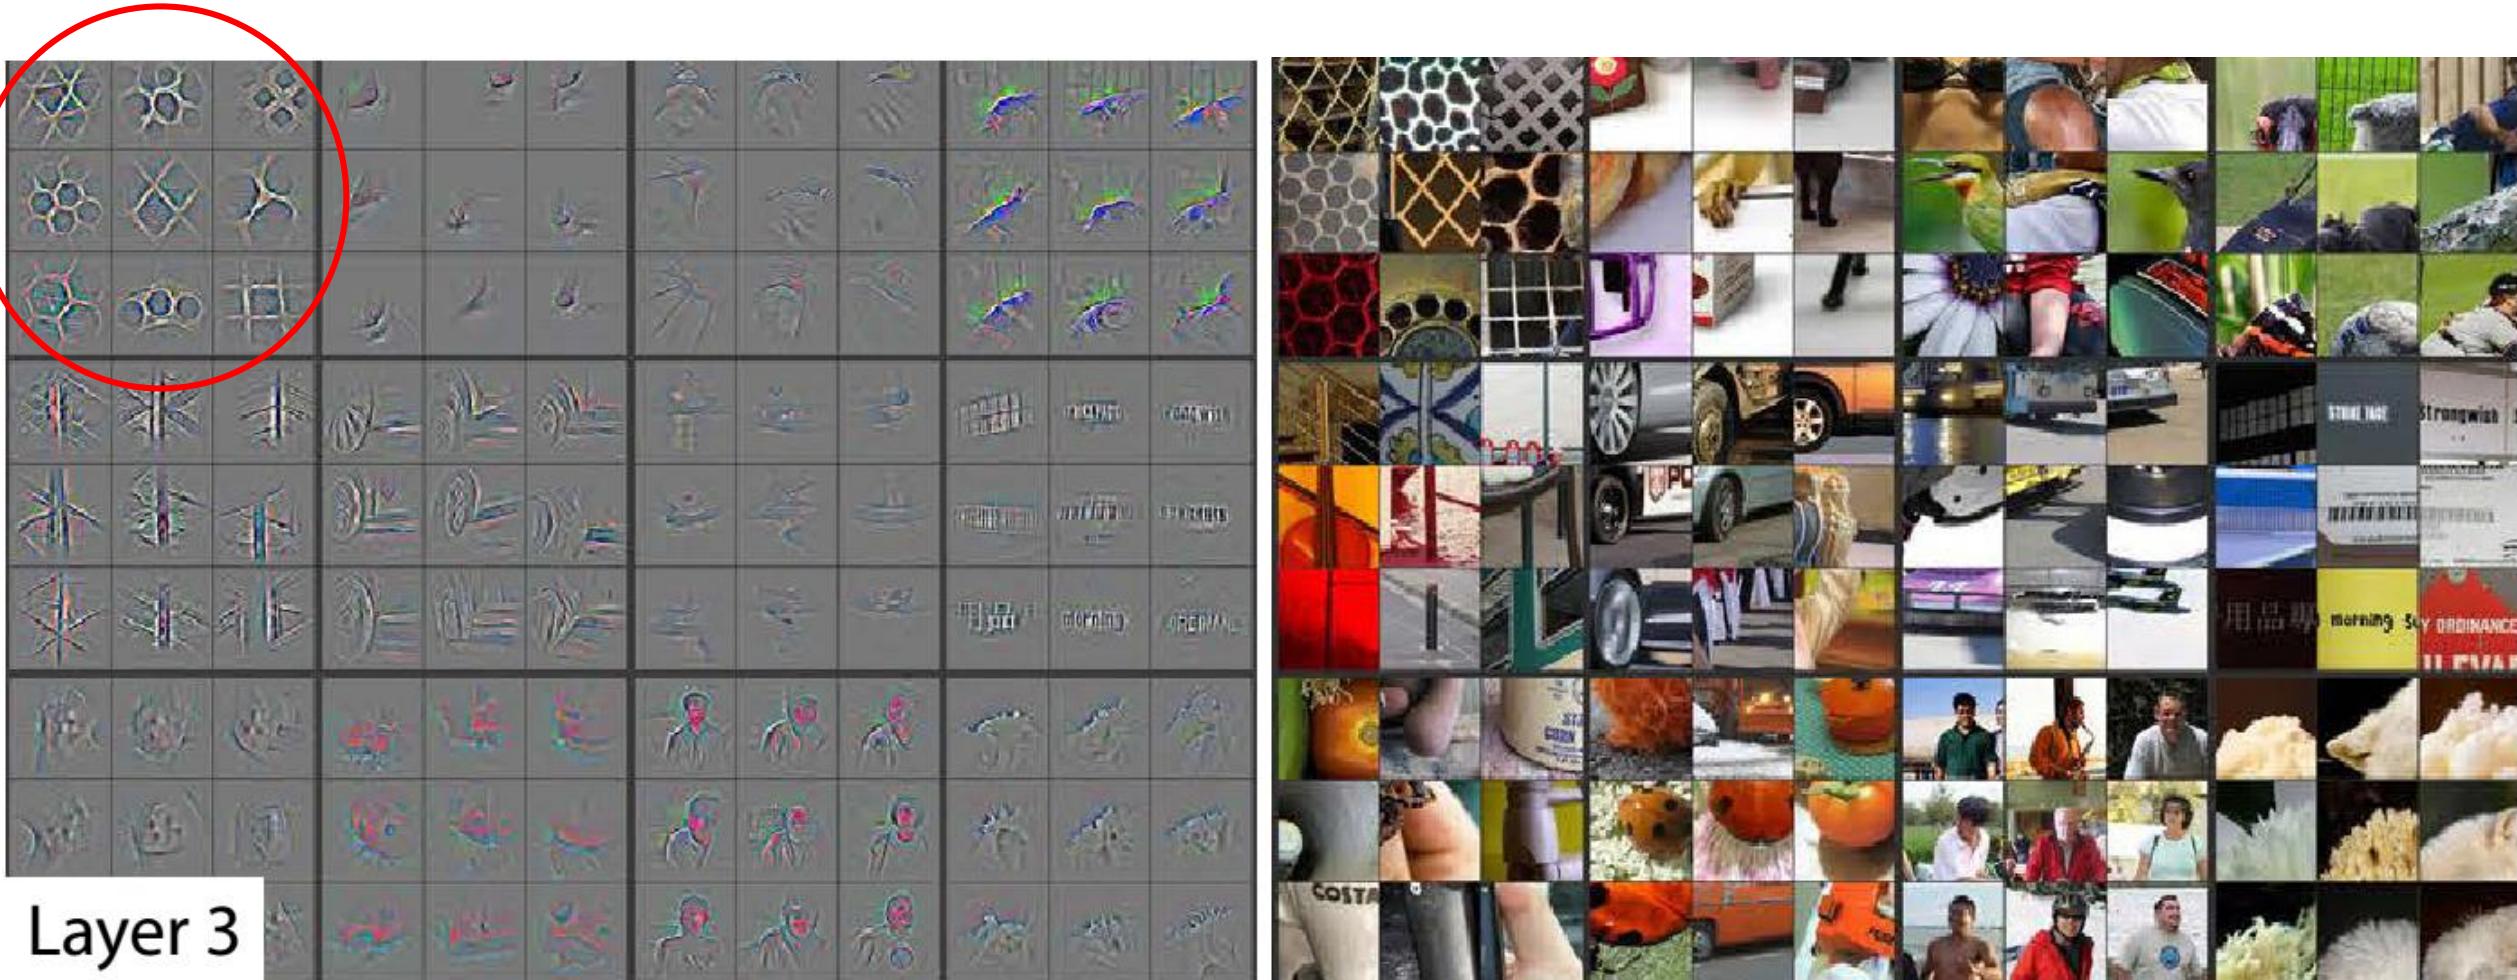
What excites feature maps?

- “Given a random feature map what are the top 9 activations?”



What excites feature maps?

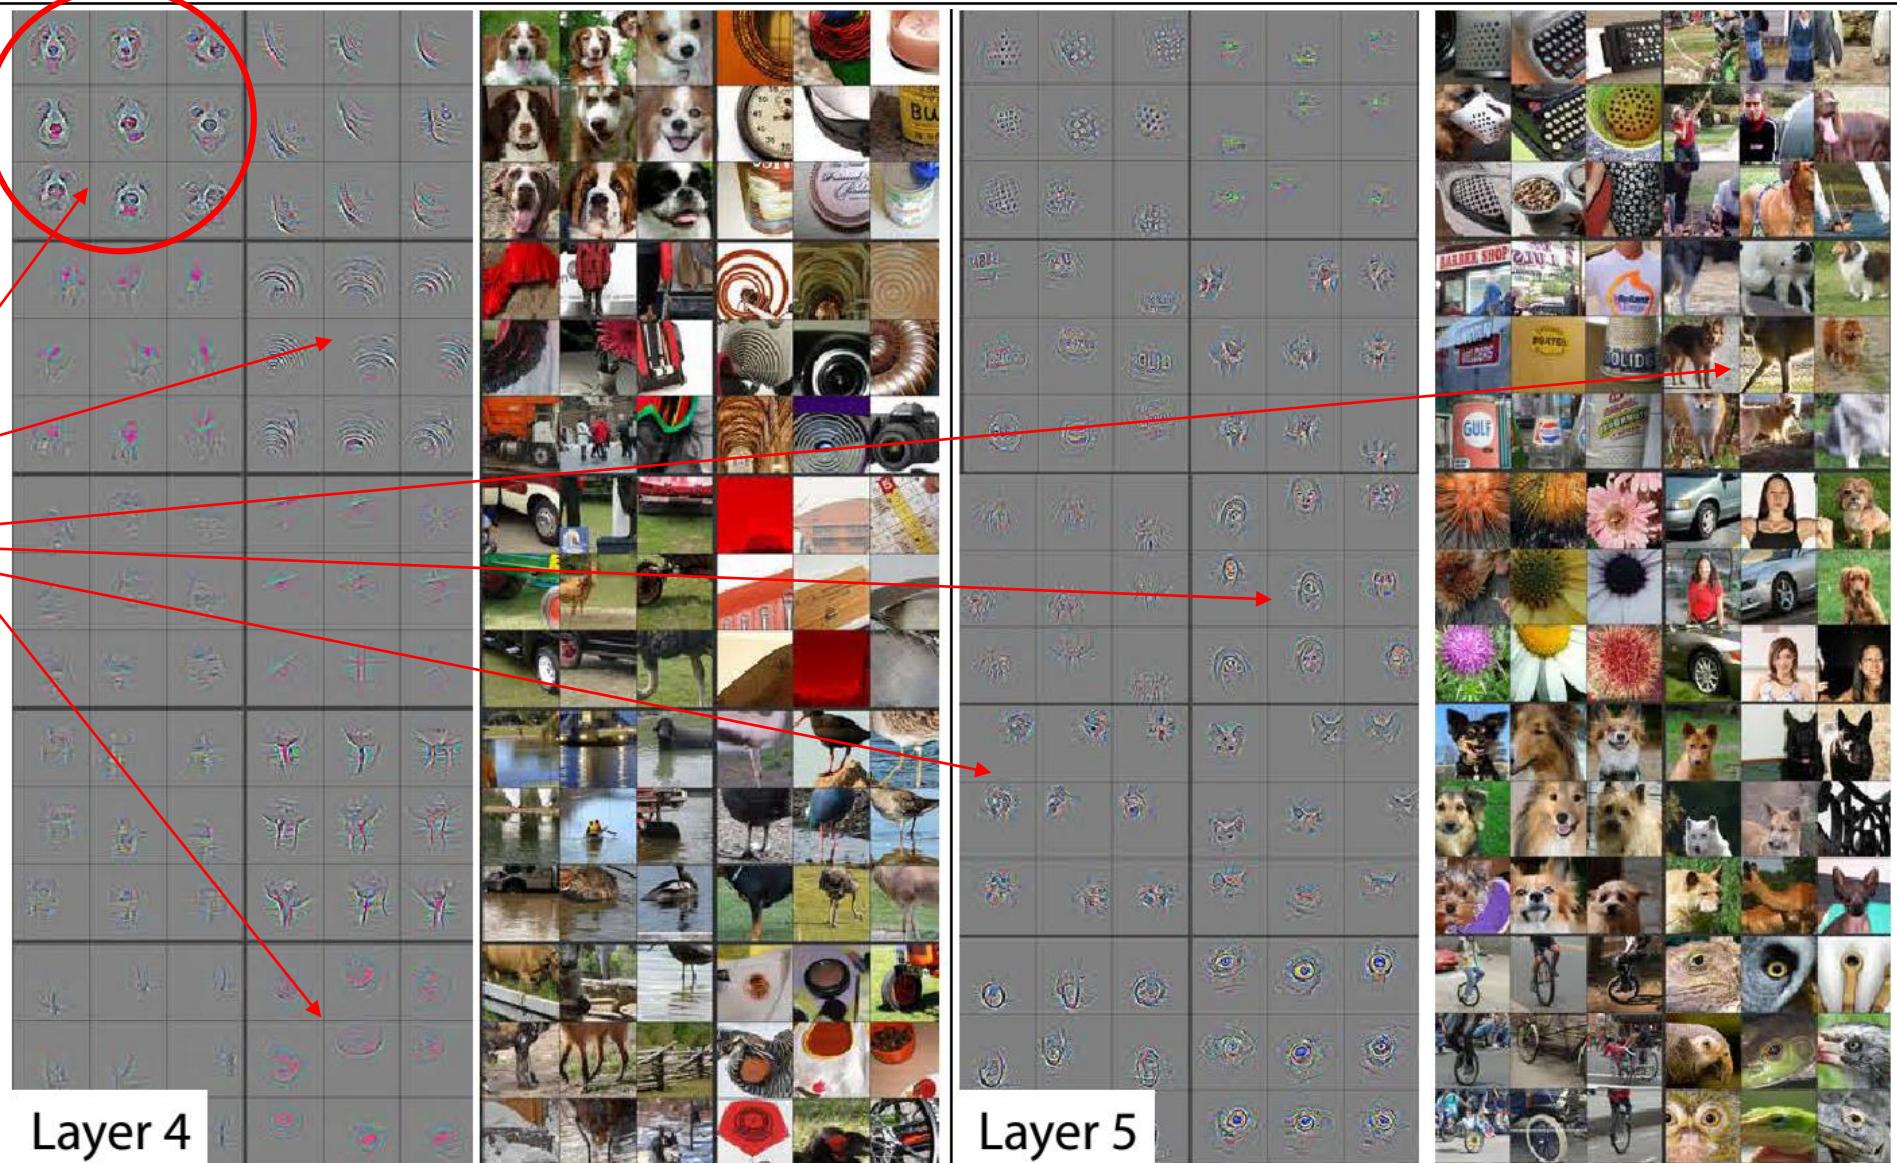
Similar activations from lower level visual patterns



What excites feature maps? [Zeiler2014]

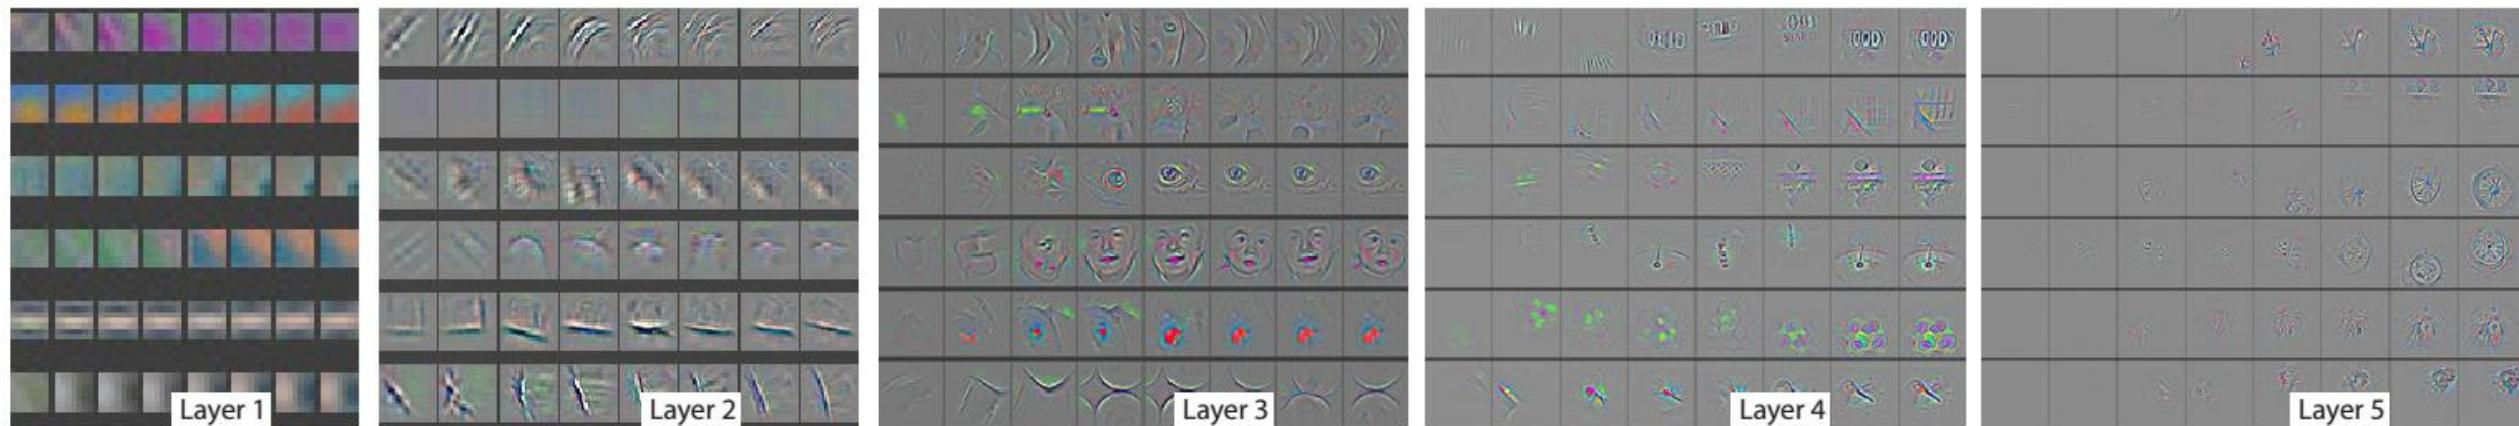
Similar activations from semantically similar pictures

Visual patterns become more and more intricate and specific (greater invariance)

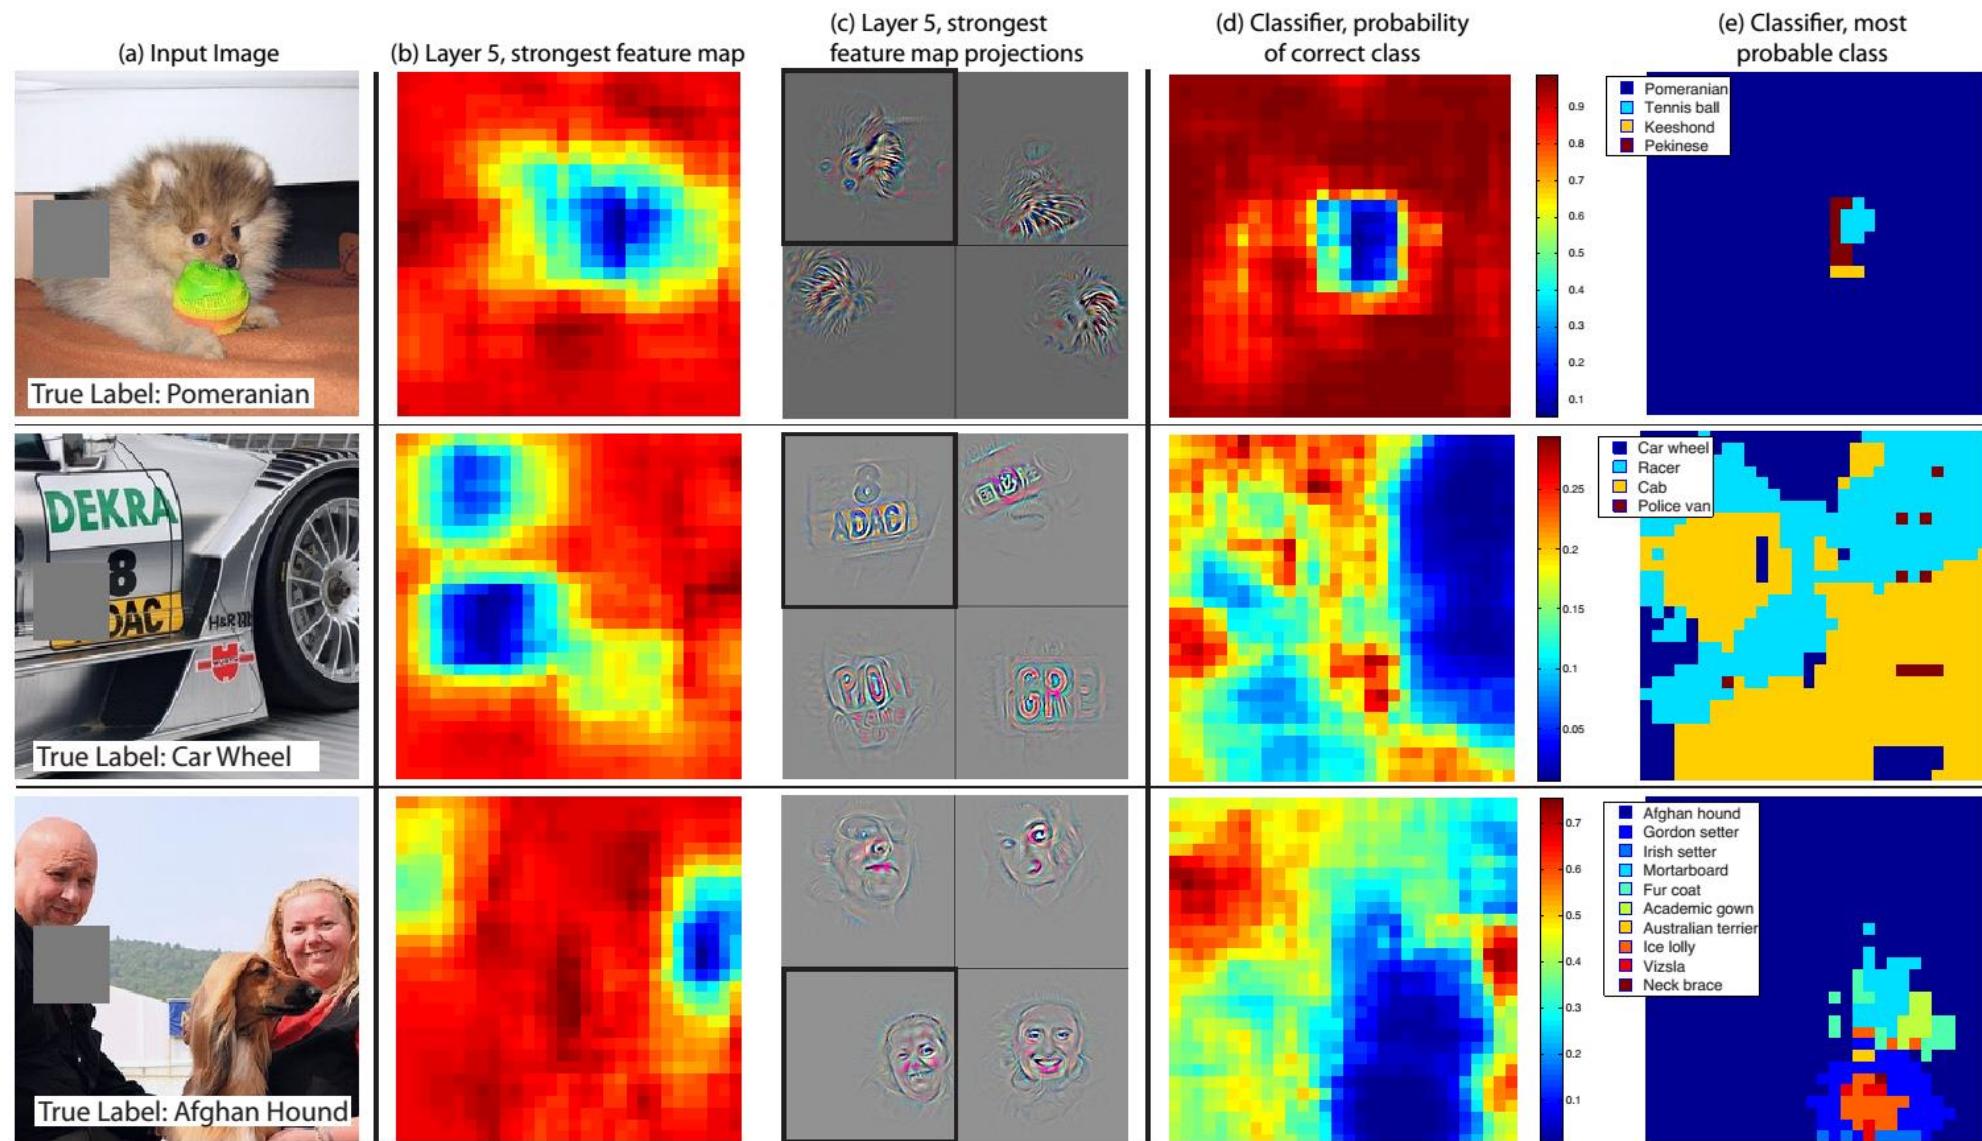


Feature evolution over training

- For a particular neuron (that generates a feature map)
- Pick the strongest activation during training
- For epochs 1, 2, 5, 10, 20, 30, 40, 64



But does a Convnet really learn the object?



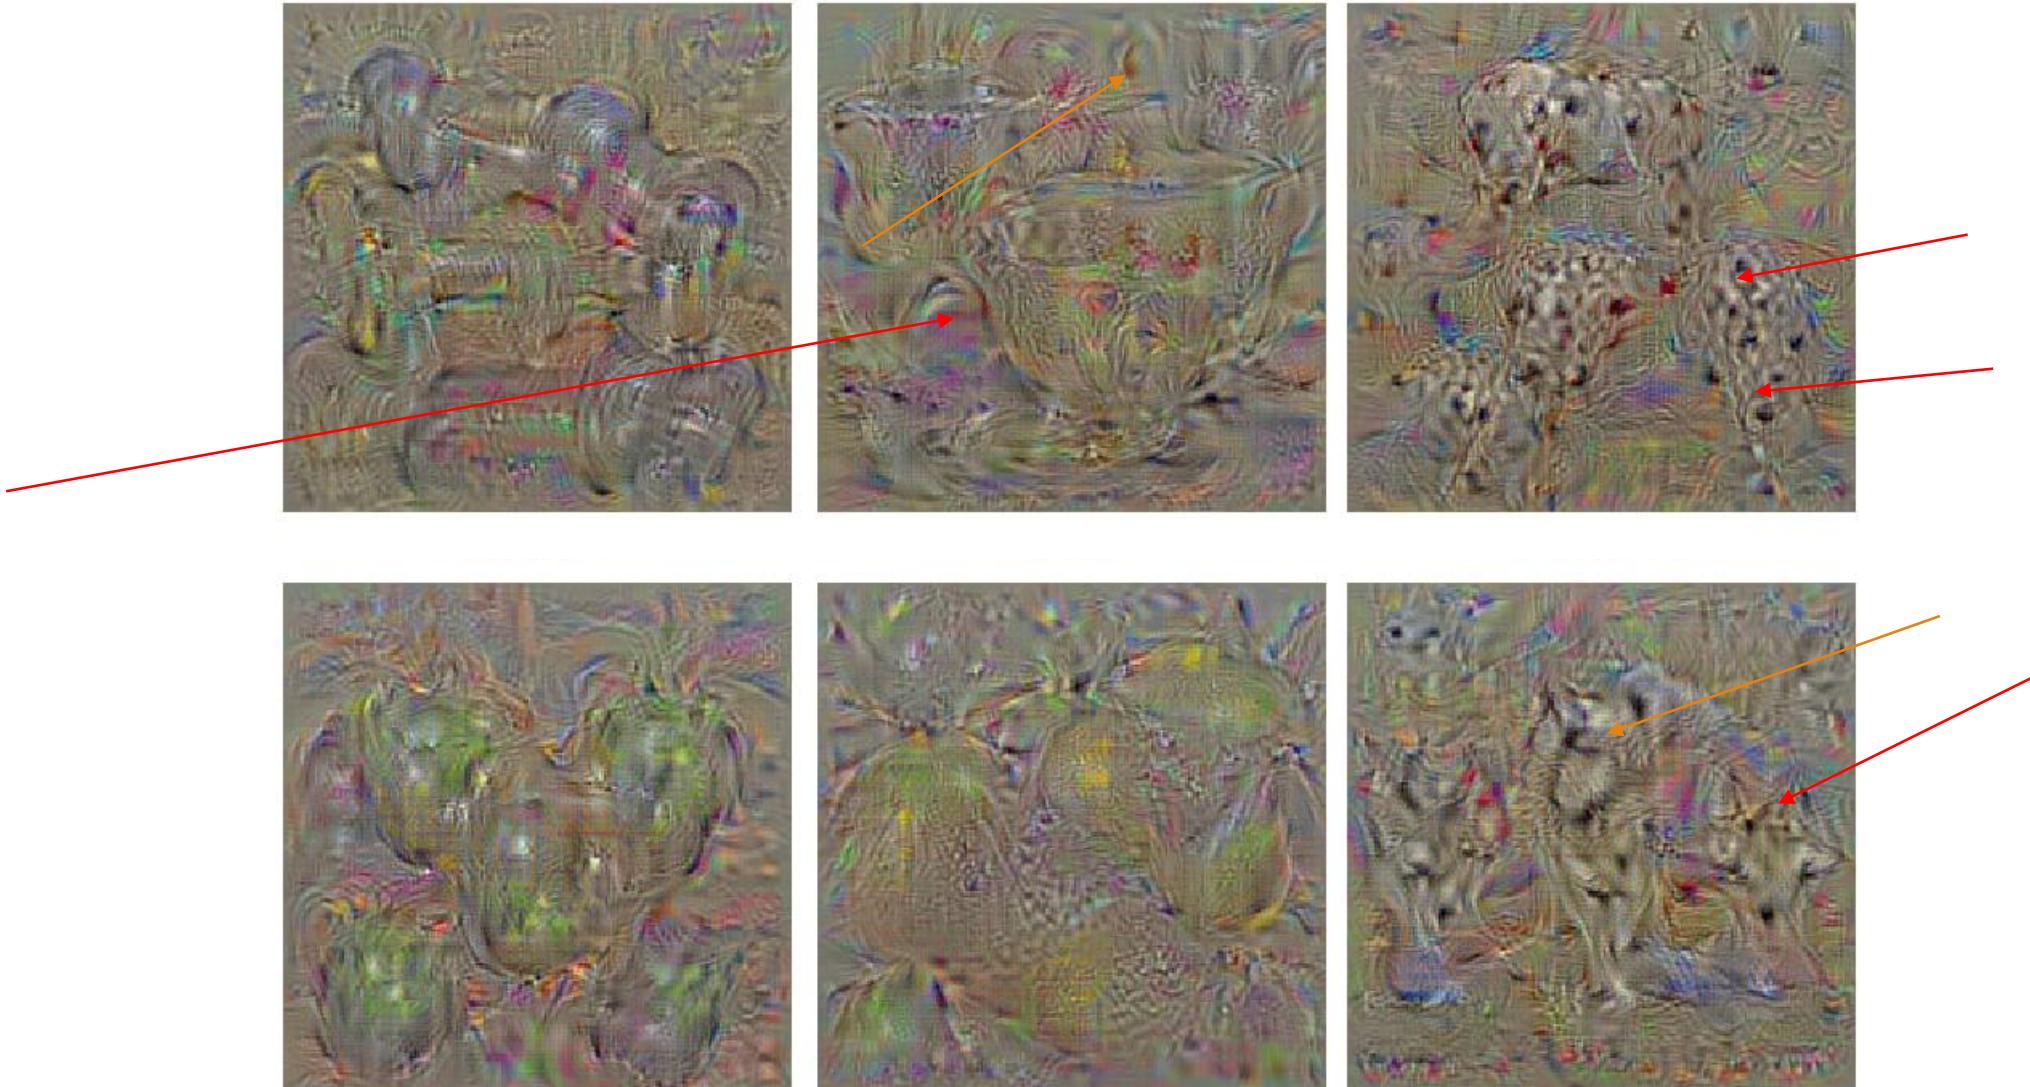
What is a “Convnet dog”, however? [Simonyan2014]

- We could understand what is “dog” is finding what image would have the “dog score”

$$\arg \max_I S_c(I; \theta) - \lambda |I|^2$$

- The parameters θ are fixed during the training
- Optimization is done with respect to the image I
- Initialized with the “average training image”

Maximum-scoring class images



Maximum-scoring class images



Class-specific image saliency

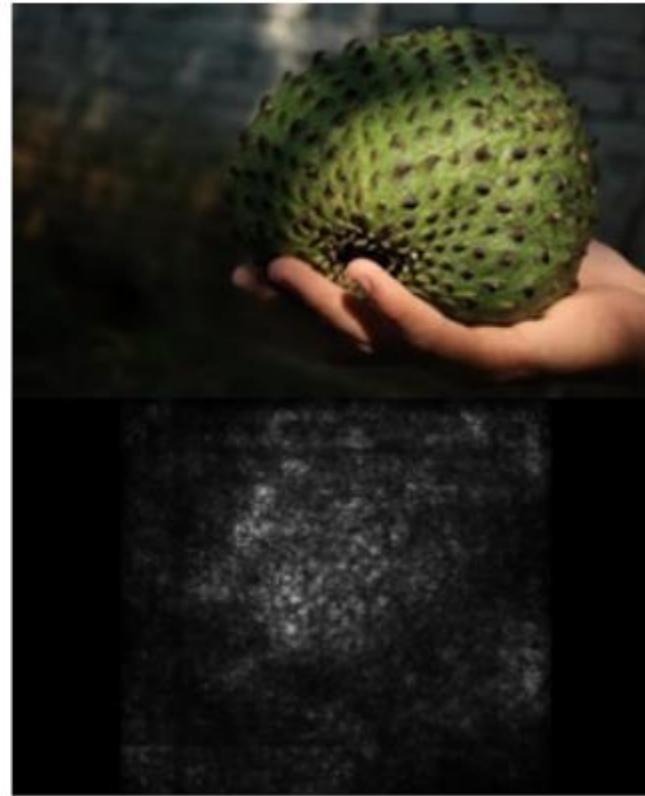
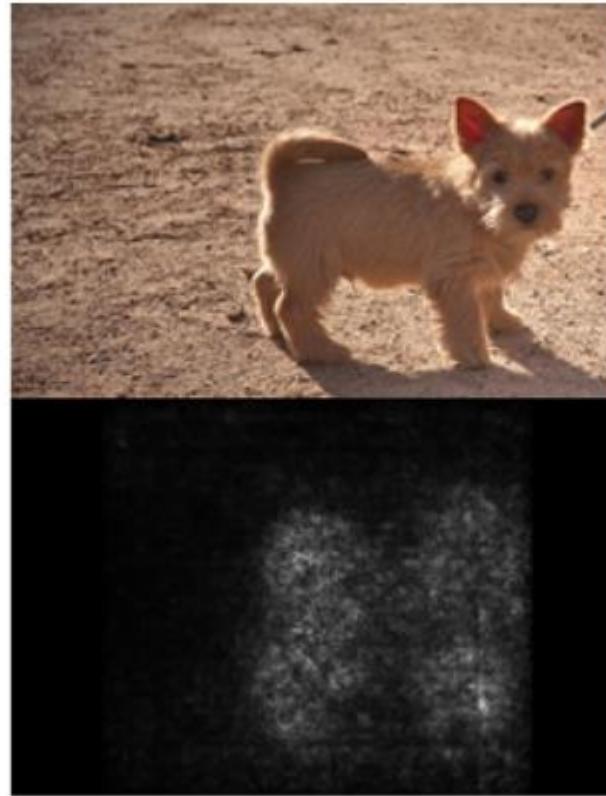
- Given the “monkey” class, what are the most “monkey-ish” parts in my image?
- Approximate S_c around an initial point I_0 with the first order Taylor expansion

$$S_c(I)|_{I_0} \approx w^T I + b, \text{ where } w = \frac{\partial S_c}{\partial I}|_{I_0} \text{ from backpropagation}$$

- Solution is locally optimal



Examples

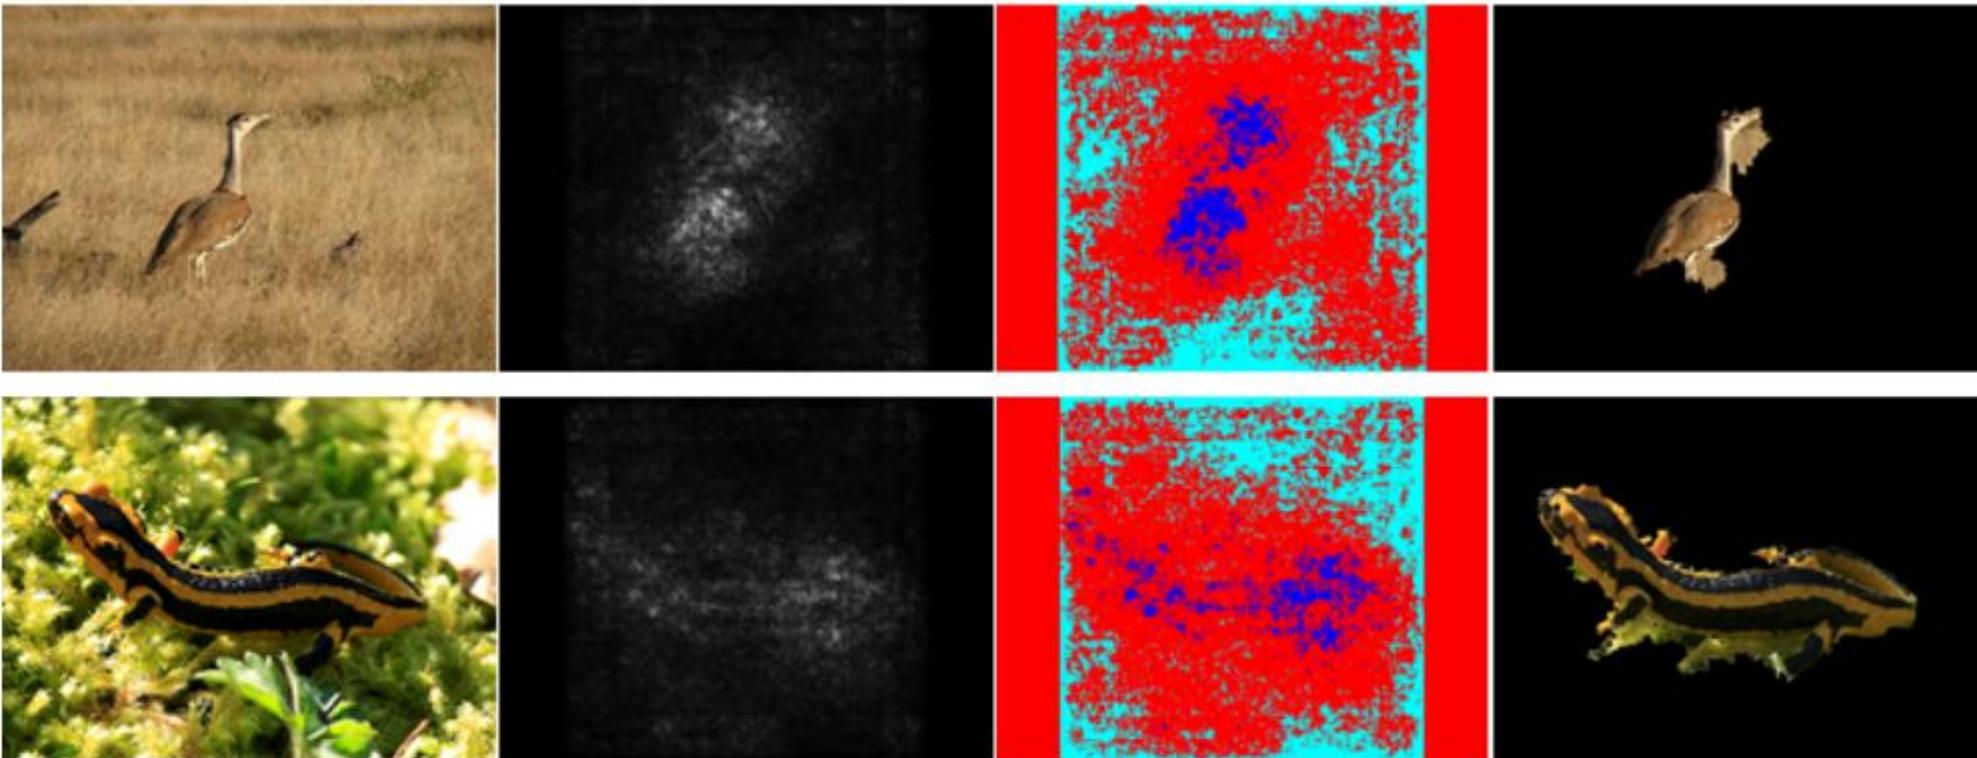


Examples

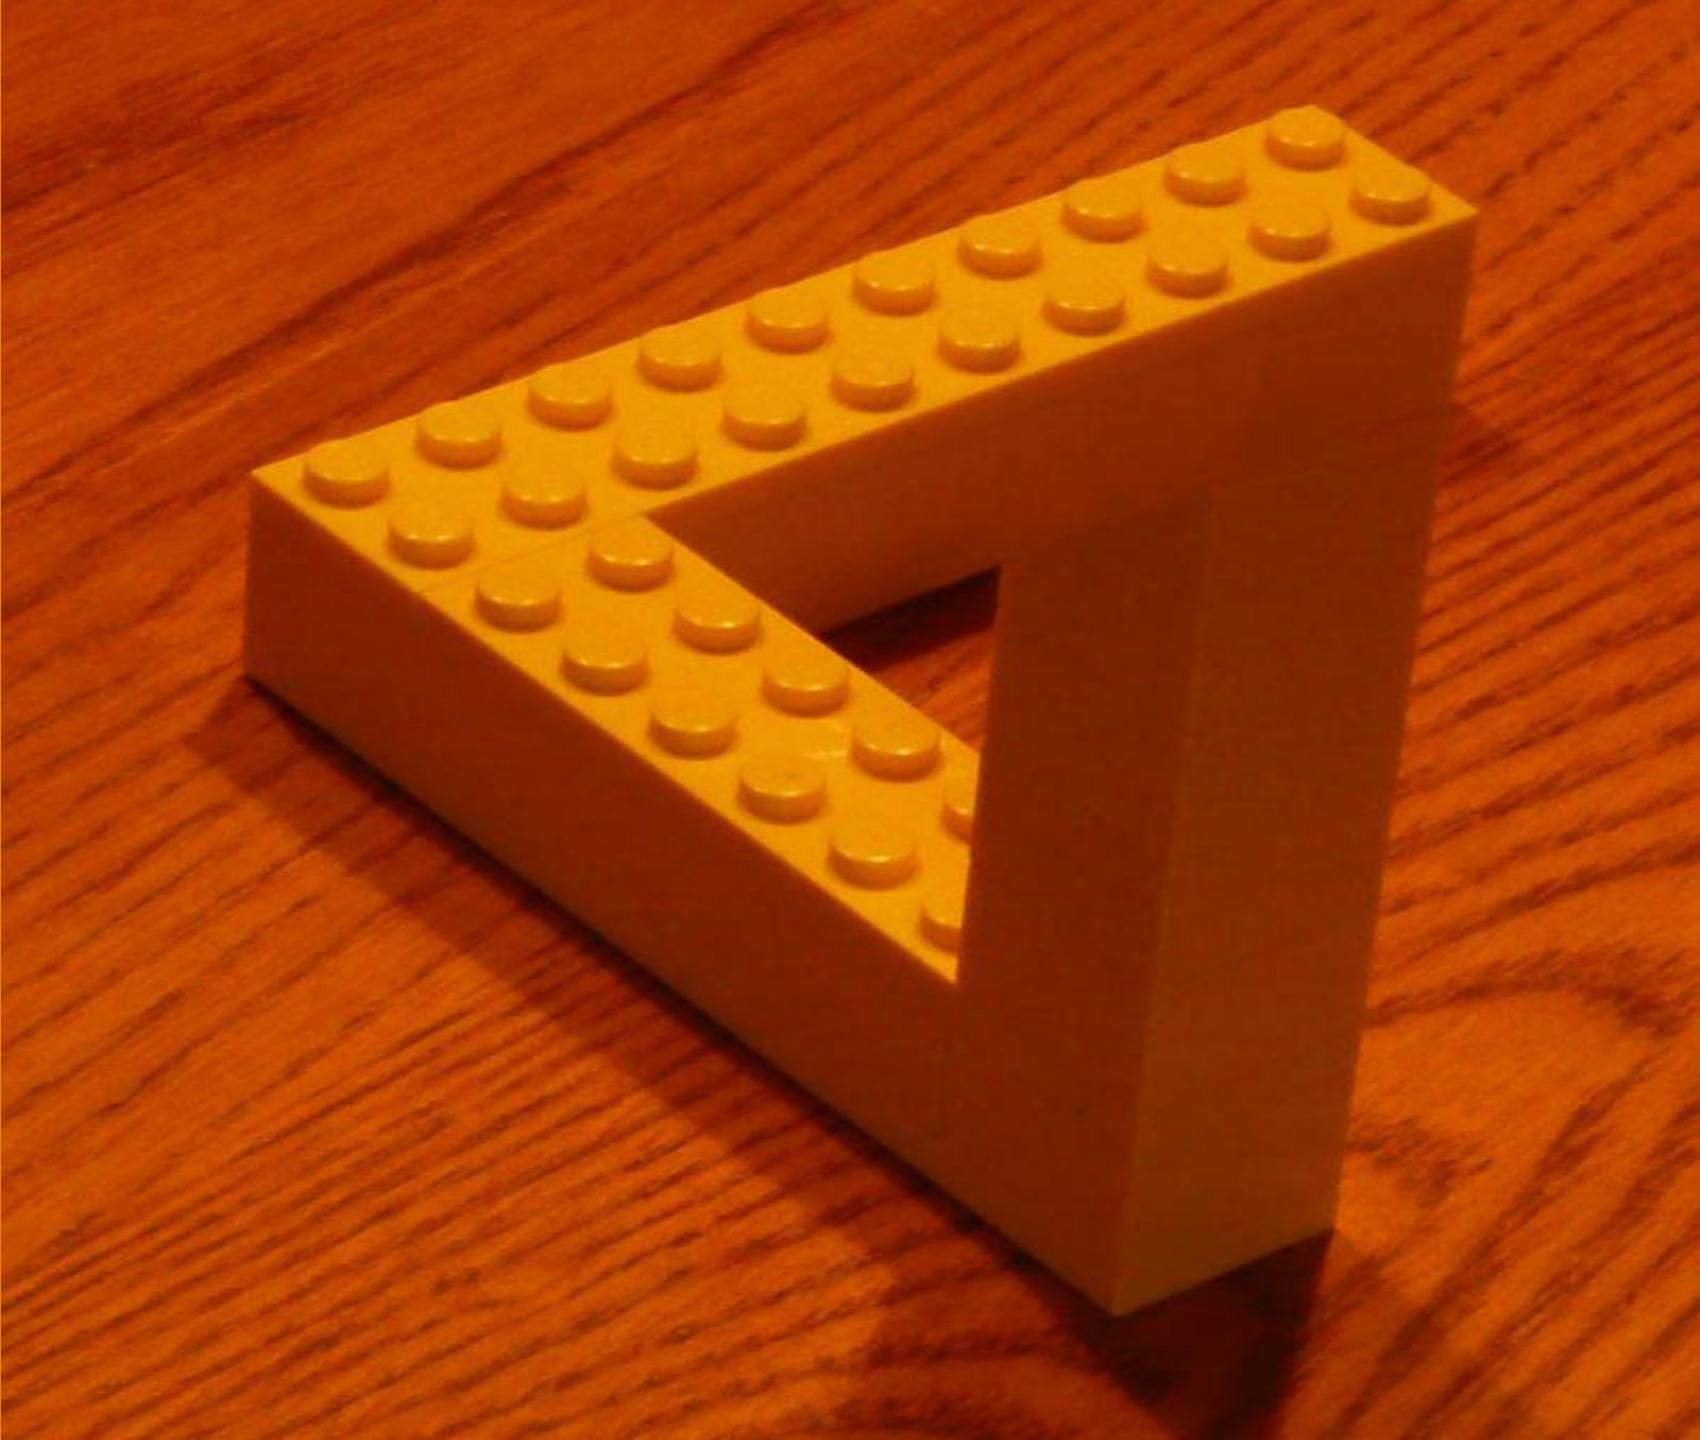


Object localization using class saliency

Grabcut



Fooling a Convnet



Fooling images

- What if we follow a similar procedure but with a different goal
- Generate “visually random” images
 - Images that make a lot of sense to a Convnet but no sense at all to us
- Or, assume we make very small changes to a picture (invisible to the naked eye)
 - Is a convnet always invariant to these changes?
 - Or could it be fooled?

Smoothness assumption

- We assume our classifiers enjoy **local generalization**
- Assume an image containing a cat lies at coordinates x
- Assume also a small change r around x , such that $x + r < \varepsilon$
 - ε is a very small constant
- Is the smoothness assumption reasonable?
- Or can we “break” it by some adversarial examples
 - E.g. if we correctly classify x as “*Argan goat*”, with the right r make the convnet see a “*BMW i6*”
 - The $x + r$ would be adversarial examples



Generating adversarial examples

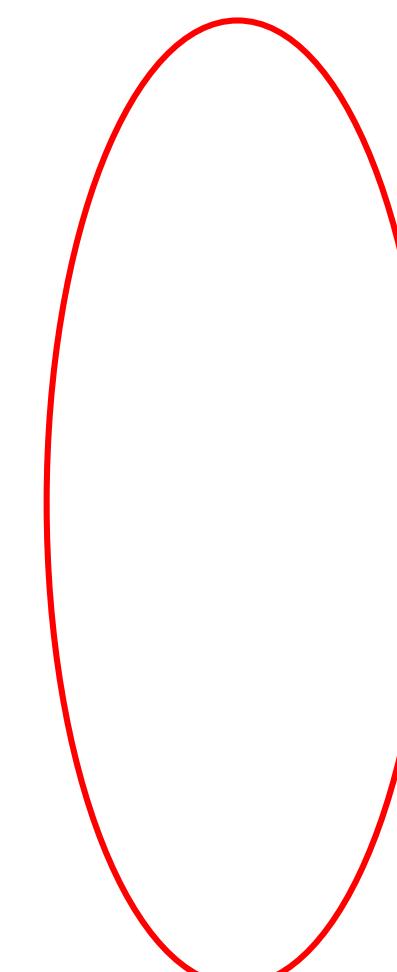
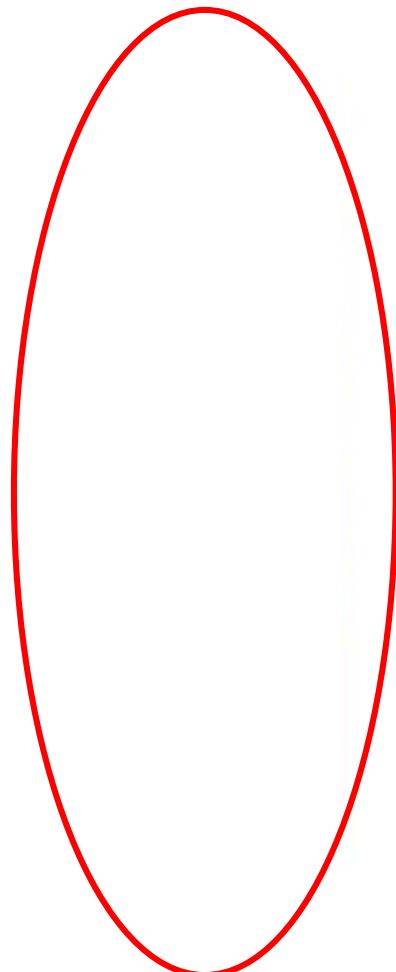
- If $f: \mathcal{R}^m \rightarrow \{1, \dots, k\}$ our goal can be mathematically described as

$$\min \|r\|^2$$

$$s.t. \quad f(x + r) = l, x + r \in [0, 1]^m$$

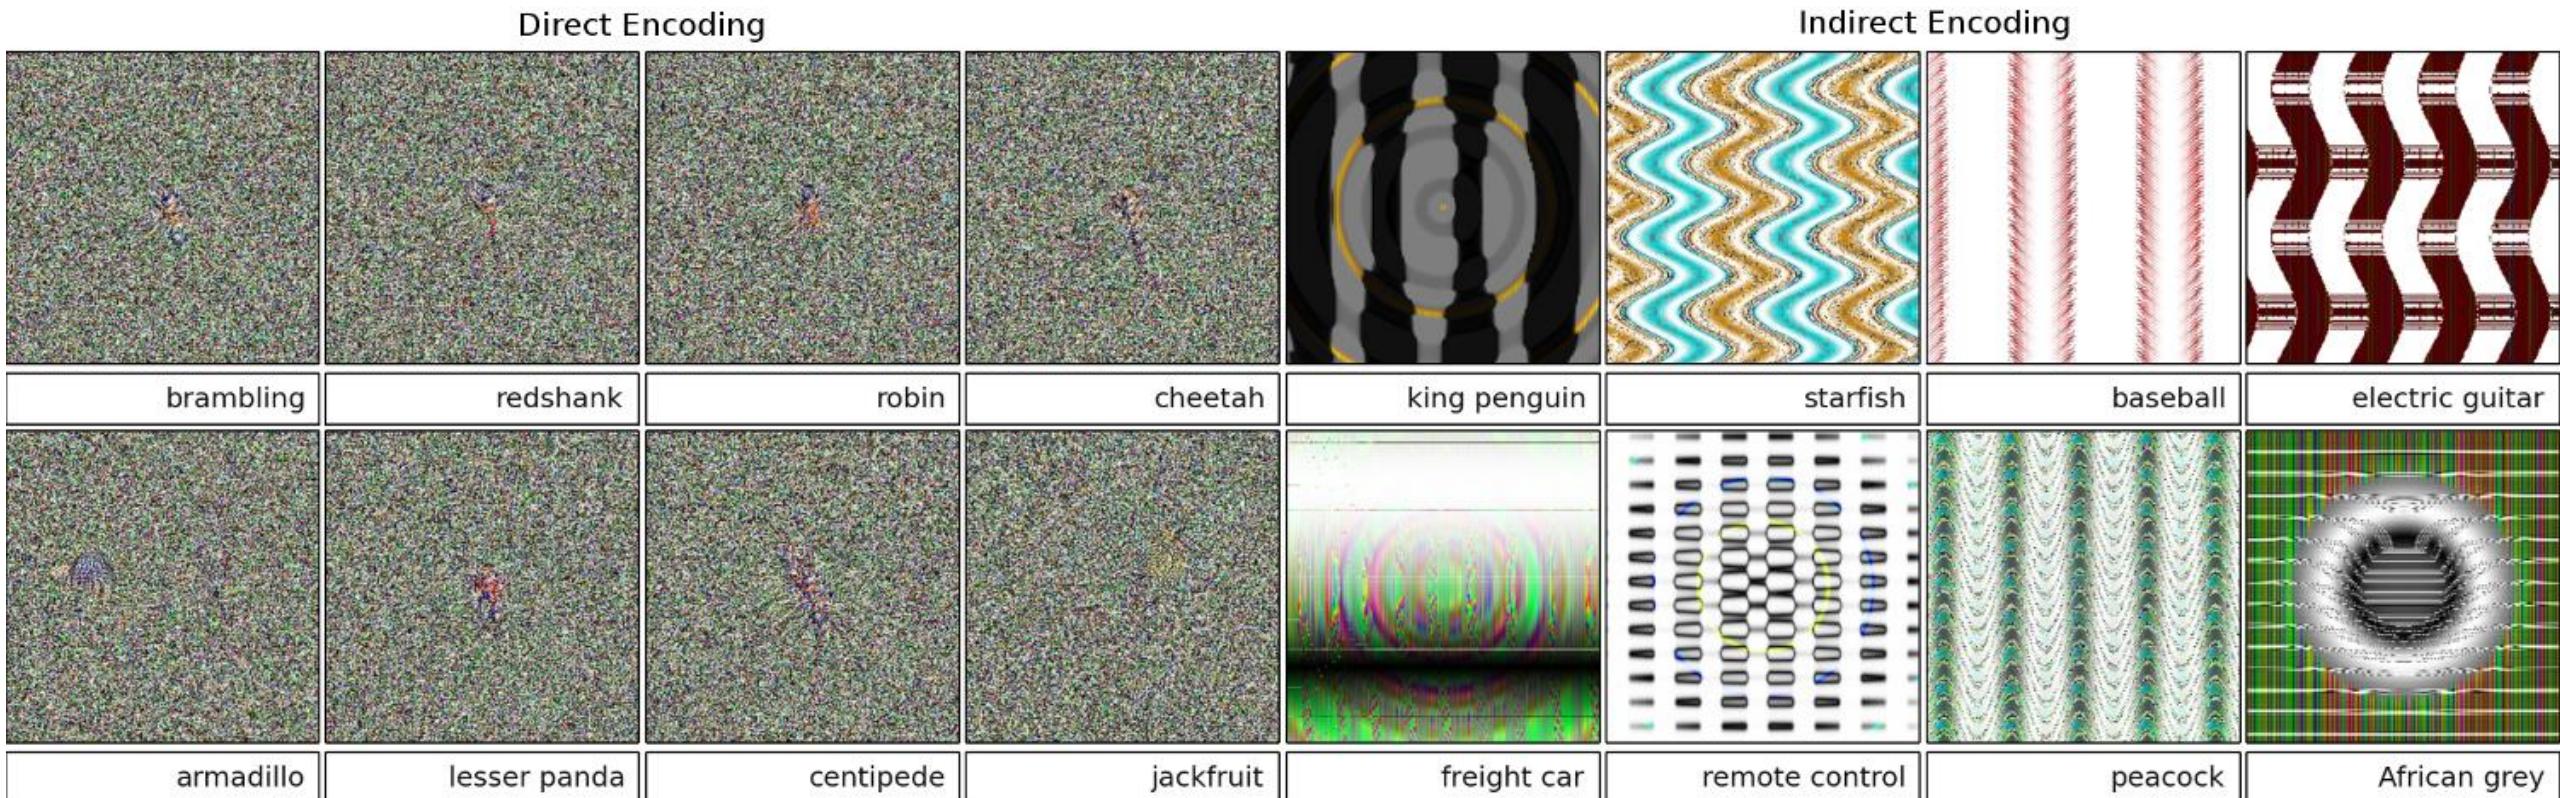
- The goal is to optimize the distortion r so that the predicted label l is different than the original label

Adversarial images



Predicted as Ostrich, *Struthio camelus*

More adversarial images



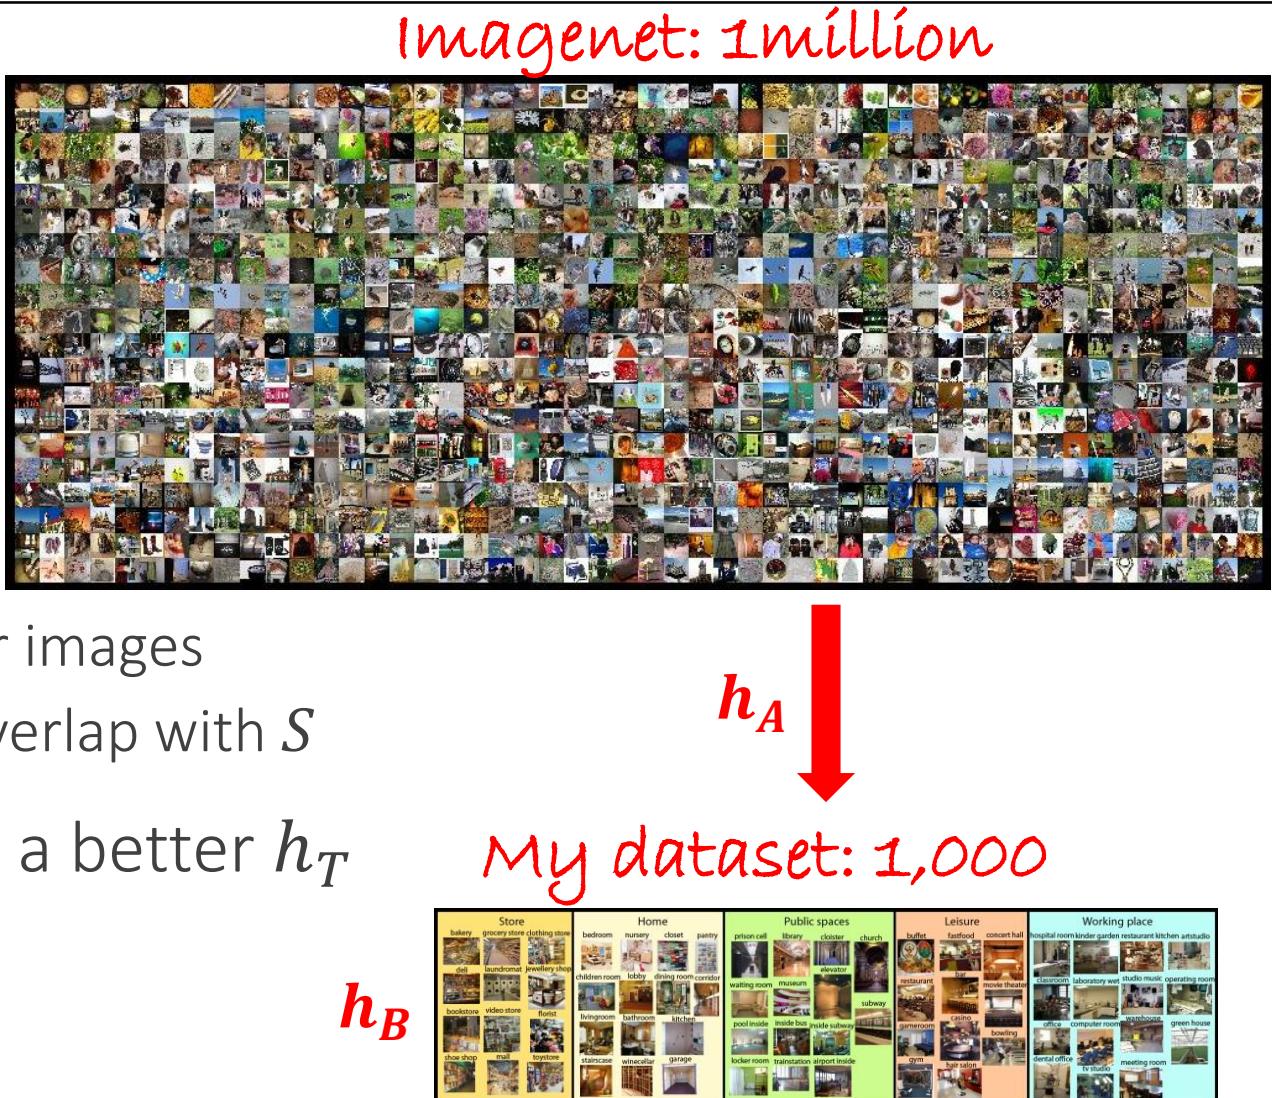
Knowledge transfer

CNNs and dataset size

- A CNN can have millions of parameters
- What about the dataset size?
- Could we still train a CNN without overfitting problems?

Transfer learning

- Assume two datasets, T and S
- Dataset S is
 - fully annotated, plenty of images
 - We can build a model h_S
- Dataset T is
 - Not as much annotated, or much fewer images
 - The annotations of T do not need to overlap with S
- We can use the model h_S to learn a better h_T
- This is called transfer learning



Convnets are good in transfer learning

- Even if our dataset T is not large, we can train a CNN for it
- Pre-train a network on the dataset S
- Then, there are two solutions
 - Fine-tuning
 - CNN as feature extractor

Solution I: Fine-tune h_T using h_S as initialization

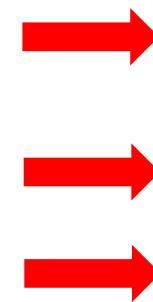
- Assume the parameters of S are already a good start near our final local optimum
- Use them as the initial parameters for our new CNN for the target dataset

$$\theta_{T,l}^{(t=0)} = \theta_{S,l} \text{ for some layers } l = 1, 2, \dots$$

- This is a good solution when the dataset T is relatively big
 - E.g. for Imagenet S with approximately 1 million images
 - For a dataset T with more than a few thousand images should be ok
- What layers to initialize and how?

Initializing h_T with h_S

- Classifier layer to loss
 - The loss layer essentially is the “classifier”
 - Same labels → keep the weights from h_S
 - Different labels → delete the layer and start over
 - When too few data, fine-tune only this layer
- Fully connected layers
 - Very important for fine-tuning
 - Sometimes you need to completely delete the last before the classification layer if datasets are very different
 - Capture more semantic, “specific” information
 - Always try first when fine-tuning
 - If you have more data, fine-tune also these layers



Classifier layer fc8

Fully connected layer fc7

Fully connected layer fc6

Convolutional Layer 5

Convolutional Layer 4

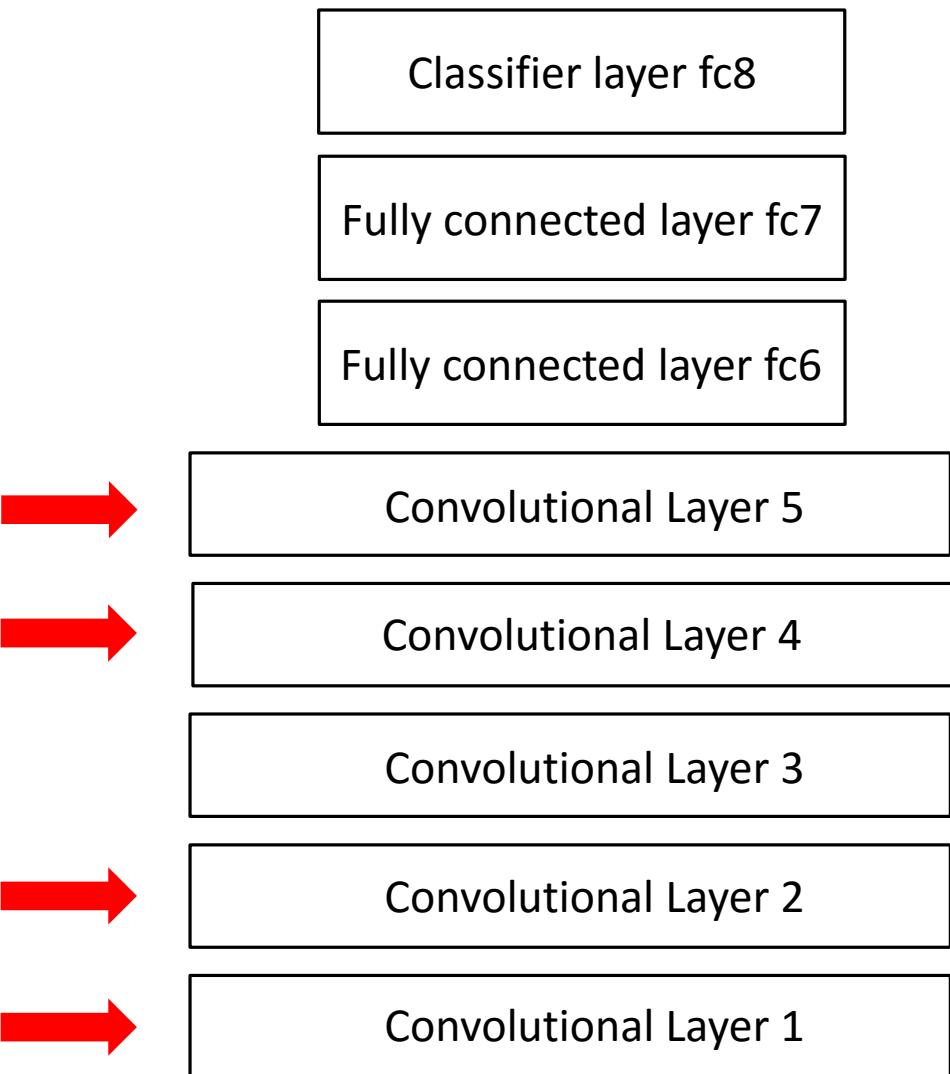
Convolutional Layer 3

Convolutional Layer 2

Convolutional Layer 1

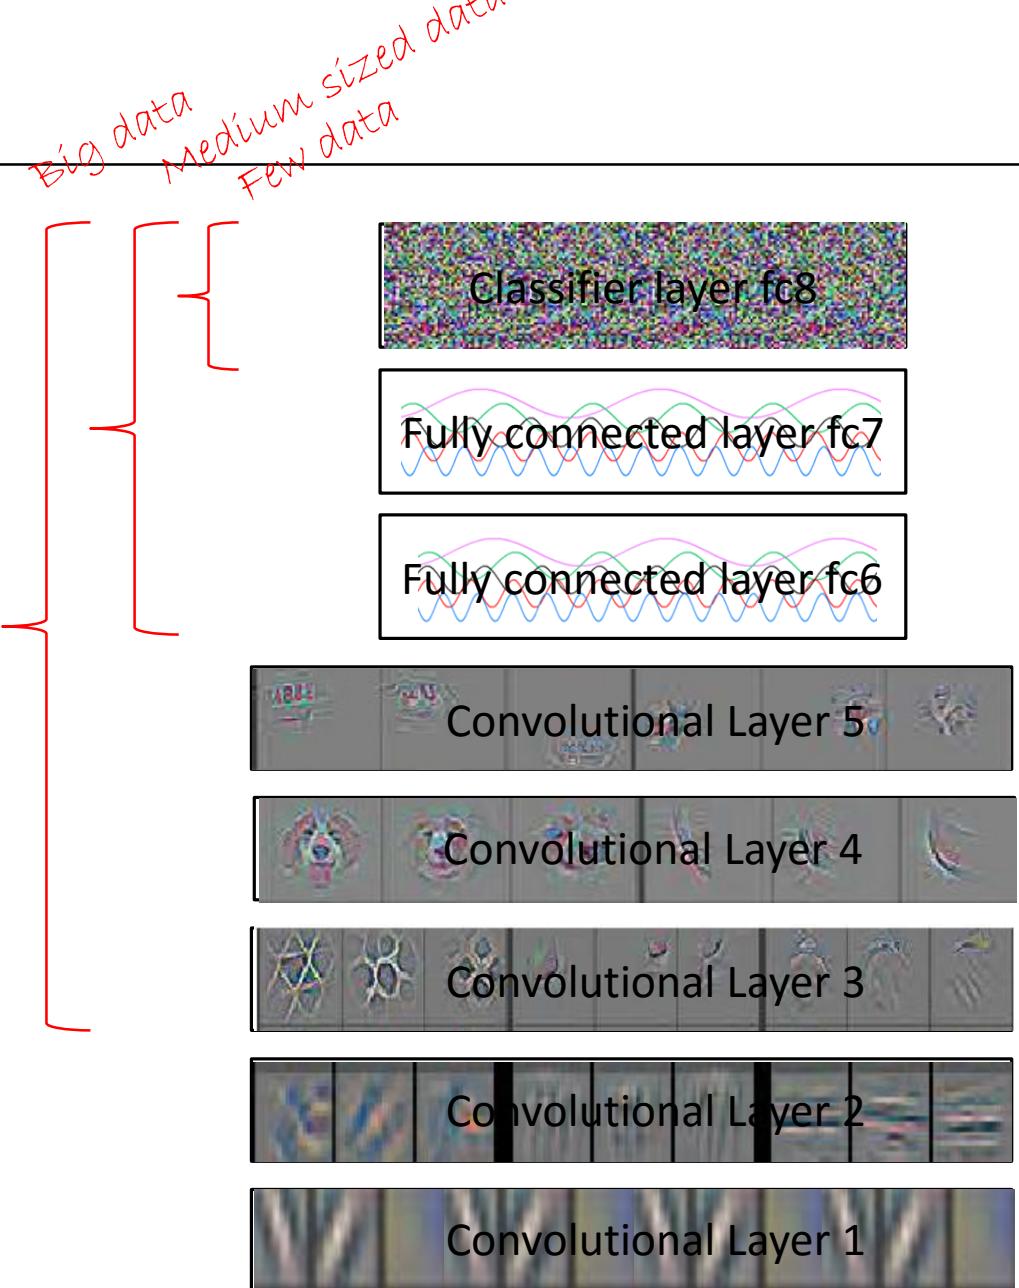
Initializing h_T with h_S

- Upper convolutional layers (conv4, conv5)
 - Mid-level spatial features (face, wheel detectors ...)
 - Can be different from dataset to dataset
 - Capture more generic information
 - Fine-tuning pays off
 - Fine-tune if dataset is big enough
- Lower convolutional layers (conv1, conv2)
 - They capture low level information
 - This information does not change usually
 - Probably, no need to fine-tune but no harm trying



How to fine-tune?

- For layers initialized from h_S use a mild learning rate
 - Remember: your network is already close to a near optimum
 - If too aggressive, learning might diverge
 - A learning rate of 0.001 is a good starting choice (assuming 0.01 was the original learning rate)
- For completely new layers (e.g. loss) use aggressive learning rate
 - If too small, the training will converge very slowly
 - Remember: the rest of the network is near a solution, this layer is very far from one
 - A learning rate of 0.01 is a good starting choice
- If datasets are very similar, fine-tune only fully connected layers
- If datasets are different and you have enough data, fine-tune all layers

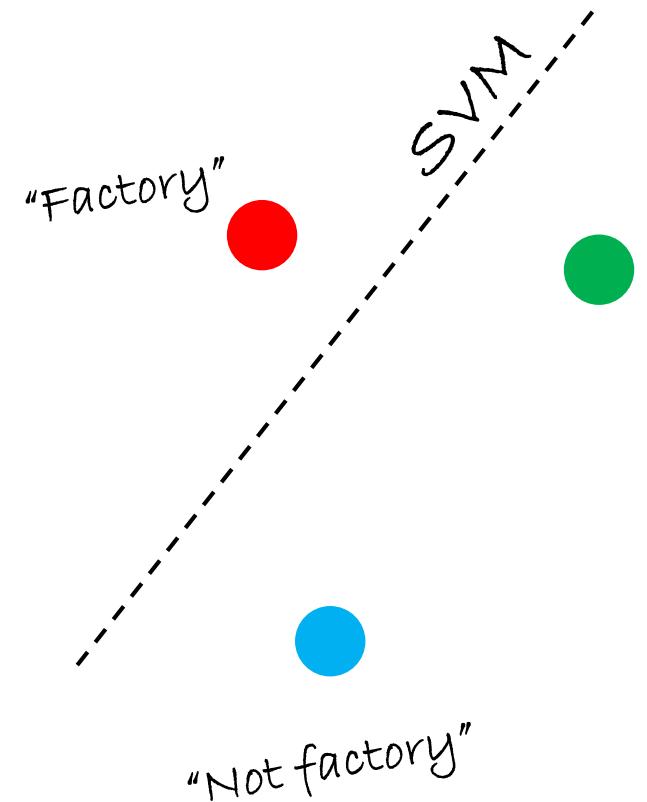
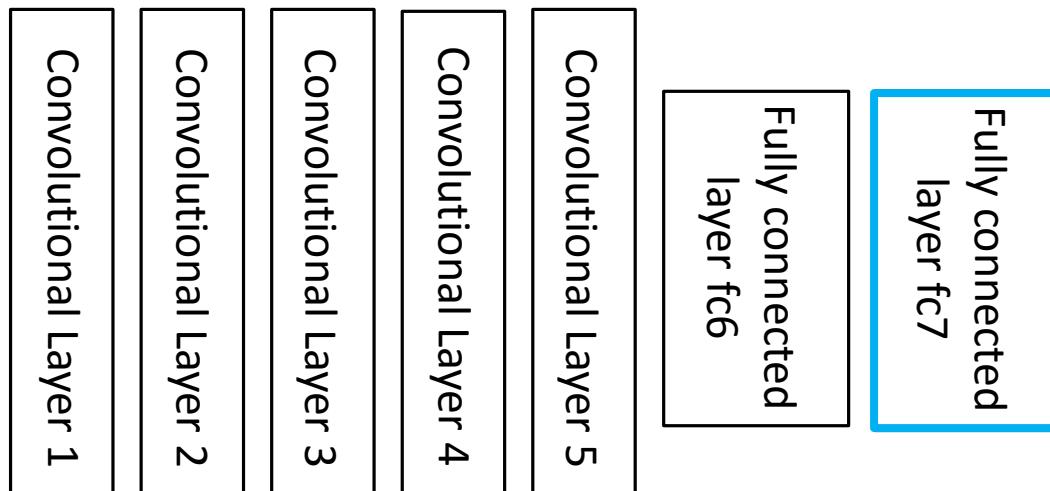


Solution II: Use h_S as a feature extractor for h_T

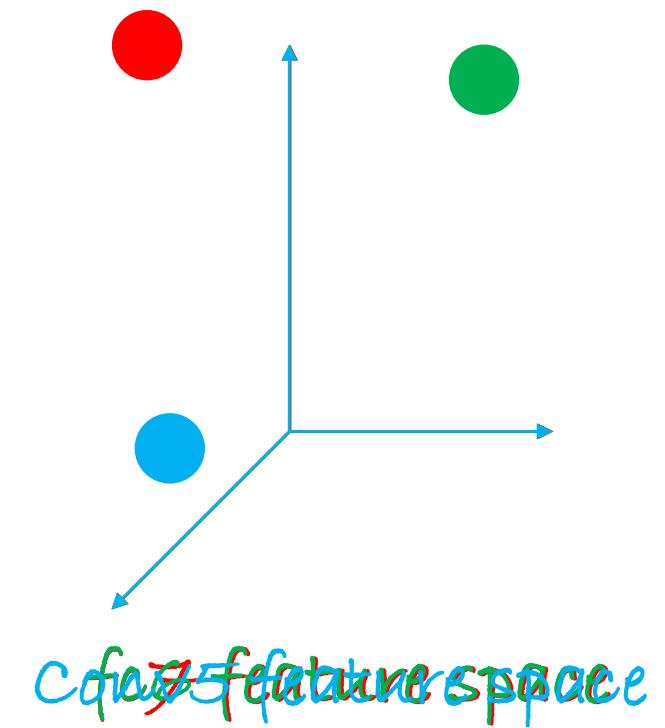
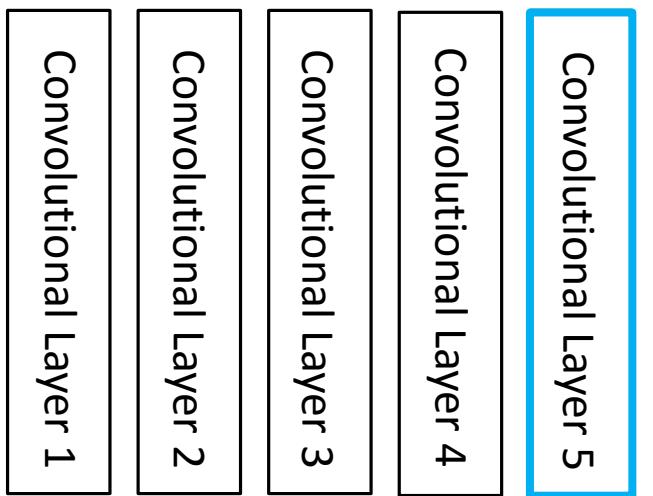
- Essentially similar to a case of solution I
 - but train only the loss layer
- Essentially use the network as a pretrained feature extractor
- This is a good solution if the dataset T is small
 - Any fine-tuning of layer might cause overfitting
- Or when we don't have the resources to train a deep net
- Or when we don't care for the best possible accuracy

Off-the-shelf classifiers on deep features

- E.g. extract “deep features” from the final layer



Deep features from different layers



Which layer?

Table 6. Analysis of the discriminative information contained in each layer of feature maps within our ImageNet-pretrained convnet. We train either a linear SVM or softmax on features from different layers (as indicated in brackets) from the convnet. Higher layers generally produce more discriminative features.

	Cal-101 (30/class)	Cal-256 (60/class)
SVM (1)	44.8 ± 0.7	24.6 ± 0.4
SVM (2)	66.2 ± 0.5	39.6 ± 0.3
SVM (3)	72.3 ± 0.4	46.0 ± 0.3
SVM (4)	76.6 ± 0.4	51.3 ± 0.1
SVM (5)	86.2 ± 0.8	65.6 ± 0.3
SVM (7)	85.5 ± 0.4	71.7 ± 0.2
Softmax (5)	82.9 ± 0.4	65.7 ± 0.5
Softmax (7)	85.4 ± 0.4	72.6 ± 0.1

Higher layer features are capture more semantic information. Good → for higher level classification

Lower layer features capture more basic information (texture, etc). Good ← for image-to-image comparisons, image retrieval

Summary

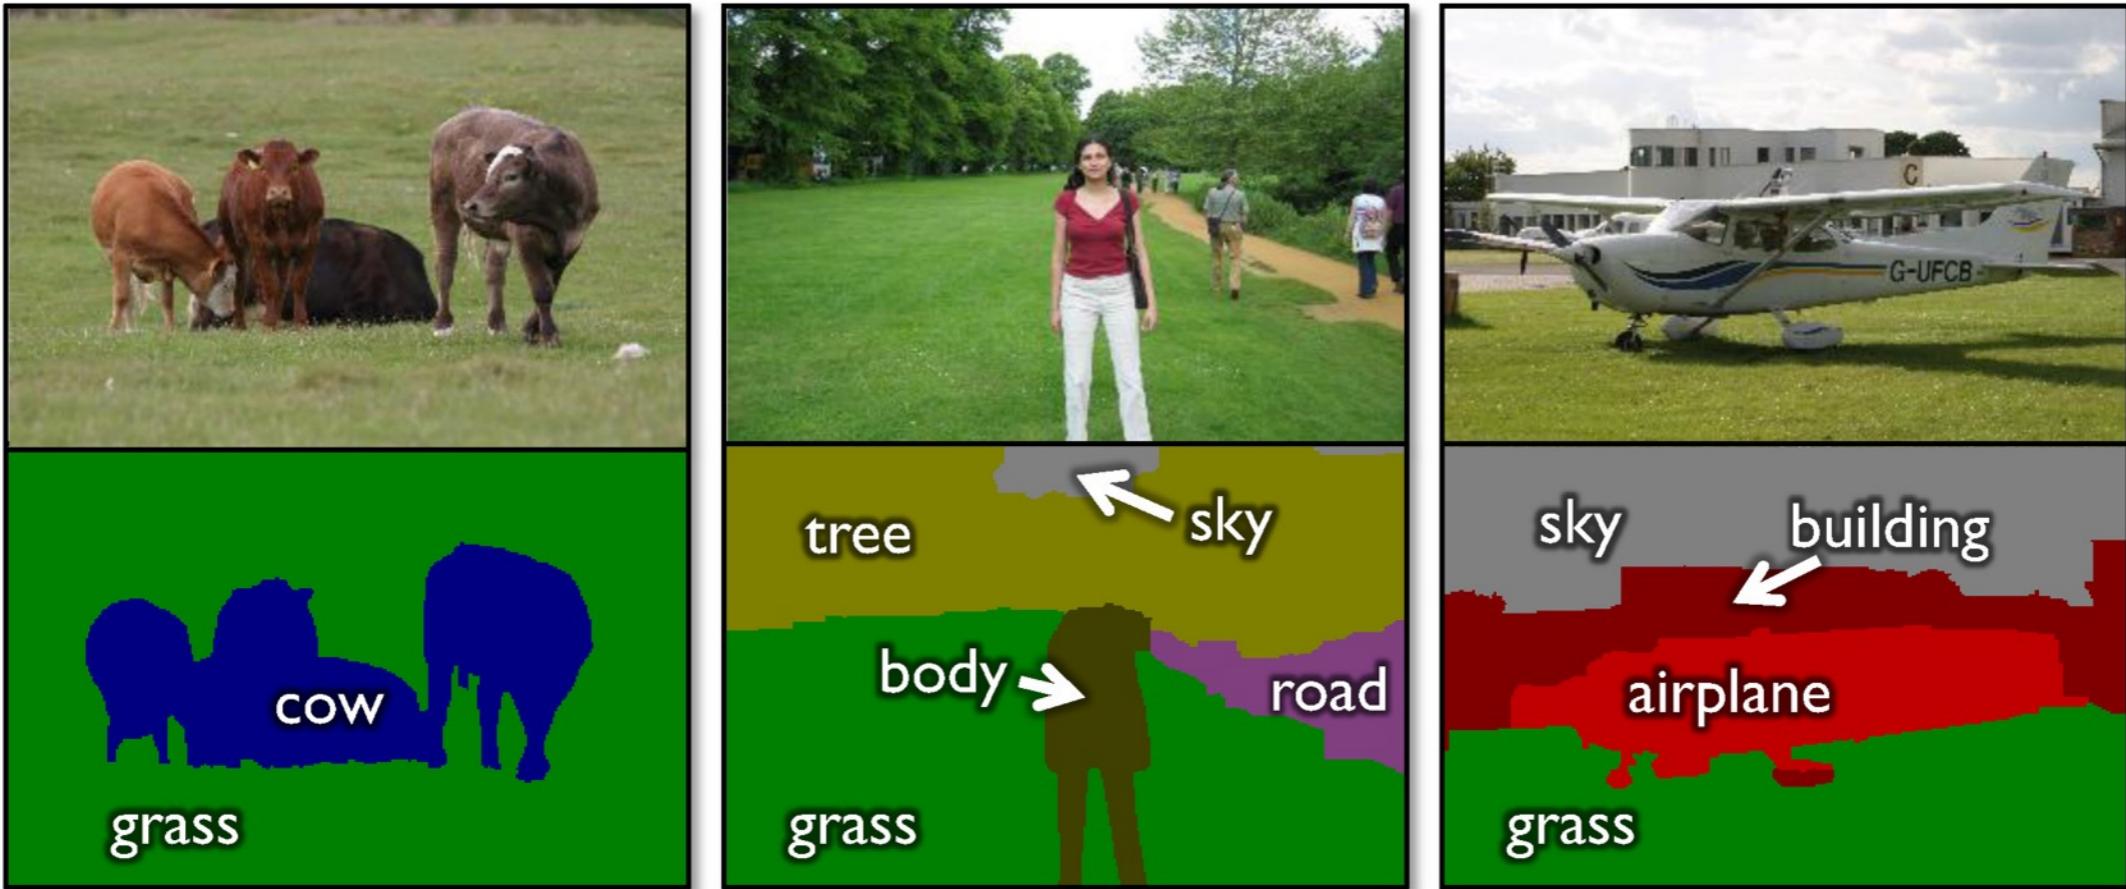
- What do convolutions look like?
- Build on the visual intuition behind Convnets
- Deep Learning Feature maps
- Transfer Learning

Next lecture

- Convolutional Neural Networks for Object Detection and Segmentation
- Convolutional Neural Networks and Structured Prediction

References

- [Zeiler2014] Visualizing and Understanding Convolution Networks, Zeiler and Fergus, ECCV 2014



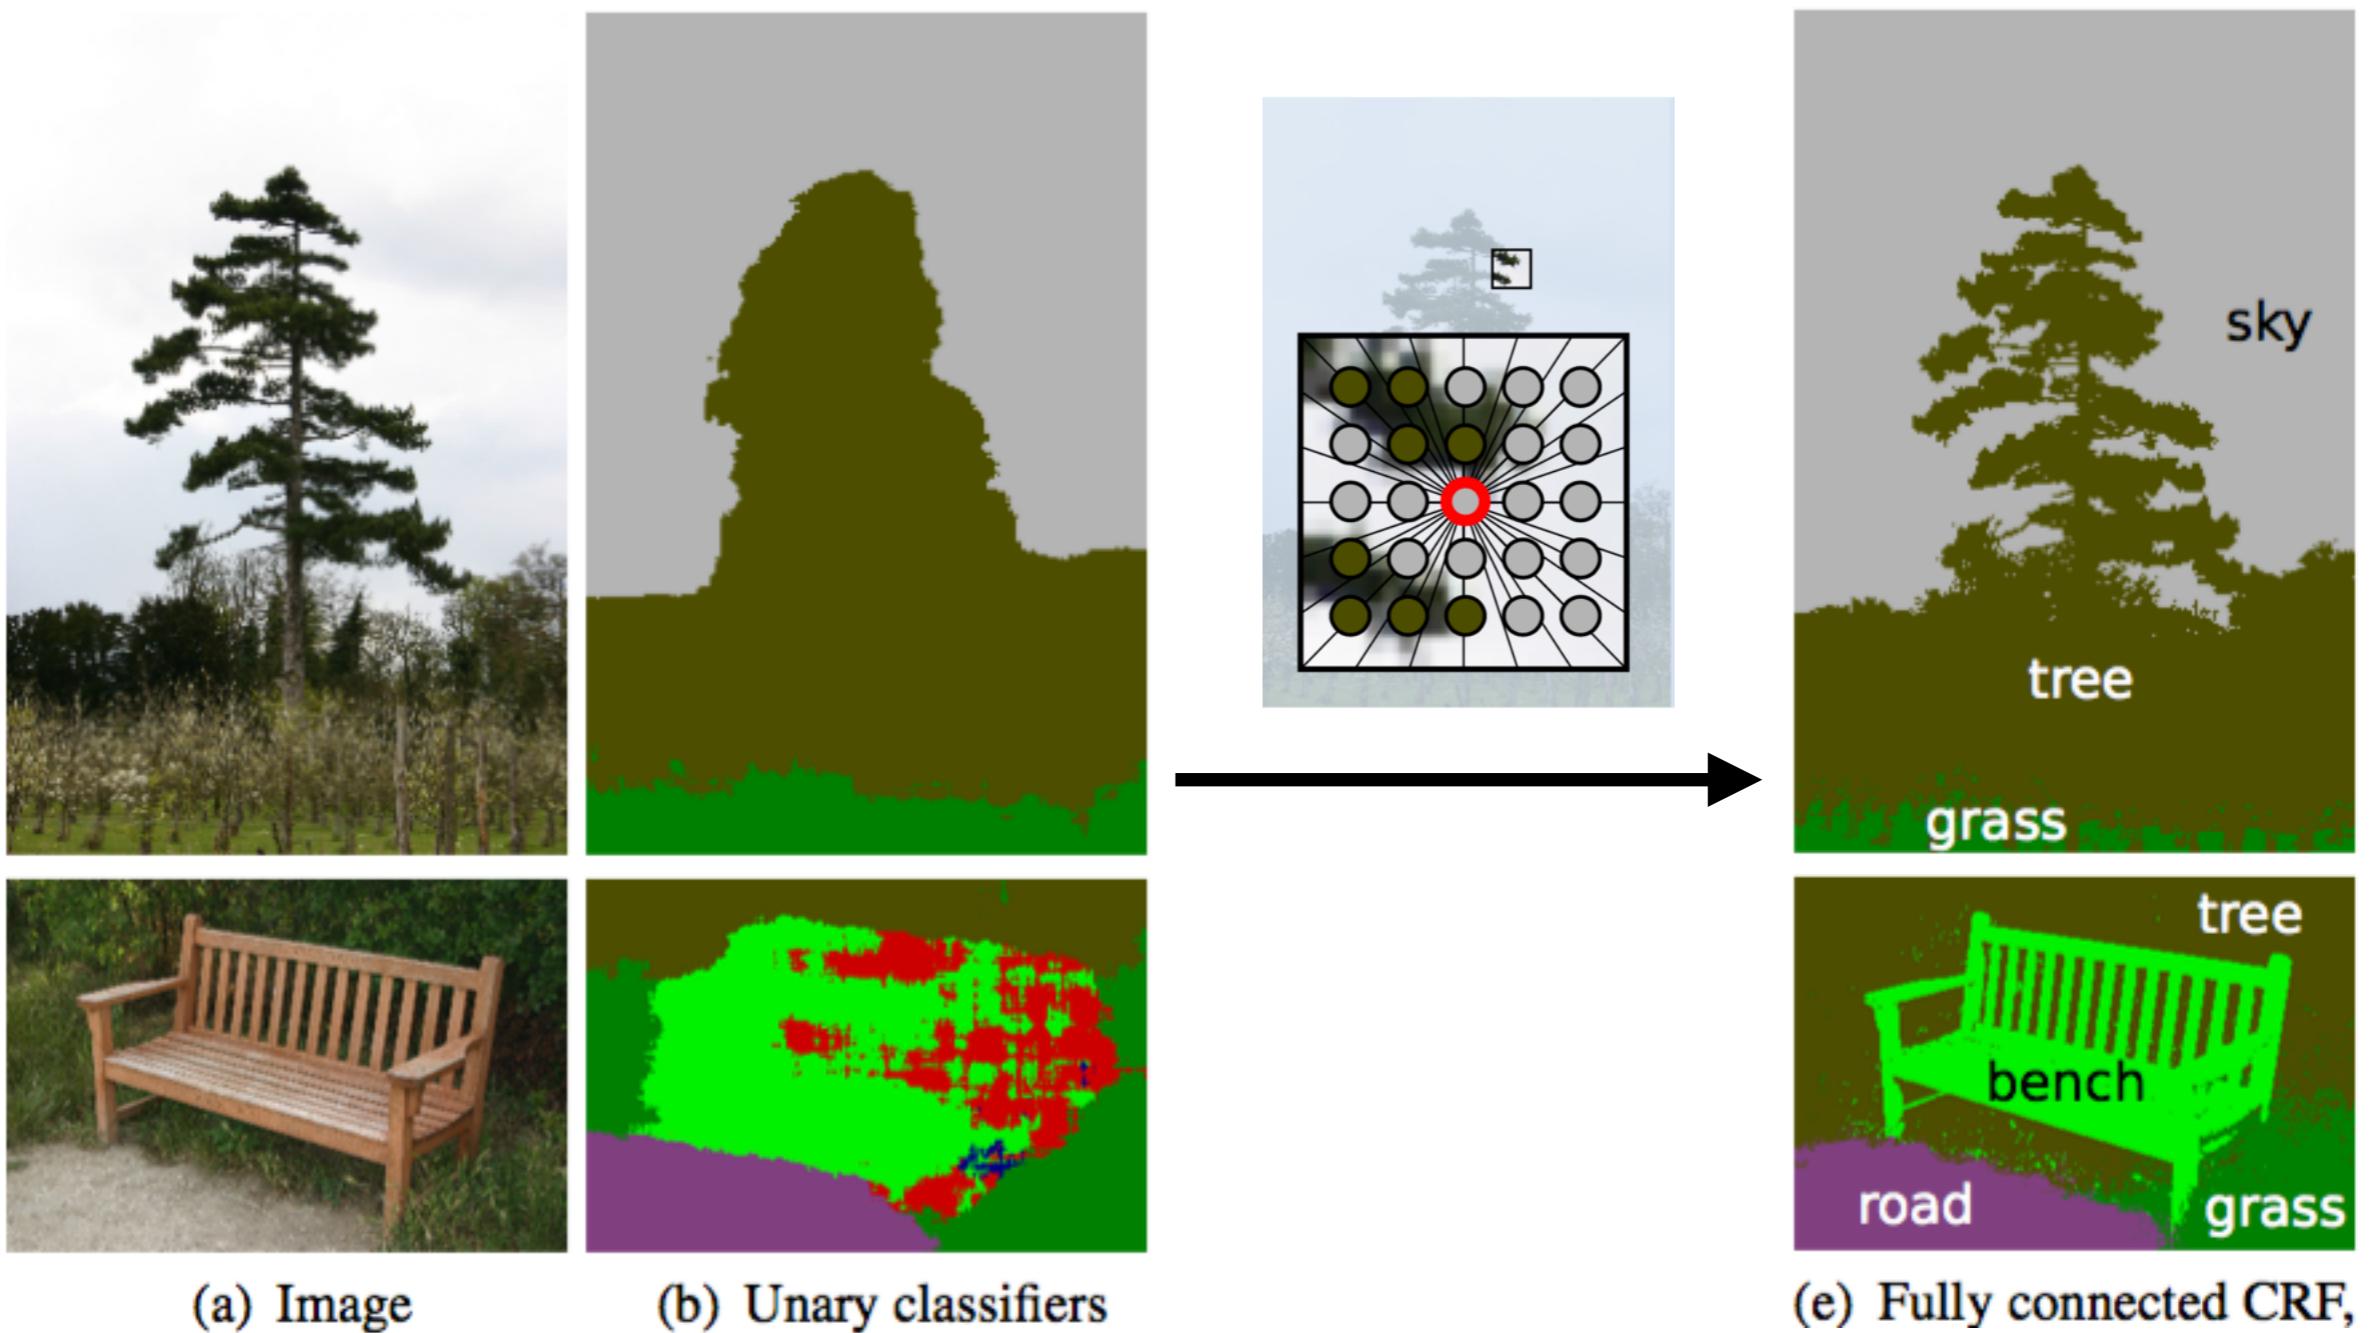
object classes	building	grass	tree	cow	sheep	sky	airplane	water	face	car
bicycle	flower	sign	bird	book	chair	road	cat	dog	body	boat

[Shotton et al., 2009]

Image Segmentation using Conditional Random Fields

Deep Learning @ UvA
Patrick Putzky

Outlook

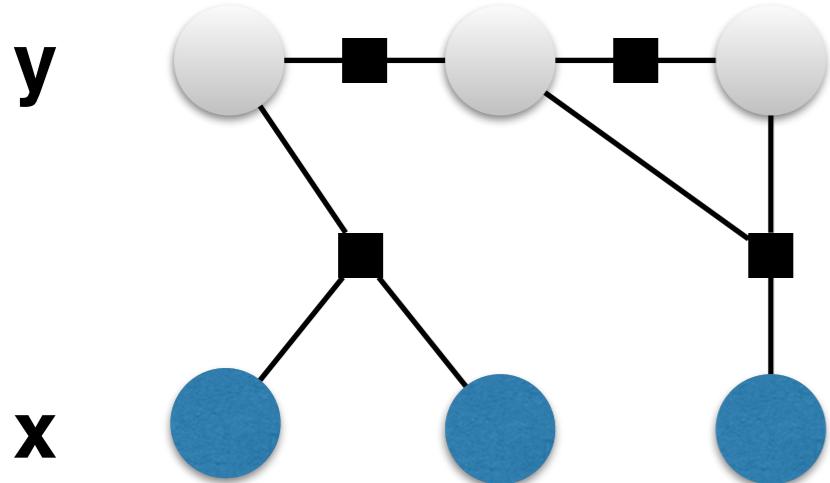


[Krähenbühl & Koltun, 2011]

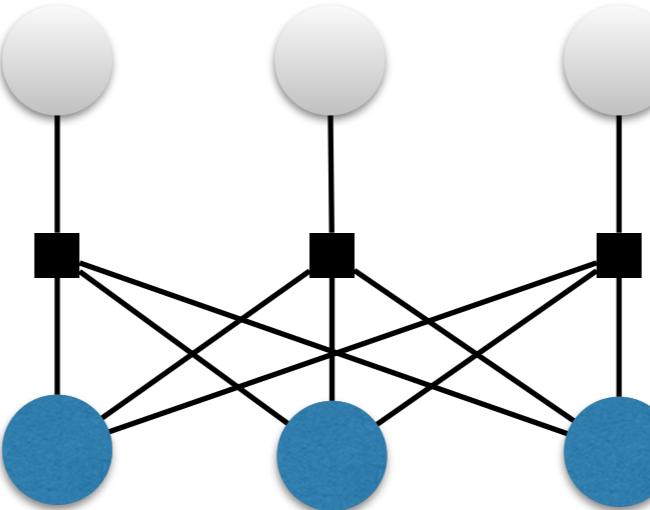
Conditional Random Fields

$$p(\mathbf{y}|\mathbf{x}) = \frac{1}{Z(\mathbf{x})} \exp(-E(\mathbf{y}|\mathbf{x}))$$

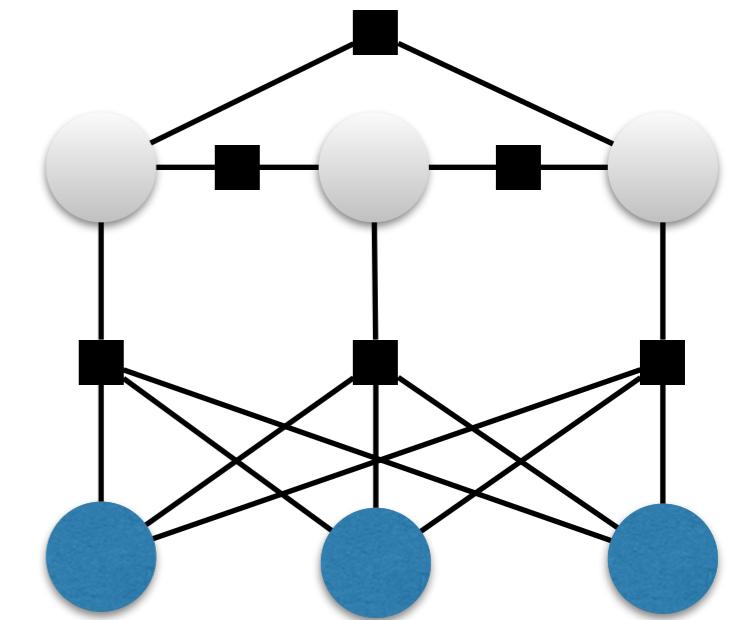
General form



Unary form



Pairwise CRF



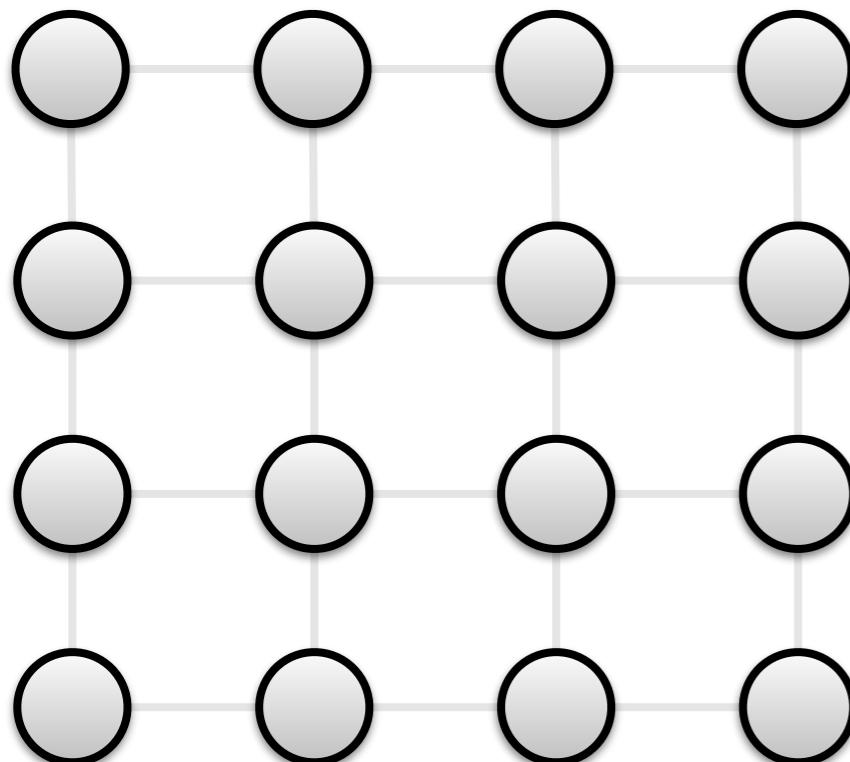
$$E(\mathbf{y}|\mathbf{x}) = \sum_{c \in \mathcal{C}_G} \psi_c(\mathbf{y}|\mathbf{x})$$

$$E(\mathbf{y}|\mathbf{x}) = \sum_i \psi_u(y_i|\mathbf{x})$$

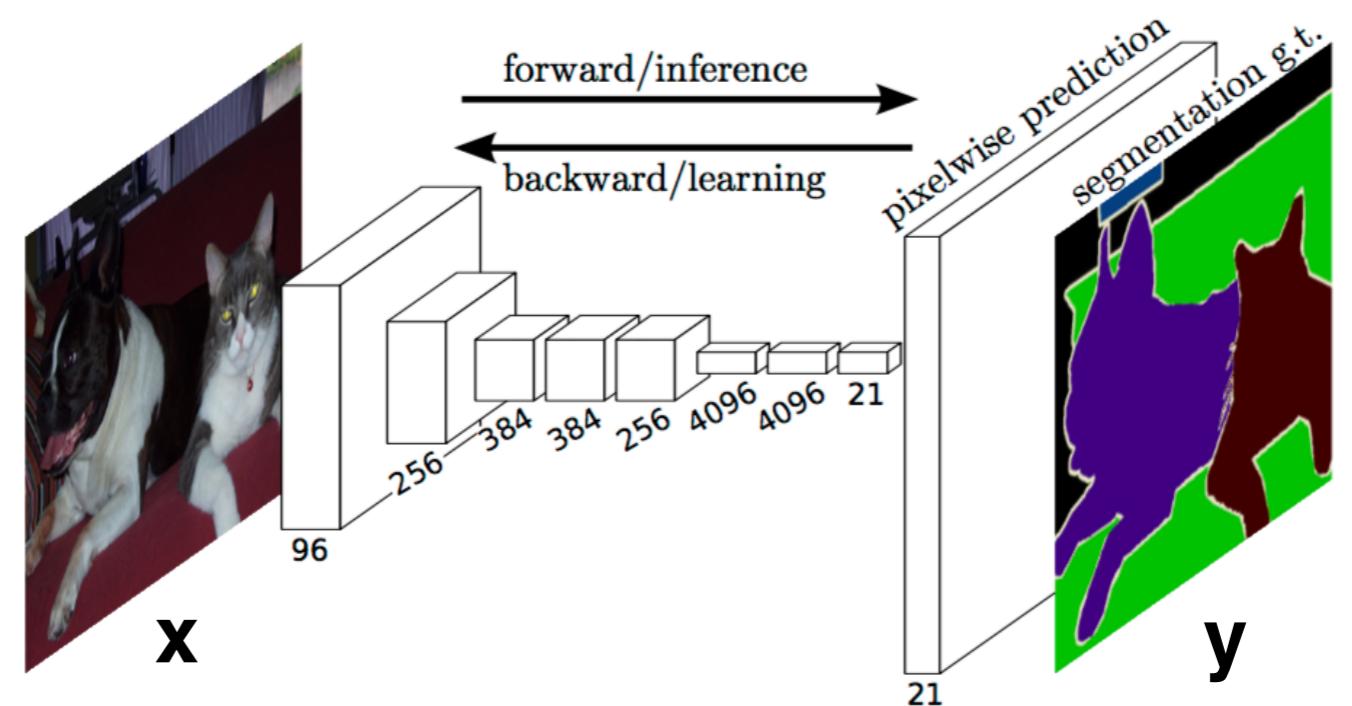
$$E(\mathbf{y}|\mathbf{x}) = \sum_i \psi_u(y_i|\mathbf{x}) + \sum_{i < j} \psi_p(y_i, y_j|\mathbf{x})$$

Unary models

$$E(\mathbf{y}|\mathbf{x}) = \sum_i \psi_u(y_i|\mathbf{x})$$



Example: Fully Convolutional Networks

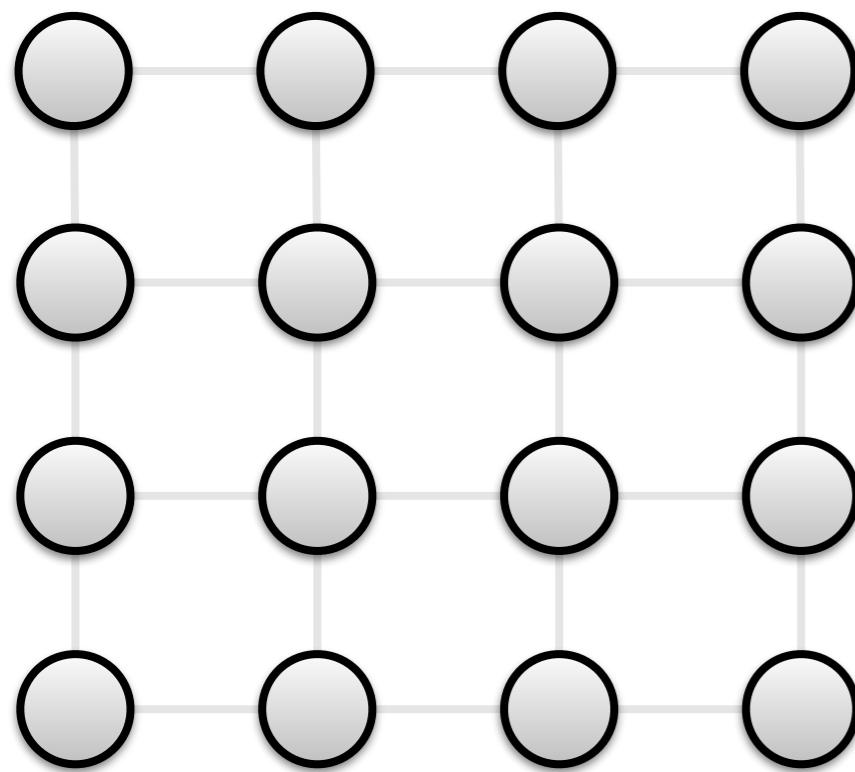


- Per pixel predictions
- No interactions between neighbouring class predictions

[Long et al., 2015]

Unary models

Fully Convolutional Networks: Fail Case



Input Image

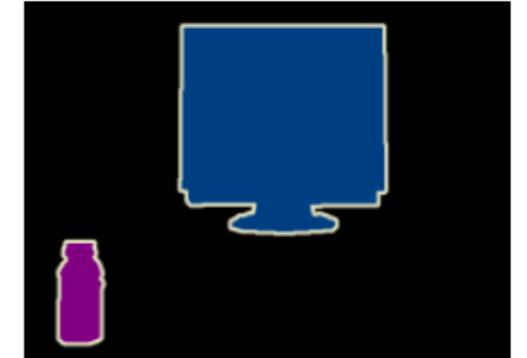


Aeroplane	Bottle	Bus
Cat	Dog	Horse
Person	Train	TV/Monitor

FCN-8s



Ground Truth

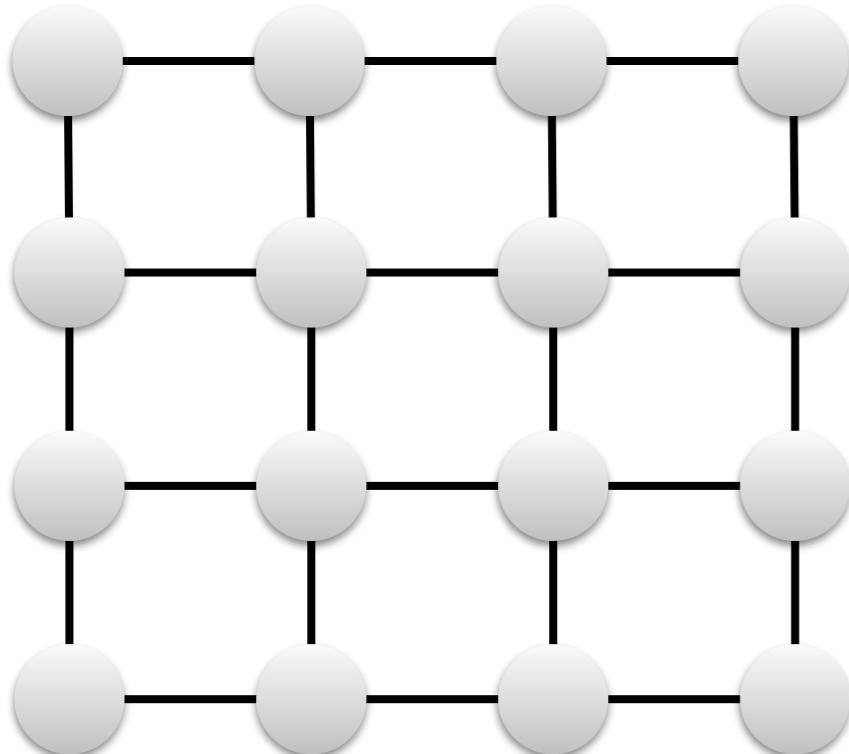


[Zheng et al., 2015]

- Per pixel predictions
- No interactions between neighbouring class predictions
- **No object coherency**

Adjacency CRFs

$$E(\mathbf{y}|\mathbf{x}) = \sum_i \underbrace{\psi_u(y_i|\mathbf{x})}_{\text{unary potential}} + \sum_{j \in \mathcal{N}(i)} \underbrace{\psi_p(y_i, y_j|\mathbf{x})}_{\text{pairwise potential}}$$

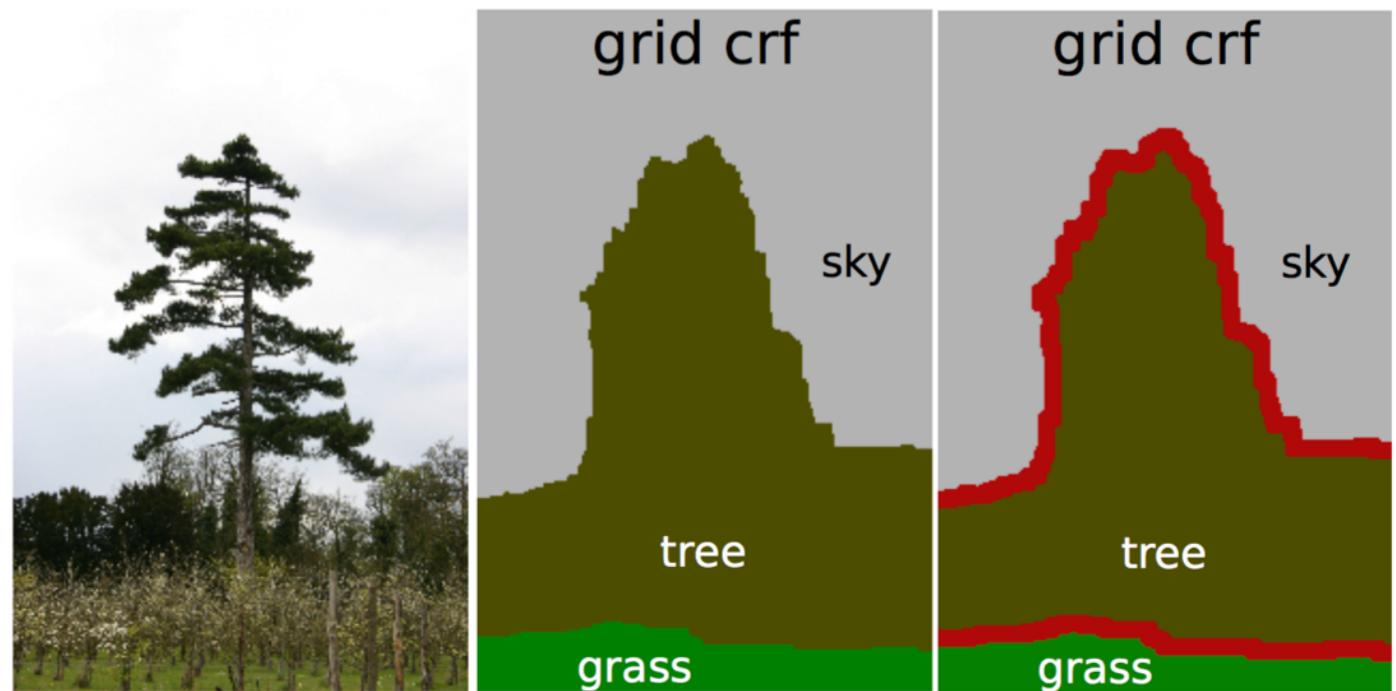
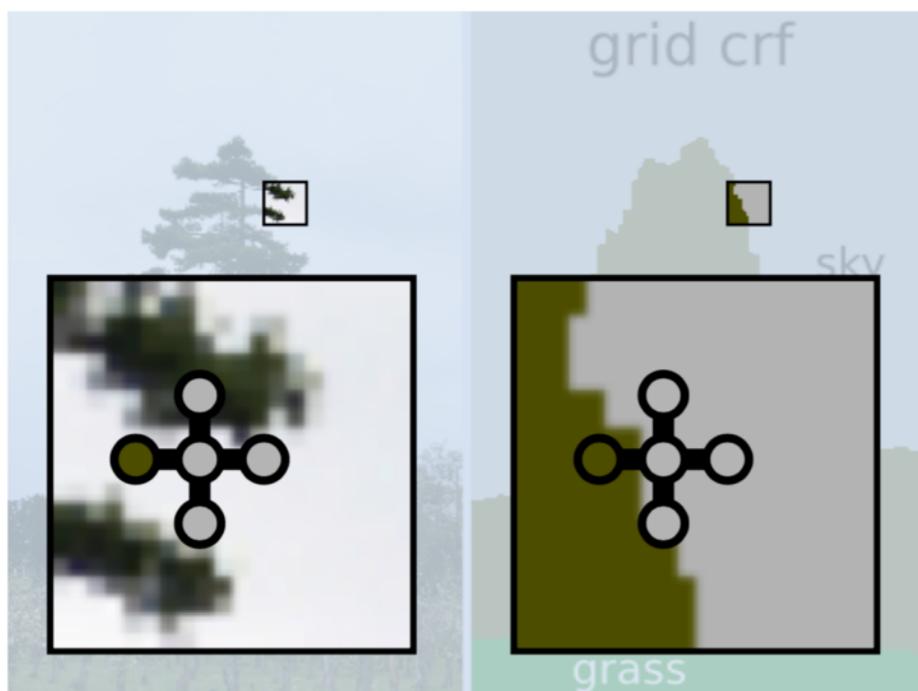


- define a neighbourhood graph
- arises naturally in images
- efficient inference
- **only local interactions**
- **graph is not trainable**

Adjacency CRFs

Popular: Potts model

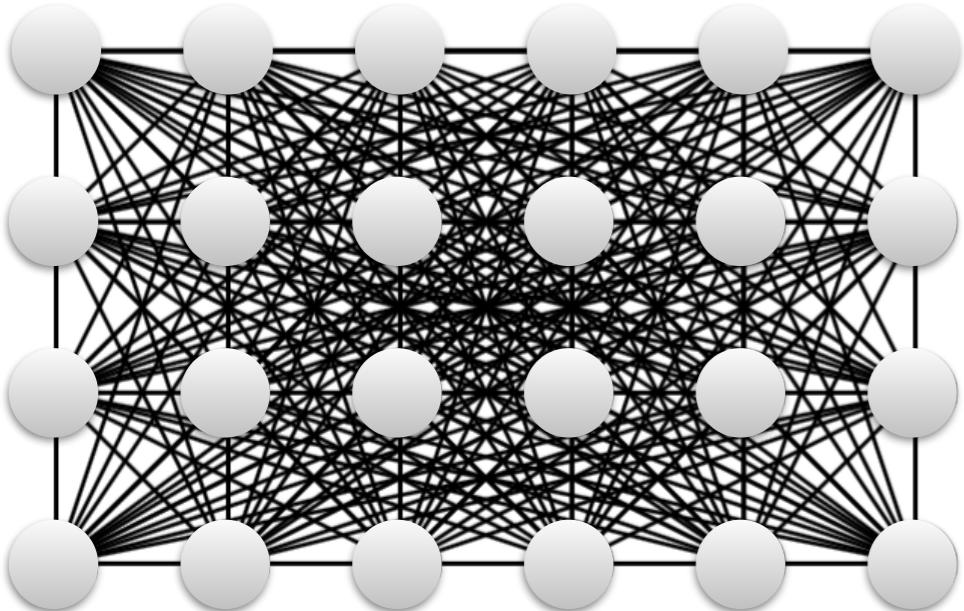
$$\psi_p(y_i, y_j | \mathbf{x}) = 1_{[y_i \neq y_j]} (w_1 \exp(-\beta \|\mathbf{x}_i - \mathbf{x}_j\|^2) + w_2)$$



- **No distinction between classes**
- **Shrinking bias**
- Limited edge-awareness

Fully connected CRFs

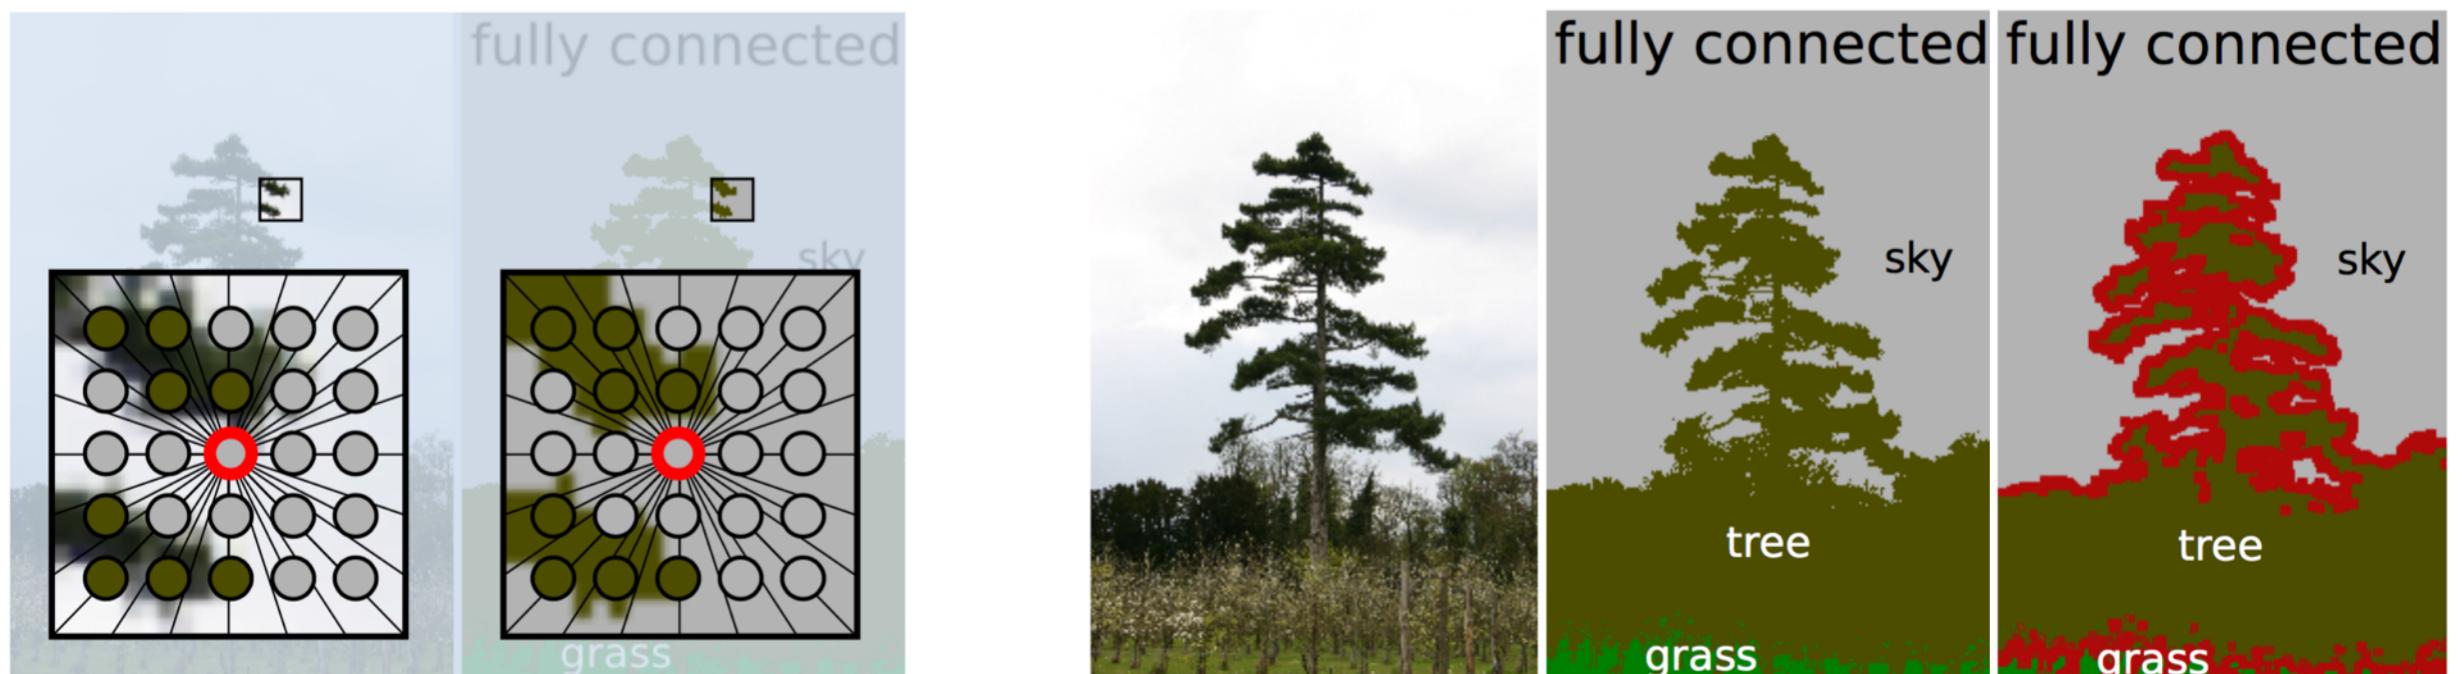
$$E(\mathbf{y}|\mathbf{x}) = \sum_i \underbrace{\psi_u(y_i|\mathbf{x})}_{\text{unary potential}} + \sum_{j \neq i} \underbrace{\psi_p(y_i, y_j|\mathbf{x})}_{\text{pairwise potential}}$$



- Edges between all pairs of nodes
- Long-range interactions

Fully connected CRFs

$$E(\mathbf{y}|\mathbf{x}) = \sum_i \underbrace{\psi_u(y_i|\mathbf{x})}_{\text{unary potential}} + \sum_{j \neq i} \underbrace{\psi_p(y_i, y_j|\mathbf{x})}_{\text{pairwise potential}}$$



- Strength of edges define connectivity
- No more shrinking bias
- **N² edges -> computationally expensive**

Gaussian Edge Potentials

Potts model

$$\psi_p(y_i, y_j | \mathbf{x}) = 1_{[y_i \neq y_j]} (w_1 \exp(-\beta \|\mathbf{x}_i - \mathbf{x}_j\|^2) + w_2)$$

Generalisation

$$\psi_p(y_i, y_j | \mathbf{x}) = \boxed{\mu(y_i, y_j)} \sum_{m=1}^M \boxed{w_m} k_m(\mathbf{f}_i, \mathbf{f}_j)$$

With Gaussian kernels

trainable parameters

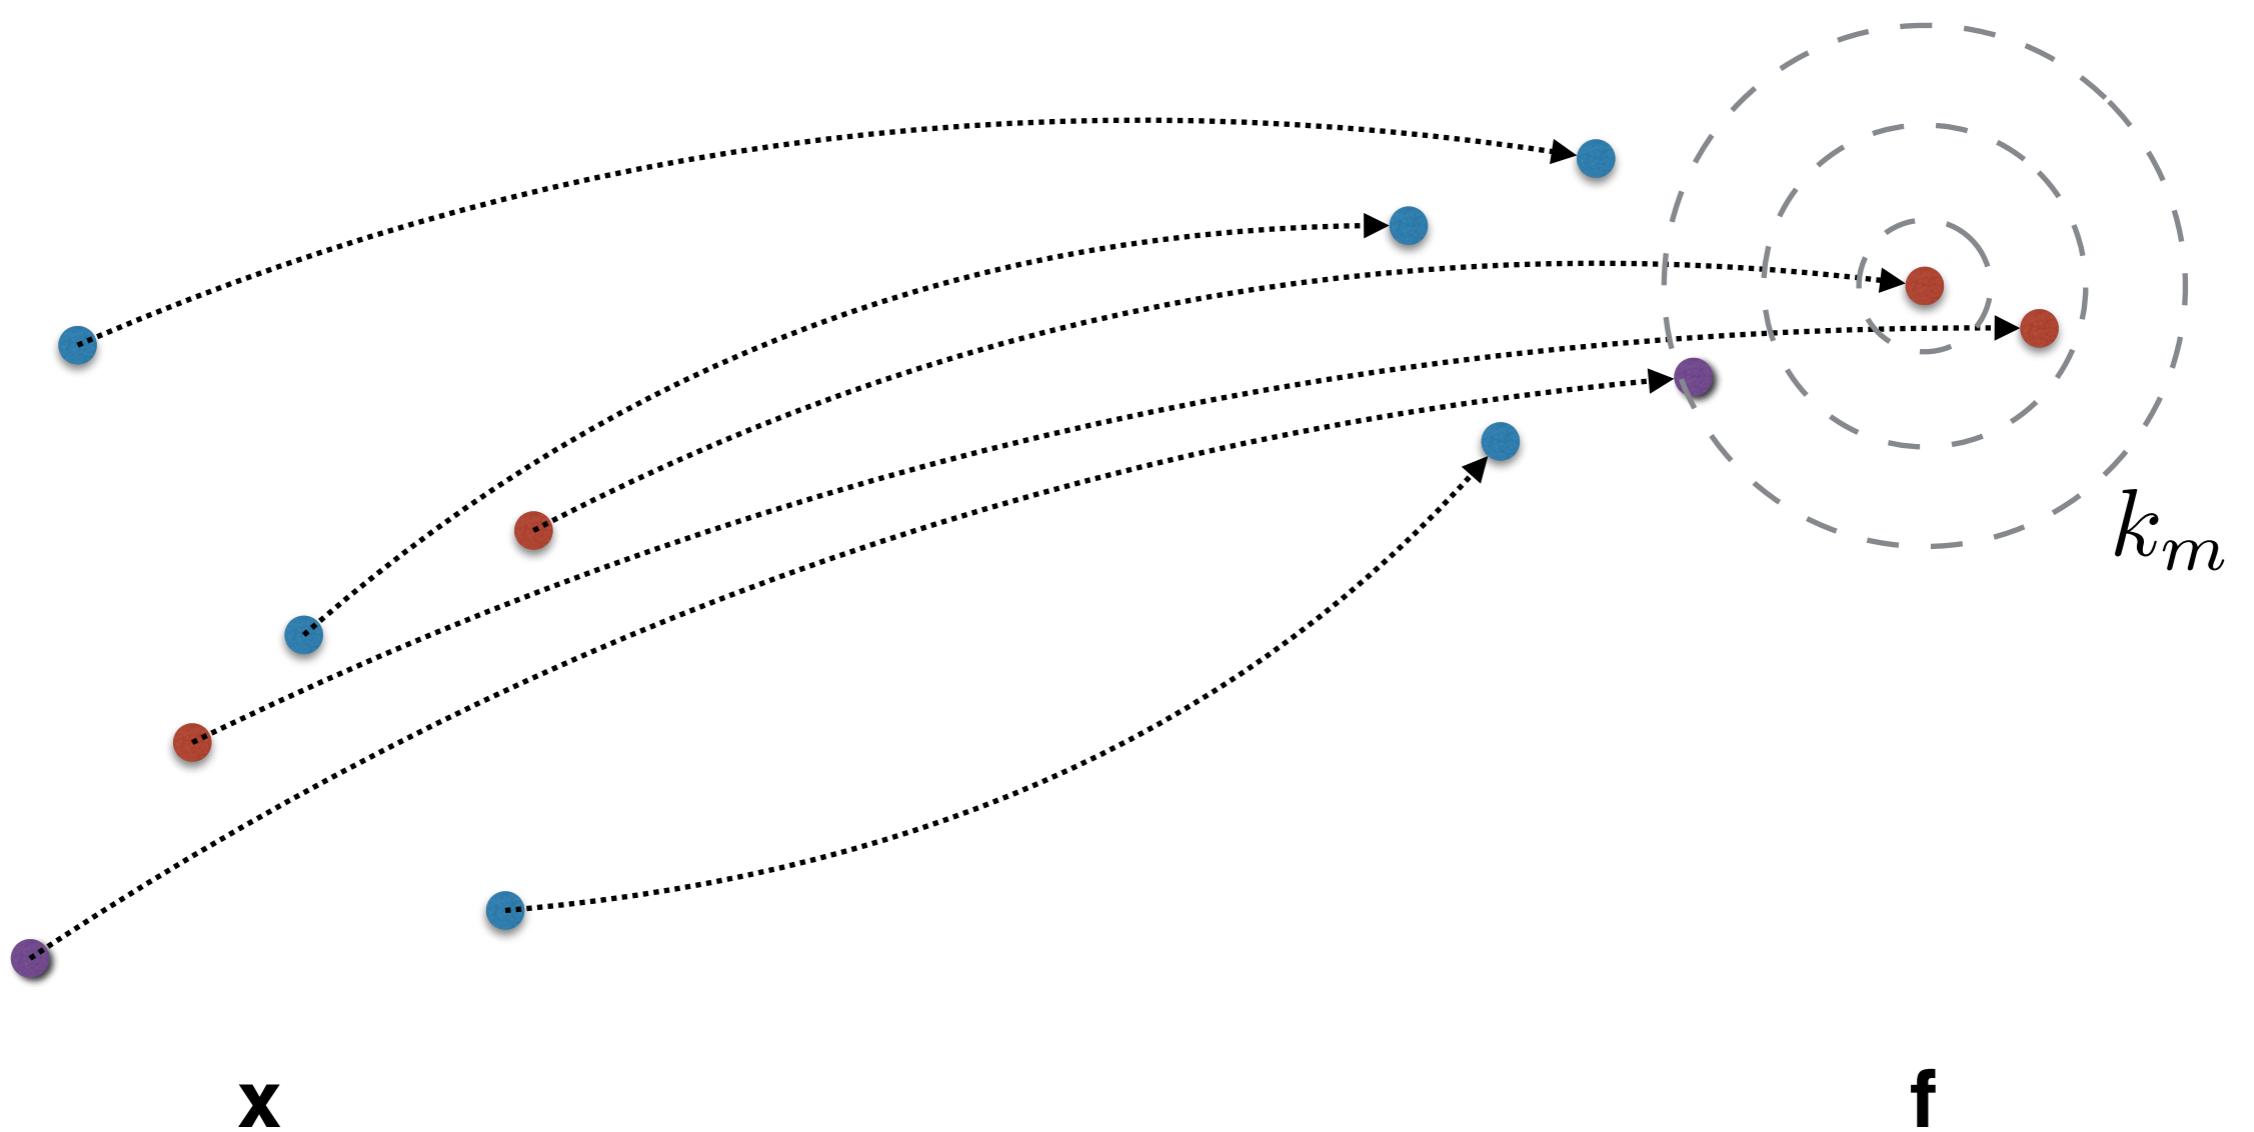
$$k_m(\mathbf{f}_i, \mathbf{f}_j) = \exp \left(-\frac{1}{2} (\mathbf{f}_i - \mathbf{f}_j)^T \boxed{\mathbf{Q}_m} (\mathbf{f}_i - \mathbf{f}_j) \right)$$

Label compatibility $\mu(y_i, y_j)$ models interactions between classes

$$\text{Potts } \mu(y_i, y_j) = 1_{[y_i \neq y_j]}$$

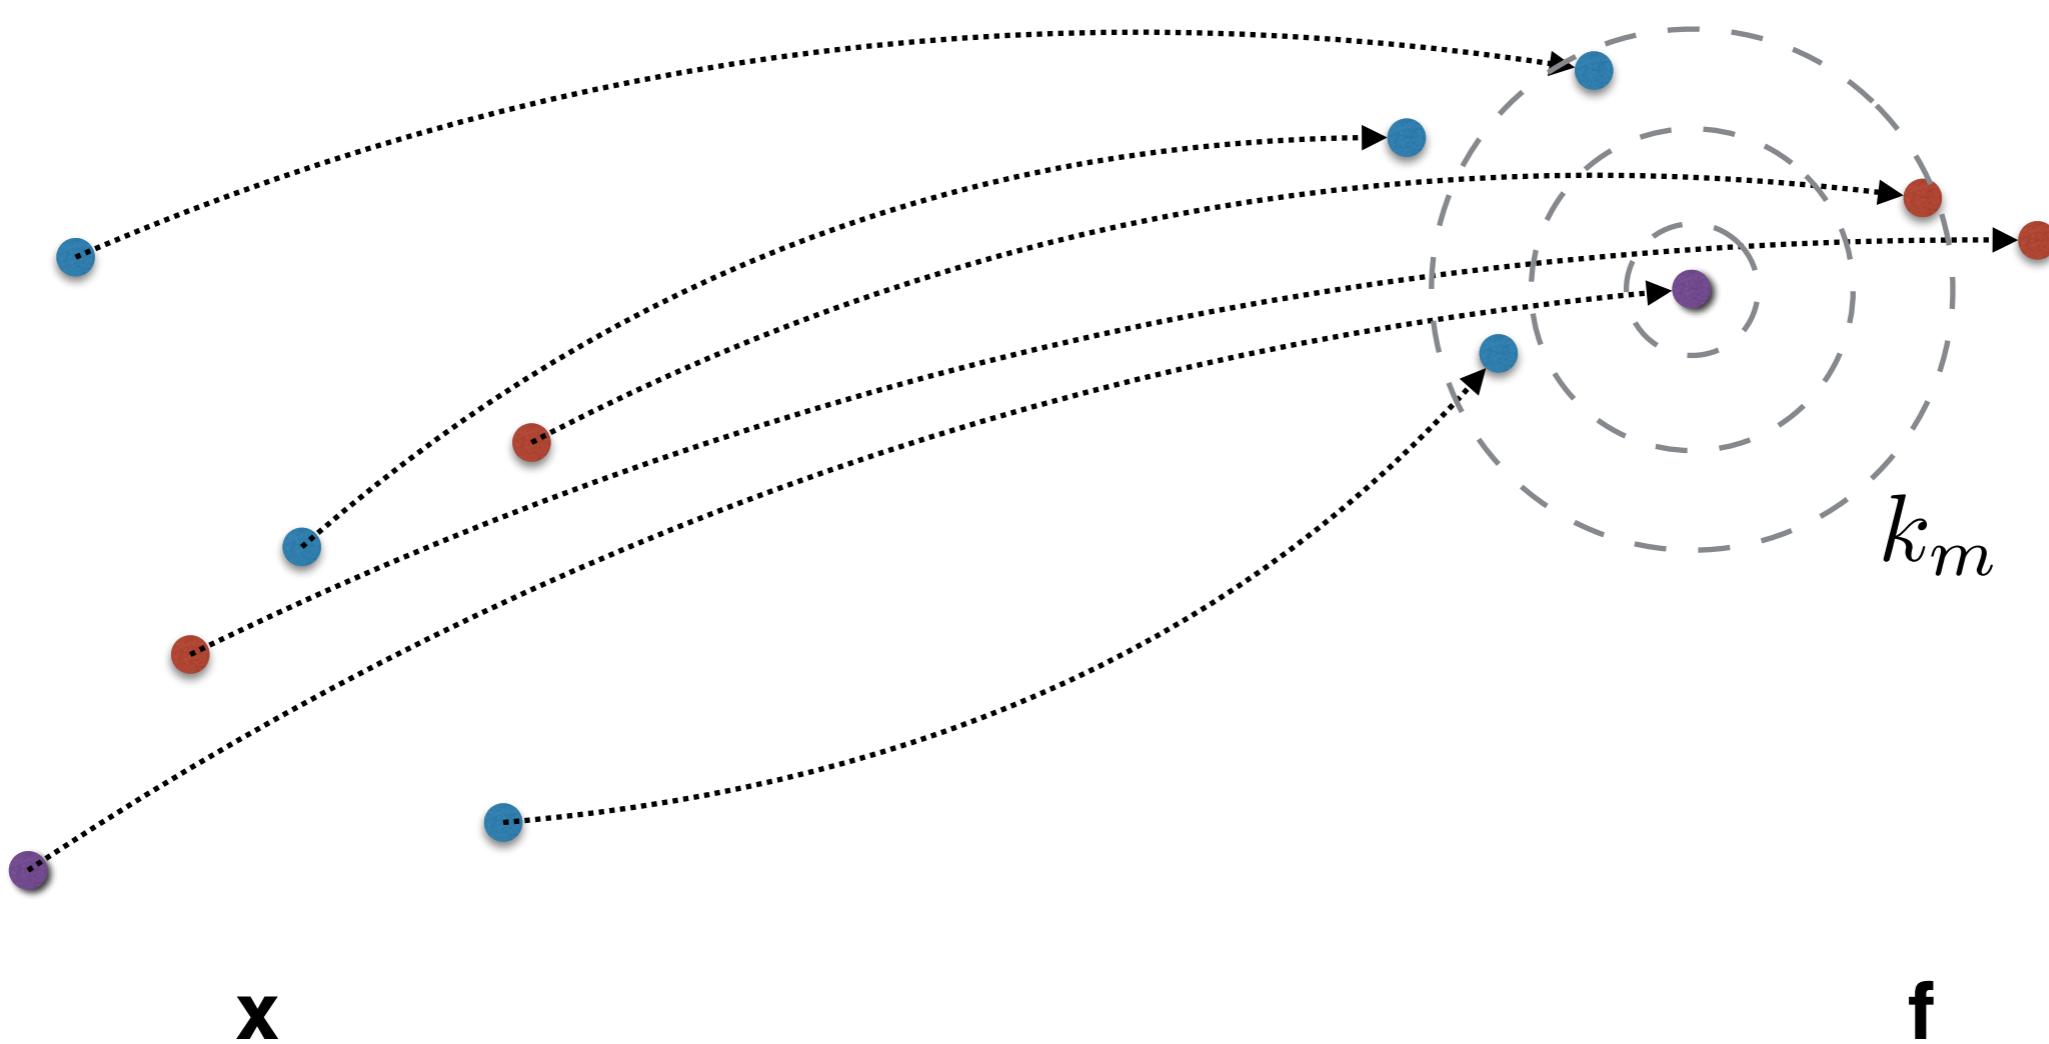
Gaussian Edge Potentials

$$k_m (\mathbf{f}_i, \mathbf{f}_j) = \exp \left(-\frac{1}{2} (\mathbf{f}_i - \mathbf{f}_j)^T \mathbf{Q}_m (\mathbf{f}_i - \mathbf{f}_j) \right)$$



Gaussian Edge Potentials

$$k_m(\mathbf{f}_i, \mathbf{f}_j) = \exp\left(-\frac{1}{2} (\mathbf{f}_i - \mathbf{f}_j)^T \mathbf{Q}_m (\mathbf{f}_i - \mathbf{f}_j)\right)$$



Edge potentials as convolutions

$$E_p(\mathbf{y}|\mathbf{x}) = \sum_i \sum_{j \neq i} \psi_p(y_i, y_j | \mathbf{x})$$

Edge potentials as convolutions

$$\begin{aligned} E_p(\mathbf{y}|\mathbf{x}) &= \sum_i \sum_{j \neq i} \psi_p(y_i, y_j | \mathbf{x}) \\ &= \sum_i \sum_{j \neq i} \mu(y_i, y_j) \sum_{m=1}^M w_m k_m(\mathbf{f}_i, \mathbf{f}_j) \end{aligned}$$

Edge potentials as convolutions

$$\begin{aligned} E_p(\mathbf{y}|\mathbf{x}) &= \sum_i \sum_{j \neq i} \psi_p(y_i, y_j | \mathbf{x}) \\ &= \sum_i \sum_{j \neq i} \mu(y_i, y_j) \sum_{m=1}^M w_m k_m(\mathbf{f}_i, \mathbf{f}_j) \\ &= \sum_{m=1}^M w_m \sum_i \sum_{j \neq i} \mu(y_i, y_j) k_m(\mathbf{f}_i, \mathbf{f}_j) \end{aligned}$$

Edge potentials as convolutions

$$\begin{aligned} E_p(\mathbf{y}|\mathbf{x}) &= \sum_i \sum_{j \neq i} \psi_p(y_i, y_j | \mathbf{x}) \\ &= \sum_i \sum_{j \neq i} \mu(y_i, y_j) \sum_{m=1}^M w_m k_m(\mathbf{f}_i, \mathbf{f}_j) \\ &= \sum_{m=1}^M w_m \sum_i \sum_{j \neq i} \mu(y_i, y_j) k_m(\mathbf{f}_i, \mathbf{f}_j) \\ &= \sum_{m=1}^M w_m \sum_i \sum_{j=1}^N \mu(y_i, y_j) k_m(\mathbf{f}_i, \mathbf{f}_j) - \mu(y_i, y_i) k_m(\mathbf{f}_i, \mathbf{f}_i) \end{aligned}$$

Edge potentials as convolutions

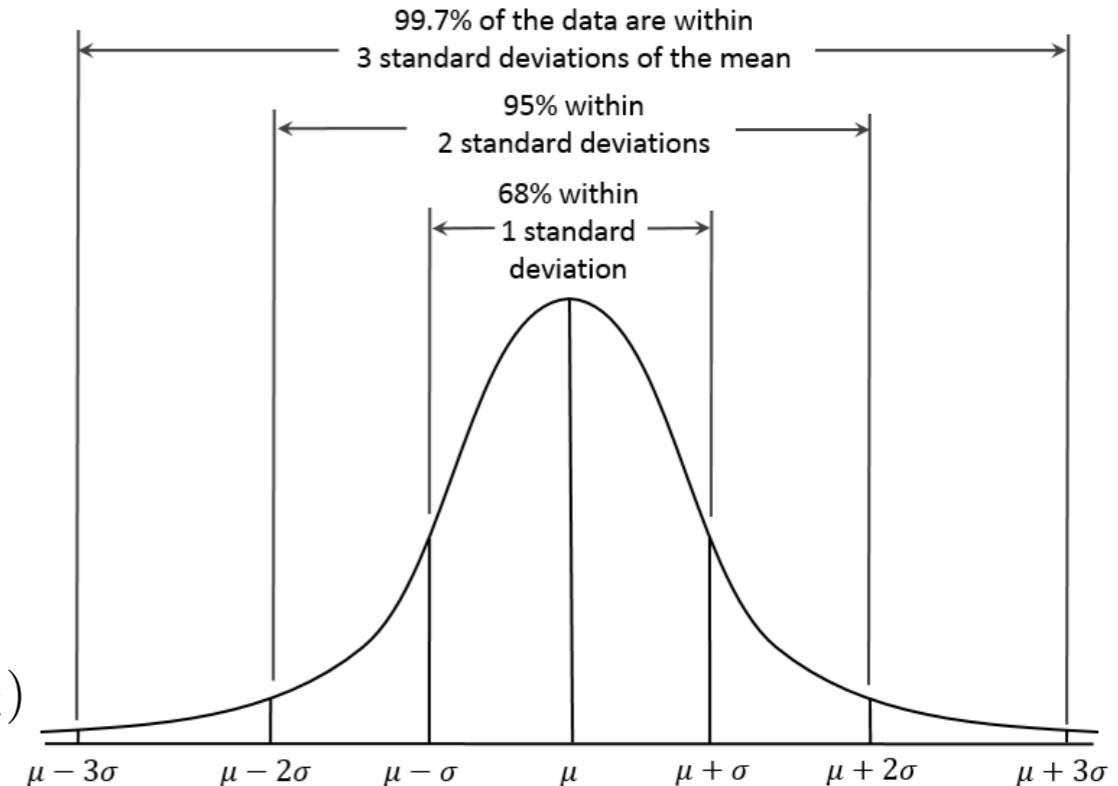
$$E_p(\mathbf{y}|\mathbf{x}) = \sum_i \sum_{j \neq i} \psi_p(y_i, y_j | \mathbf{x})$$

$$= \sum_i \sum_{j \neq i} \mu(y_i, y_j) \sum_{m=1}^M w_m k_m(\mathbf{f}_i, \mathbf{f}_j)$$

$$= \sum_{m=1}^M w_m \sum_i \sum_{j \neq i} \mu(y_i, y_j) k_m(\mathbf{f}_i, \mathbf{f}_j)$$

$$= \sum_{m=1}^M w_m \sum_i \sum_{j=1}^N \mu(y_i, y_j) k_m(\mathbf{f}_i, \mathbf{f}_j) - \mu(y_i, y_i) k_m(\mathbf{f}_i, \mathbf{f}_i)$$

$$= \sum_{m=1}^M w_m \sum_i \sum_{l=i-1}^{i-N} \mu(y_i, y_{i-l}) k_m(\mathbf{f}_i, \mathbf{f}_{i-l}) - \mu(y_i, y_i) k_m(\mathbf{f}_i, \mathbf{f}_i)$$

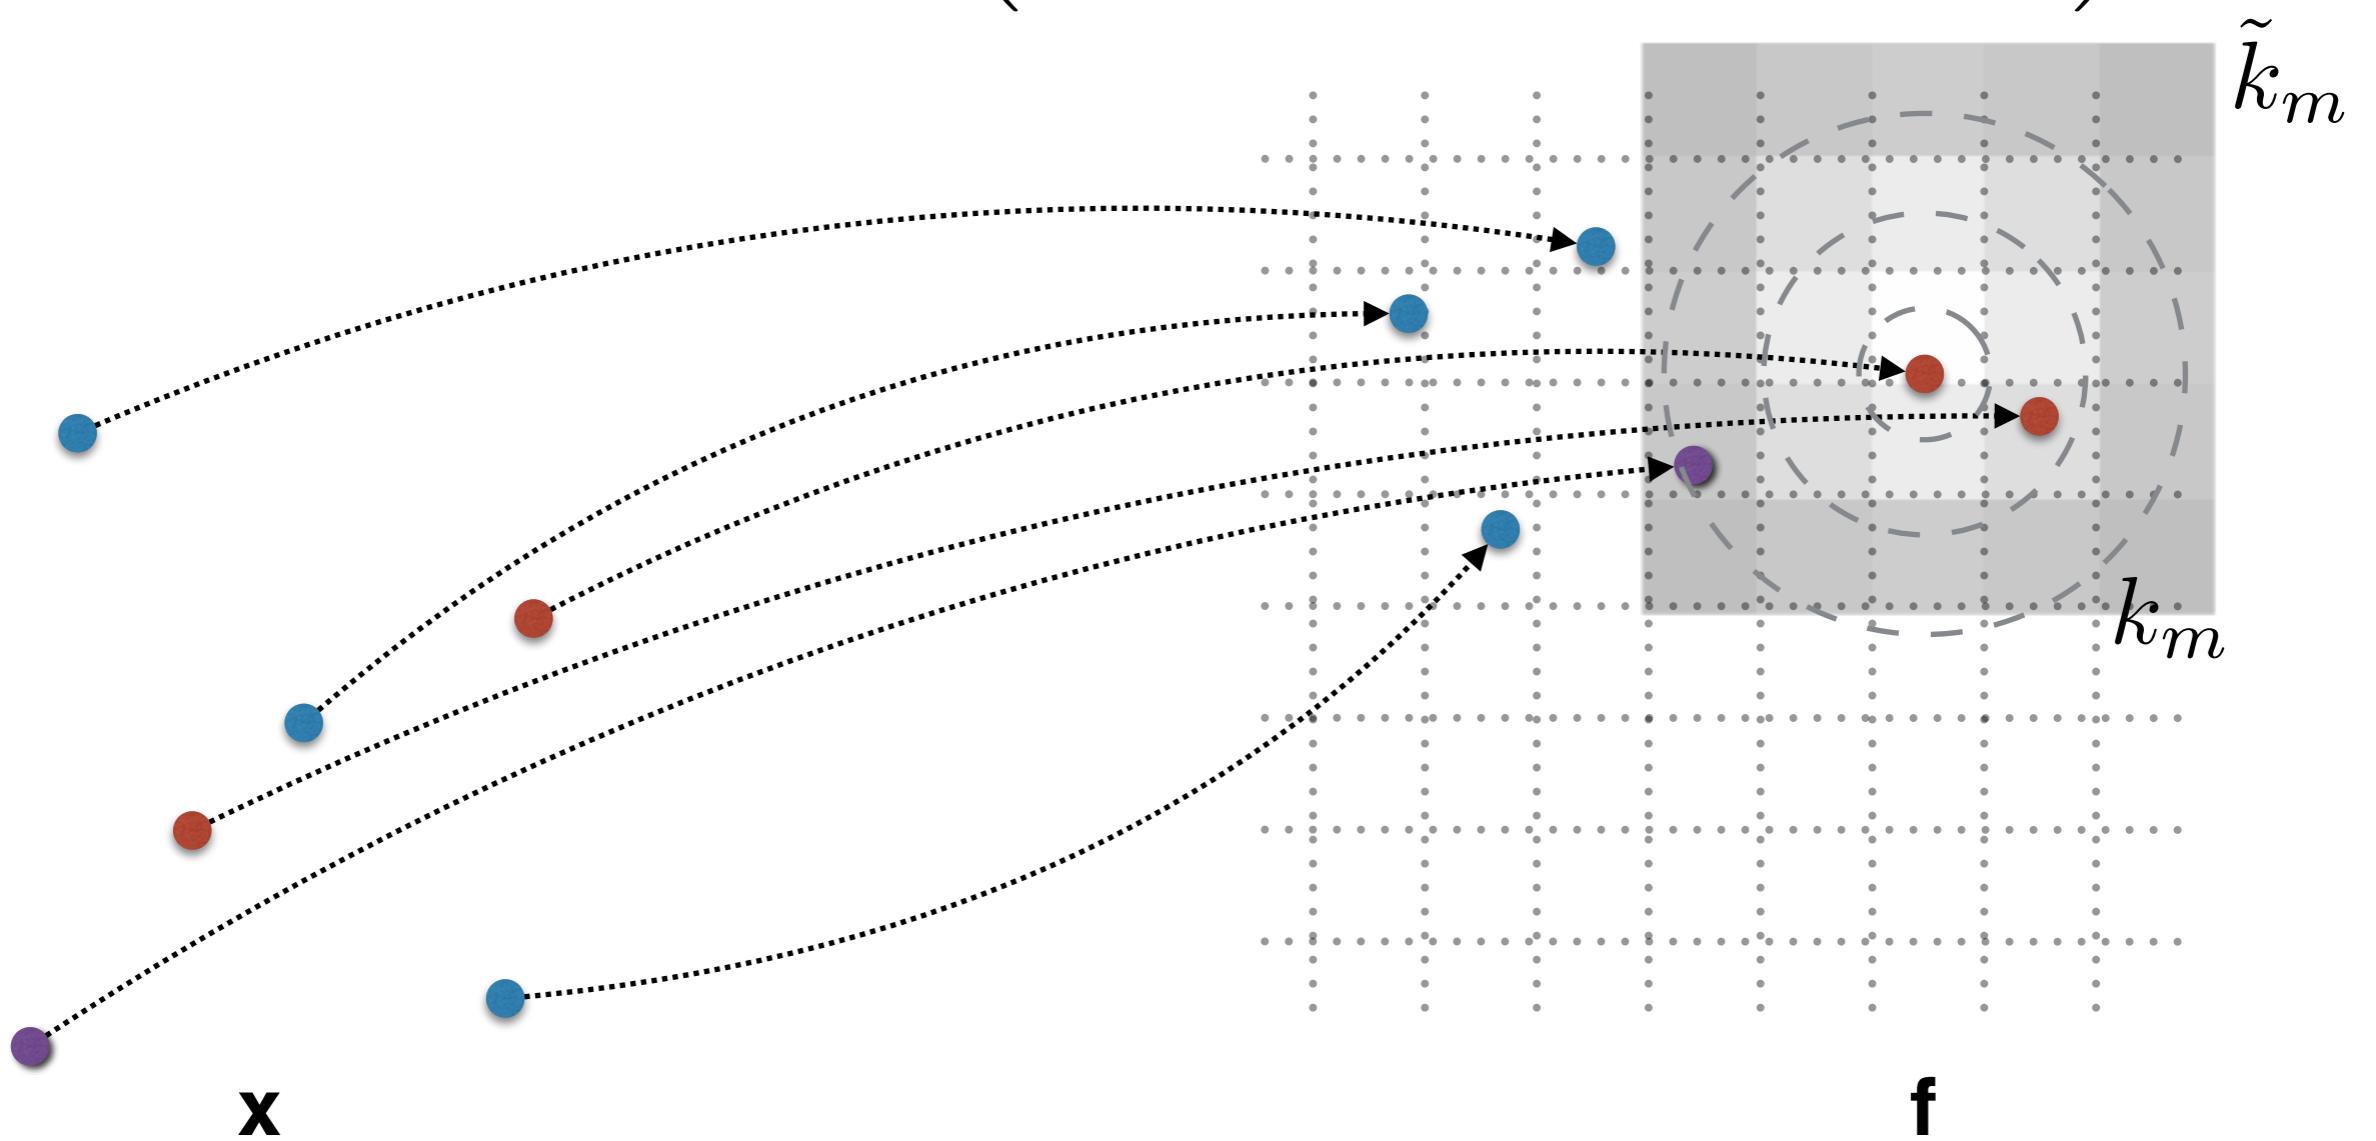


[Wikipedia]

- If k_m is a symmetric kernel the pairwise potential can be implemented as a convolution
- Gaussian has most probability mass around a close region about its mean
- Approximate this using a truncated Gaussian

Gaussian Edge Potentials

$$k_m (\mathbf{f}_i, \mathbf{f}_j) = \exp \left(-\frac{1}{2} (\mathbf{f}_i - \mathbf{f}_j)^T \mathbf{Q}_m (\mathbf{f}_i - \mathbf{f}_j) \right)$$

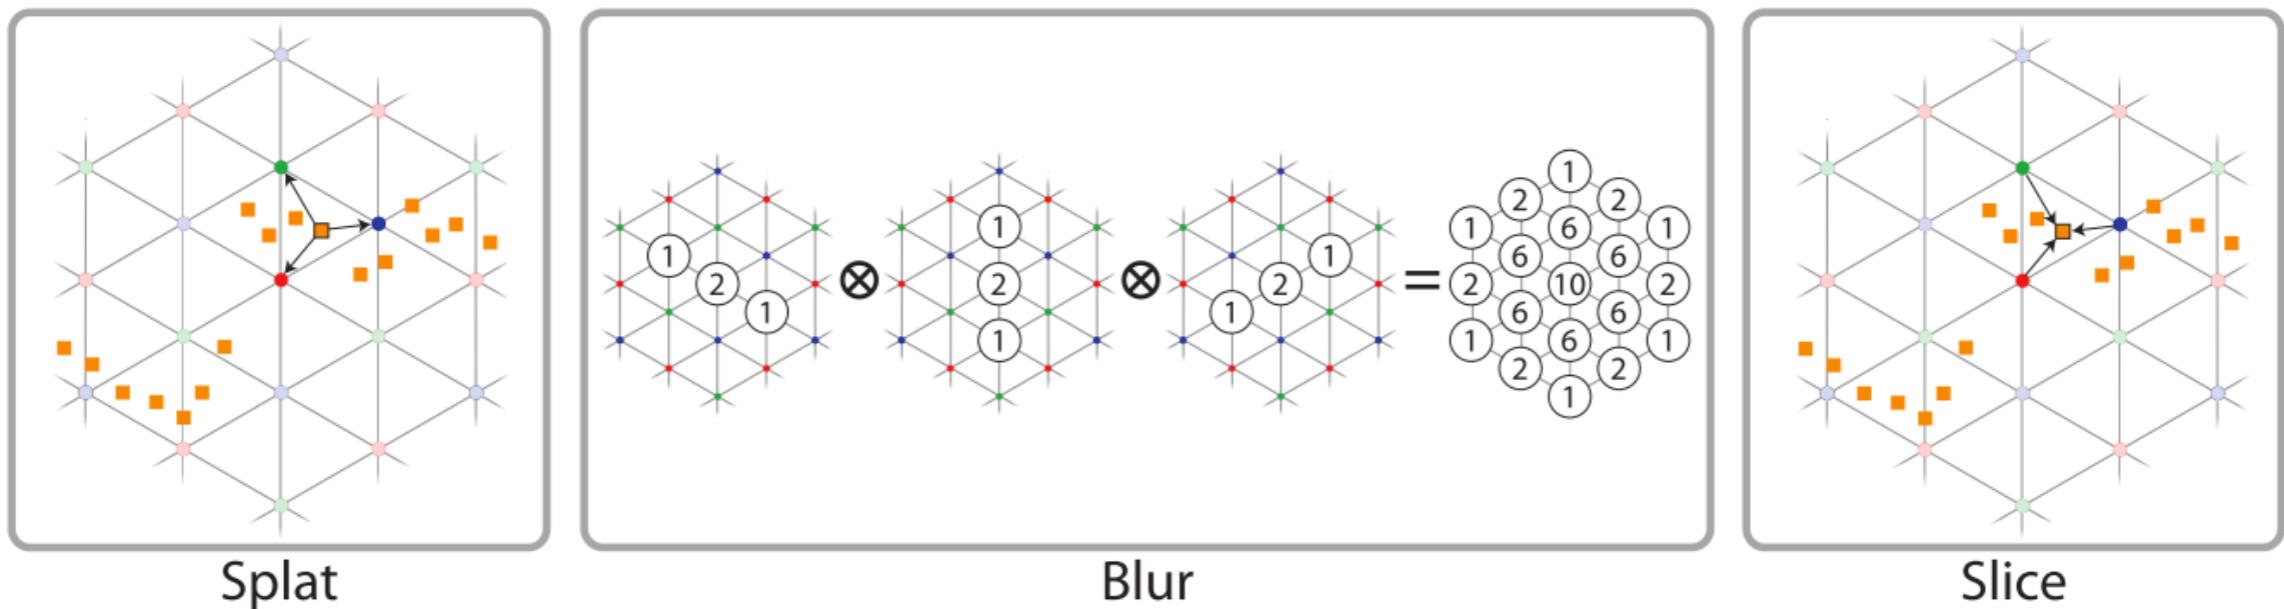


Convolution in a high dimensional space

Permutohedral Lattices for convolution in high dimensional space

Problem:

- f_i can be high dimensional
- as a result the feature space will be sparsely populated



[Adams et al., 2010]

- Convolutions in $\mathcal{O}(d^2n)$ and $\mathcal{O}(dn)$ for separable filters
- Naive convolution on a grid is $\mathcal{O}(2^d n)$

Efficient inference in dense CRFs

$$p(\mathbf{y}|\mathbf{x}) = \frac{1}{Z(\mathbf{x})} \exp \left(-\sum_i \psi_u(y_i|\mathbf{x}) - \sum_{j \neq i} \psi_p(y_i, y_j|\mathbf{x}) \right)$$

Hard to evaluate

Mean-field approximation: $Q(\mathbf{y}|\mathbf{x}) = \prod_i Q(y_i|\mathbf{x})$

Mean field updates:

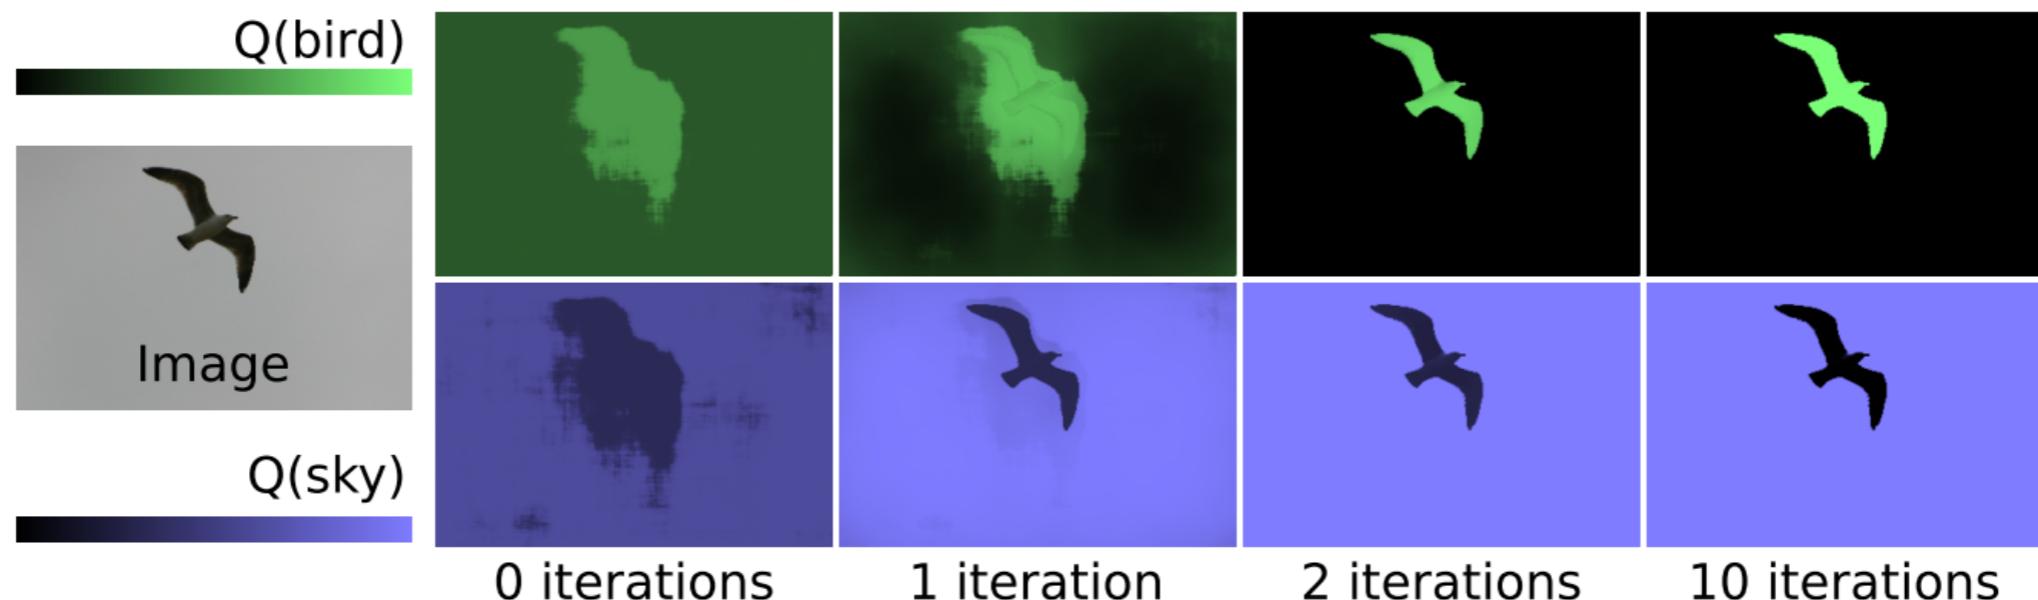
$$Q_i(y_i = l) = \frac{1}{Z_i} \exp \left(-\psi_u(y_i|\mathbf{x}) - \sum_{l' \in \mathcal{L}} \mu(l, l') \sum_{m=1}^M w_m \sum_{j \neq i} k_m(\mathbf{f}_i, \mathbf{f}_j) Q_j(l') \right)$$

- Naively implemented an update over all Q_i costs $O(N^2)$
- With the approximations from before it is $O(N)$

Efficient inference in dense CRFs

Mean field updates:

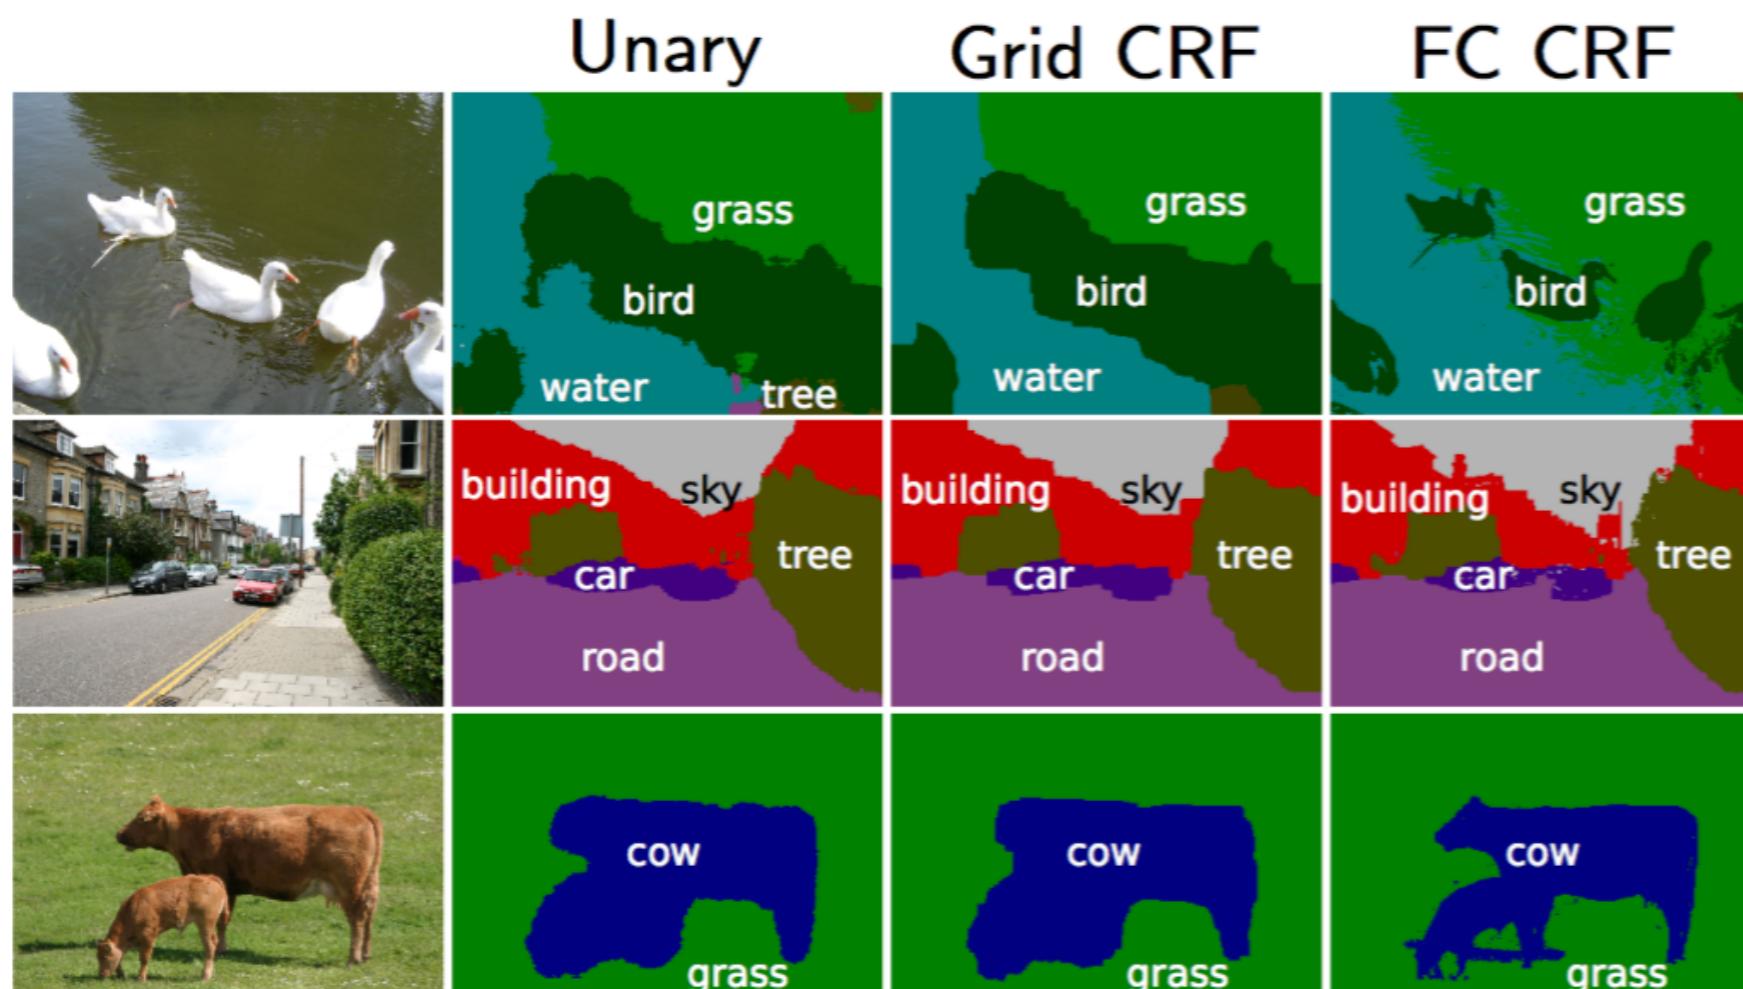
$$Q_i(y_i = l) = \frac{1}{Z_i} \exp \left(-\psi_u(y_i | \mathbf{x}) - \sum_{l' \in \mathcal{L}} \mu(l, l') \sum_{m=1}^M w_m \sum_{j \neq i} k_m(\mathbf{f}_i, \mathbf{f}_j) Q_j(l') \right)$$



(b) Distributions $Q(X_i = \text{"bird"})$ (top) and $Q(X_i = \text{"sky"})$ (bottom)

10 Iterations own 0.2 seconds

Efficient inference in dense CRFs



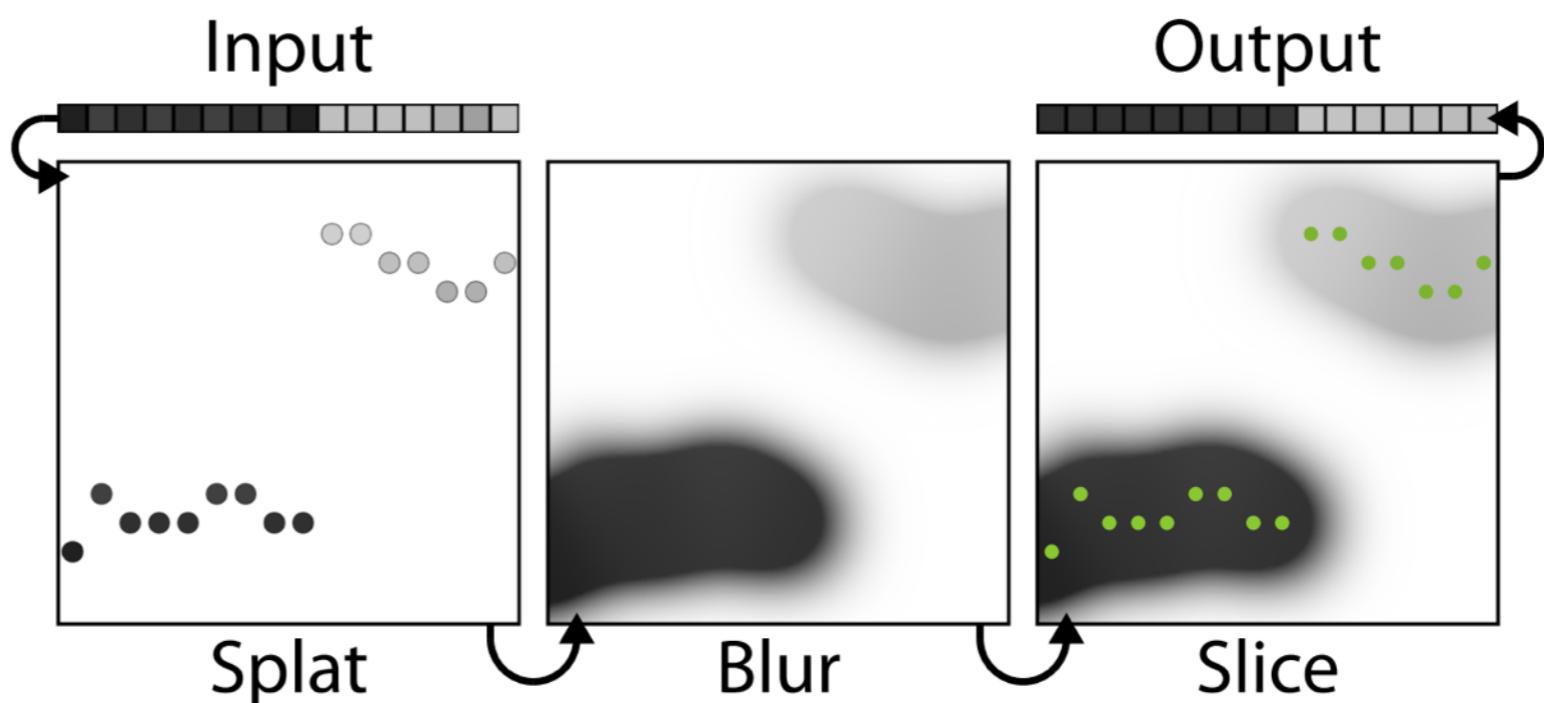
	Time	Global	Avg
Unary	-	84.0	76.6
Grid CRF	1s	84.6	77.2
FC CRF	0.2s	86.0	78.3

[Krähenbühl & Koltun, 2011]

What are these features?

$$k_m (\mathbf{f}_i, \mathbf{f}_j) = \exp \left(-\frac{1}{2} (\mathbf{f}_i - \mathbf{f}_j)^T \mathbf{Q}_m (\mathbf{f}_i - \mathbf{f}_j) \right)$$

$$\mathbf{f}_i = \begin{pmatrix} \mathbf{p}_i \\ \mathbf{x}_i \end{pmatrix}$$



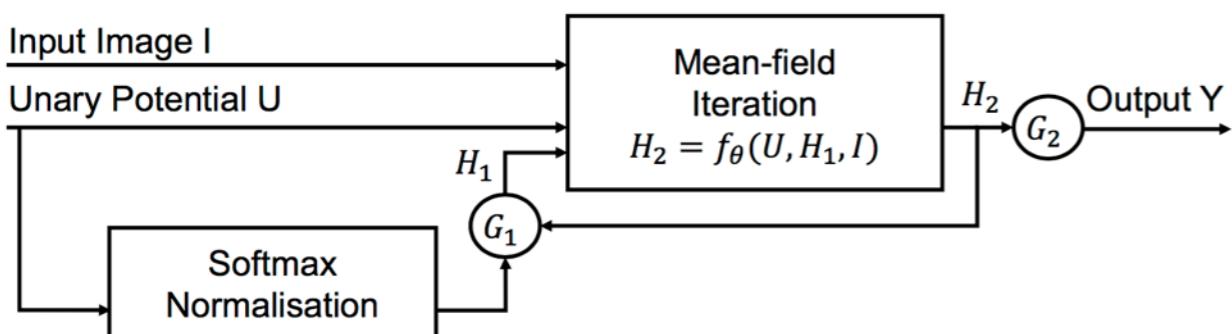
\mathbf{p}_i position vector

[Adams et al., 2010]

\mathbf{x}_i vector of pixel values

- This is known as bilateral filtering
- It is an edge-preserving approach to filtering

CRFs as RNNs

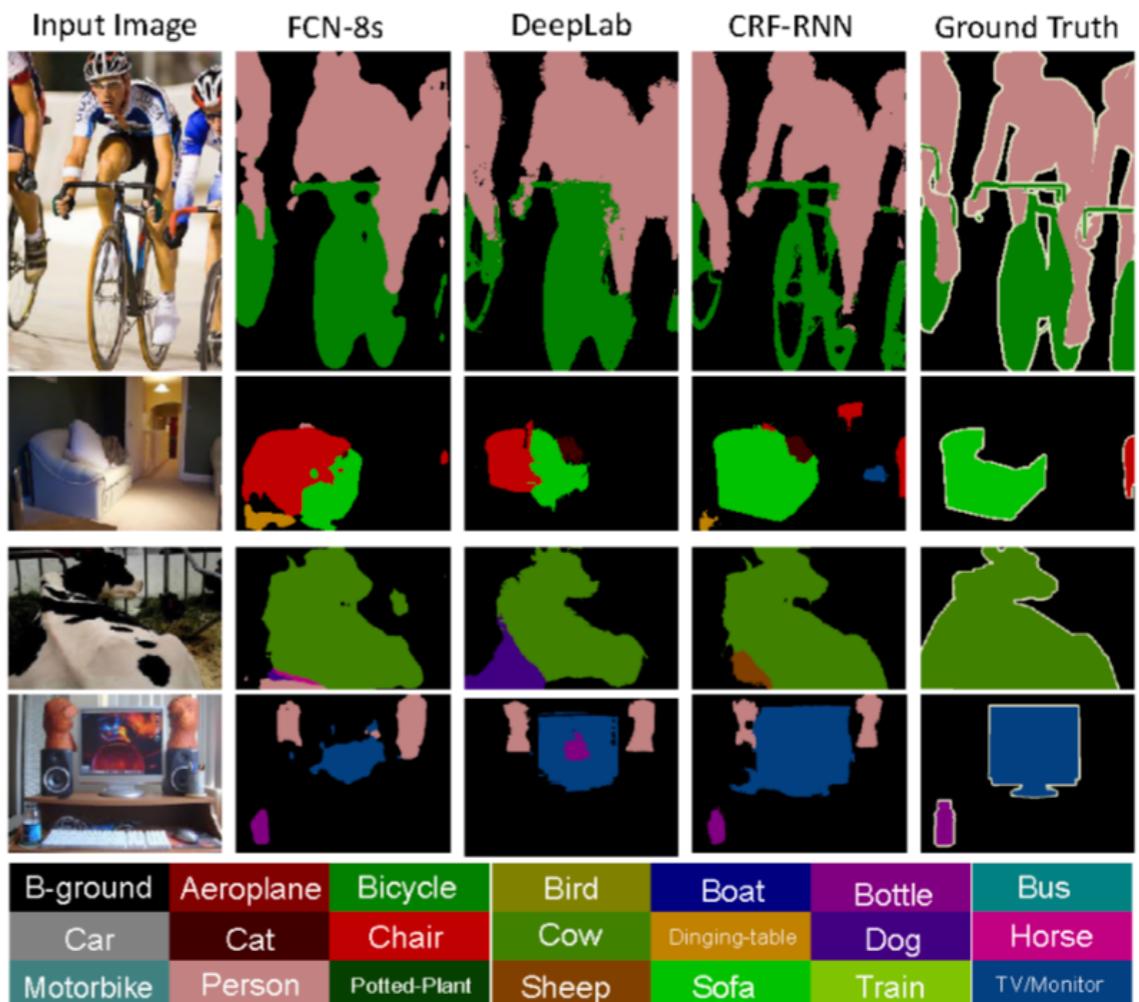


PASCAL VOC

Our approach achieved state-of-the-art comparable performance in PASCAL VOC semantic image segmentation dataset, with mean intersection-over-union (IOU) score 74.7% on VOC 2012 test set.

[Zheng et al., 2015]

- Mean field iterations as steps in an RNN
- End-to-end training



Demo: <http://www.robots.ox.ac.uk/~szheng/crfasrnndemo>

Learning the filters

+ MF-1step	+ MF-2 step	+ <i>loose</i> MF-2 step
Semantic segmentation (IoU) - CNN [16]: 72.08 / 66.95		
Gauss CRF +2.48	+3.38	+3.38 / +3.00
Learned CRF +2.93	+3.71	+3.85 / +3.37
Material segmentation (Pixel Accuracy) - CNN [11]: 67.21 / 69.23		
Gauss CRF +7.91 / +6.28	+9.68 / +7.35	+9.68 / +7.35
Learned CRF +9.48 / +6.23	+11.89 / +6.93	+11.91 / +6.93

Table 2. Improved mean-field inference with learned potentials. (top) Average IoU score on Pascal VOC12 validation/test data [20] for semantic segmentation; (bottom) Accuracy for all pixels / averaged over classes on the MINC test data [11] for material segmentation.

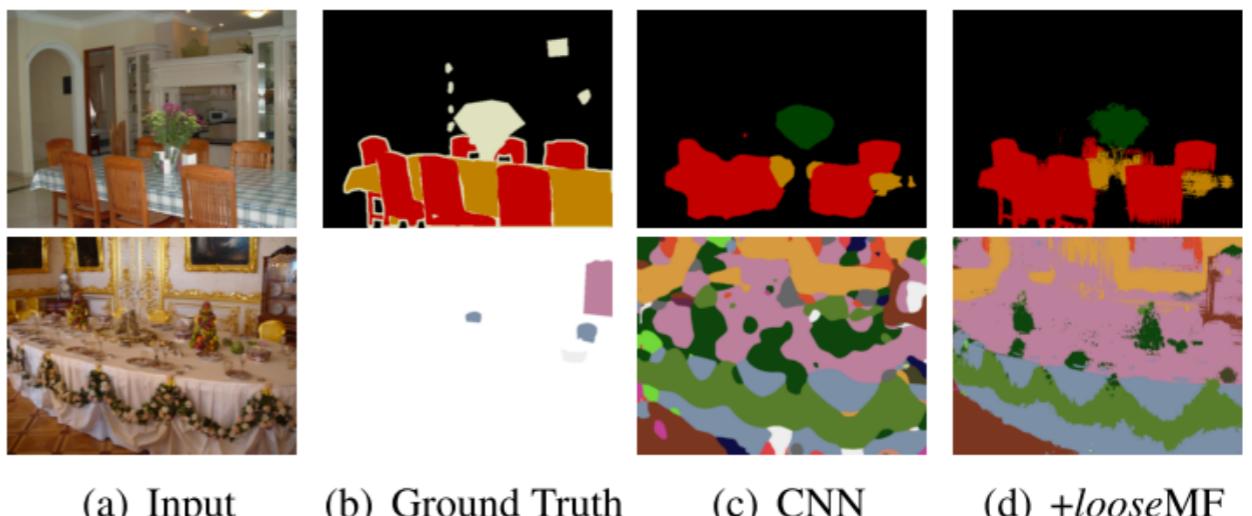


Figure 4. Segmentation results. An example result for semantic (top) and material (bottom) segmentation. (c) depicts the unary results before application of MF, (d) after two steps of *loose*-MF with a learned CRF. More examples with comparisons to Gaussian pairwise potentials are in Sec. C.5 and Sec. C.6.

Summary

- We can improve performance of segmentation algorithms by adding pairwise interactions between pixel predictions
- Approximate Inference in dense CRFs can be done efficiently
- End-to-end learning can improve performance
- Pairwise kernels can be learned and lead to potential improvements
- These methods could have an impact in other problem domains

References

- (1) Adams A, Baek J, Davis MA. Fast high-dimensional filtering using the permutohedral lattice. *Comput Graph Forum*. 2010;29(2):753-762.
- (2) Krähenbühl P, Koltun V. Efficient Inference in Fully Connected CRFs with Gaussian Edge Potentials. In: Advances in Neural Information Processing Systems.; 2011:109-117. <http://papers.nips.cc/paper/4296-efficient-inference-in-fully-connected-crfs-with-gaussian-edge-potentials>. Accessed January 8, 2016.
- (3) Kiefel M, Gehler P V, Kiefel M, Mpg T, Gehler P, Mpg T. Sparse Convolutional Networks using the Permutohedral Lattice. 2015. <http://arxiv.org/abs/1503.04949>. Accessed January 8, 2016.
- (4) Shotton J, Winn J, Rother C, Criminisi A. TextonBoost for Image Understanding: Multi-Class Object Recognition and Segmentation by Jointly Modeling Texture, Layout, and Context. *Int J Comput Vis*. 2009. <http://research.microsoft.com/apps/pubs/default.aspx?id=117885>.
- (5) Zheng S, Jayasumana S, Romera-Paredes B, et al. Conditional Random Fields as Recurrent Neural Networks. In: International Conference on Computer Vision (ICCV).; 2015:16. <http://arxiv.org/abs/1502.03240>. Accessed January 8, 2016.