

# Group Equivariant Deep Learning

Lecture 1 - Regular group convolutions

Lecture 1.5 - A brief history of G-CNNs

# G-CNNs rule!

- The right inductive bias: guaranteed equivariance  
(no loss of information)
- Performance gains that can't be obtained by data-augmentation alone  
(both local and global equivariance/invariance)
- Increased sample efficiency  
(increased weight sharing, no geometric augmentation necessary)

# G-CNNs rule!

- The right inductive bias: guaranteed equivariance (no loss of information)
- Performance gains that can't be obtained by data-augmentation alone (both local and global equivariance/invariance)
- Increased sample efficiency (increased weight sharing, no geometric augmentation necessary)

**Oriented Response Networks**

Yanzhou Zhou<sup>1</sup>, Qixing Yu<sup>1</sup>, Qianqian Guo<sup>2</sup>, and Jitendra Malik<sup>1</sup>

<sup>1</sup>University of Colorado A  
<sup>2</sup>Duke University, Department of Mathematics, NC 27708

**Abstract**

Published as a conference paper at ICLR 2020.

**B-SPLINE CNNS ON LIE GROUPS**

Dirk J. Deijen<sup>1\*</sup>  
Amsterdam Institute for Advanced Metropolitan Solutions (AIAM), University of Amsterdam, The Netherlands

**Roto-Translation-Equivariant Convolutional Neural Network Application to Histopathology**

Marion W. Uffinger<sup>1</sup>, Dirk J. Deijen<sup>1</sup>, Joost P.W. Pijnacker<sup>1</sup>

<sup>1</sup>Department of Pathology, Erasmus MC, University Medical Center, Rotterdam, The Netherlands

**Group-Equivariant Convolutional Network**

Eric S. Cahn<sup>1</sup>  
University of Amsterdam

Max Welling<sup>1</sup>  
University of Amsterdam  
University of Oxford, United Kingdom  
Cavendish Laboratory for Advanced Research

**Abstract**

We introduce Group-Equivariant Convolutional Networks (G-CNNs), a natural generalization of convolutional neural networks that encode data symmetrically by exploiting symmetries. G-CNNs use G-equivariant convolution layers that enjoy a substantially higher degree of weight sharing than standard convolution layers. This allows us to encode the same data with a much smaller number of parameters. Group-equivariant layers can also be combined with other kinds of layers, such as standard fully connected layers. Our proposed G-CNNs achieve state-of-the-art results on CIFAR10 and ImageNet MNIST.

**1. Introduction**

Deep convolutional neural networks (CNNs), commonly known as neural networks, have proven to be very powerful models of sensory data such as images, video, and audio. Although a strong dependence on empirical evidence supports the notion that deep convolutional neural networks (CNNs) are superior to other types of neural networks, there is no theoretical proof that they are better than other types of neural networks. In this paper, we propose a new type of convolutional layer that is equivariant to transformations from a specific Lie group with a suitable exponential map. Incorporating equivariance to a given group requires implementing only the group exponential and logarithm maps, enabling rapid prototyping. Leveraging the simplicity and generality of our model, we apply the same model architecture to both image and point cloud data, demonstrating the capacity to learn equivariant representations of each part of the image, a convolution layer can be trained independently of the other parts, while preserving the capacity to learn meaningful associations between them.

**Generalizing Convolutional Neural Networks for Equivariance to Lie Groups on Arbitrary Continuous Data**

Marc Finzi<sup>1</sup>, Samuel Stanton<sup>1</sup>, Paul Isola<sup>1</sup>, Andrew Gordon Wilson<sup>2</sup>  
New York University

**Abstract**

The inductive biases of group equivariant networks lead to linear gains in the training of a wide variety of applications, such as 3D point clouds. A particular focus has been on rotation and translation equivariant CNNs for point clouds. Here we present a generalization of E(2)-equivariant convolutional networks to the framework of Lie groups. The theory of Steerable CNNs leads us to examine the convolution kernels which depend on non-commutative

**Figure 1: Many instances of spatial data do not fall under important symmetry groups. We present a more comprehensive survey due to the need to support a generalization of equivariant networks.**


**1. Introduction**

Symmetry pervades the natural world. The same law of gravitation governs a game of checkers, the orbits of our planets, and the formation of galaxies. It is precisely because of the role of the universe that we can hope to understand it. Once we started to understand the symmetries inherent in physical laws, we could predict behavior in galaxies billions of light years away by studying our own local region of space and time. One of the first tools to achieve their full potential, it is essential to incorporate our knowledge of symmetry and equivariance into the design of algorithms and architectures. An example of the principle of equivariance in deep learning is convolutional neural networks [LeCun et al., 1998]; when an input (e.g., an image) is translated, the output of a convolutional layer is translated in the same way.

Group theory provides a mathematical language for symmetry and equivariance. Convolutional layers are equivariant under the action of some group. Equivariant networks


# A brief history of G-CNNs


<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>



Cesa-Lang-Weiler 2022  
 $G = \mathbb{R}^d \rtimes H$  with  $H$  compact

<https://quva-lab.github.io/escnn/>





<http://arxiv.org/abs/2001.09637>

## Group Equivariant Convolutional Networks

Taco S. Cohen  
University of Amsterdam

Max Welling  
University of Amsterdam  
University of California Irvine  
Canadian Institute for Advanced Research

T.S.COHEN@UVA.NL

M.WELLING@UVA.NL

### Abstract

We introduce Group equivariant Convolutional Neural Networks (G-CNNs), a natural generalization of convolutional neural networks that reduces sample complexity by exploiting symmetries. G-CNNs use G-convolutions, a new type of layer that enjoys a substantially higher degree of weight sharing than regular convolution layers. G-convolutions increase the expressive capacity of the network without increasing the number of parameters. Group convolution layers are easy to use and can be implemented with negligible computational overhead for discrete groups generated by translations, reflections and rotations. G-CNNs achieve state of the art results on CIFAR10 and rotated MNIST.

### 1. Introduction

Deep convolutional neural networks (CNNs, convnets) have proven to be very powerful models of sensory data such as images, video, and audio. Although a strong theory of neural network design is currently lacking, a large amount of empirical evidence supports the notion that both *convolutional weight sharing* and *depth* (among other factors) are important for good predictive performance.

Convolutional weight sharing is effective because there is a *translation symmetry* in most perception tasks: the label function and data distribution are both approximately invariant to shifts. By using the same weights to analyze or model each part of the image, a convolution layer uses far fewer parameters than a fully connected one, while preserving the capacity to learn many useful transformations.

*Proceedings of the 33<sup>rd</sup> International Conference on Machine Learning*, New York, NY, USA, 2016. JMLR: W&CP volume 48. Copyright 2016 by the author(s).

Convolution layers can be used effectively in a *deep* network because all the layers in such a network are *translation equivariant*: shifting the image and then feeding it through a number of layers is the same as feeding the original image through the same layers and then shifting the resulting feature maps (at least up to edge-effects). In other words, the symmetry (translation) is preserved by each layer, which makes it possible to exploit it not just in the first, but also in higher layers of the network.



In this paper we show how convolutional networks can be generalized to exploit larger groups of symmetries, including rotations and reflections. The notion of equivariance is key to this generalization, so in section 2 we will discuss this concept and its role in deep representation learning. After discussing related work in section 3, we recall a number of mathematical concepts in section 4 that allow us to define and analyze the G-convolution in a generic manner.

In section 5, we analyze the equivariance properties of standard CNNs, and show that they are equivariant to translations but may fail to equivariant with more general transformations. Using the mathematical framework from section 4, we can define G-CNNs (section 6) by analogy to standard CNNs (the latter being the G-CNN for the translation group). We show that G-convolutions, as well as various kinds of layers used in modern CNNs, such as pooling, arbitrary pointwise nonlinearities, batch normalization and residual blocks are all equivariant, and thus compatible with G-CNNs. In section 7 we provide concrete implementation details for group convolutions.

In section 8 we report experimental results on MNIST-rot and CIFAR10, where G-CNNs achieve state of the art results (2.28% error on MNIST-rot, and 4.19% resp. 6.46% on augmented and plain CIFAR10). We show that replacing planar convolutions with G-convolutions consistently improves results without additional tuning. In section 9 we provide a discussion of these results and consider several extensions of the method, before concluding in section 10.



Tensor field networks:  
station- and translation-equivariant neural  
networks for 3D point clouds



The big empirical success of group equivariant networks has led to recent years the growth of interest in their applications to equivariant learning and to a particular focus on developing more robust representations on CNNs. This paper gives a broad overview of the literature on equivariant CNNs. The theory of Steerable CNNs is also explained, as well as the connection to the convolution kernels which depend on more general groups.

Abstract

The big empirical success of group equivariant networks has led to recent years the growth of interest in their applications to equivariant learning and to a particular focus on developing more robust representations on CNNs. This paper gives a broad overview of the literature on equivariant CNNs. The theory of Steerable CNNs is also explained, as well as the connection to the convolution kernels which depend on more general groups.

Abstract

The main idea of Group equivariant Convolutional Neural Networks (G-CNNs) is to extend the notion of convolutional neural networks to groups. This makes it possible to express symmetries in G-CNNs via G-convolutions, a new type of layer that enjoys a substantially higher degree of weight sharing than regular convolution layers. G-convolutions are particularly well suited for equivariant learning, as they are able to learn features that are invariant under group actions. Group convolution layers are also able to learn features that are invariant under group actions. G-CNNs achieve state of the art results on CIFAR10 and rotated MNIST.

Abstract

The main idea of Group equivariant Convolutional Neural Networks (G-CNNs) is to extend the notion of convolutional neural networks to groups. This makes it possible to express symmetries in G-CNNs via G-convolutions, a new type of layer that enjoys a substantially higher degree of weight sharing than regular convolution layers. G-convolutions are particularly well suited for equivariant learning, as they are able to learn features that are invariant under group actions. G-CNNs achieve state of the art results on CIFAR10 and rotated MNIST.

Abstract

Summary: Symmetry is a fundamental concept in physics and mathematics. It is also a key feature of many natural and man-made systems. A group action is a continuous transformation of a set by a group. A group equivariant function is a function that respects the group action. In this paper, we develop a general framework for group equivariant functions. We show that group equivariant functions are universal approximators. We also prove that group equivariant functions are stable under composition. Finally, we show that group equivariant functions are universal approximators of group equivariant functions. This provides a unified framework for group equivariant functions and their applications in machine learning and computer vision.

Abstract

Summary: Symmetry is a fundamental concept in physics and mathematics. It is also a key feature of many natural and man-made systems. A group action is a continuous transformation of a set by a group. A group equivariant function is a function that respects the group action. In this paper, we develop a general framework for group equivariant functions. We show that group equivariant functions are universal approximators. We also prove that group equivariant functions are stable under composition. Finally, we show that group equivariant functions are universal approximators of group equivariant functions. This provides a unified framework for group equivariant functions and their applications in machine learning and computer vision.

Abstract

Summary: Symmetry is a fundamental concept in physics and mathematics. It is also a key feature of many natural and man-made systems. A group action is a continuous transformation of a set by a group. A group equivariant function is a function that respects the group action. In this paper, we develop a general framework for group equivariant functions. We show that group equivariant functions are universal approximators. We also prove that group equivariant functions are stable under composition. Finally, we show that group equivariant functions are universal approximators of group equivariant functions. This provides a unified framework for group equivariant functions and their applications in machine learning and computer vision.

Abstract

Summary: Symmetry is a fundamental concept in physics and mathematics. It is also a key feature of many natural and man-made systems. A group action is a continuous transformation of a set by a group. A group equivariant function is a function that respects the group action. In this paper, we develop a general framework for group equivariant functions. We show that group equivariant functions are universal approximators. We also prove that group equivariant functions are stable under composition. Finally, we show that group equivariant functions are universal approximators of group equivariant functions. This provides a unified framework for group equivariant functions and their applications in machine learning and computer vision.

Abstract

Summary: Symmetry is a fundamental concept in physics and mathematics. It is also a key feature of many natural and man-made systems. A group action is a continuous transformation of a set by a group. A group equivariant function is a function that respects the group action. In this paper, we develop a general framework for group equivariant functions. We show that group equivariant functions are universal approximators. We also prove that group equivariant functions are stable under composition. Finally, we show that group equivariant functions are universal approximators of group equivariant functions. This provides a unified framework for group equivariant functions and their applications in machine learning and computer vision.


Abstract

Summary: Symmetry is a fundamental concept in physics and mathematics. It is also a key feature of many natural and man-made systems. A group action is a continuous transformation of a set by a group. A group equivariant function is a function that respects the group action. In this paper, we develop a general framework for group equivariant functions. We show that group equivariant functions are universal approximators. We also prove that group equivariant functions are stable under composition. Finally, we show that group equivariant functions are universal approximators of group equivariant functions. This provides a unified framework for group equivariant functions and their applications in machine learning and computer vision.

Abstract

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>




Cesa-Lang-Weiler 2022  
 $G = \mathbb{R}^d \rtimes H$  with  $H$  compact

<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>




Cesa-Lang-Weiler 2022  
 $G = \mathbb{R}^d \rtimes H$  with  $H$  compact

<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>



<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>

NNs →

Published as a conference paper at ICLR 2018

## HEXAConv

Emiel Hoogeboom\*, Jorn W.T. Peters\* & Taco S. Cohen  
University of Amsterdam  
[e.hoogeboom, j.w.t.peters, t.s.cohen]@uva.nl

Max Welling  
University of Amsterdam & CIFAR  
m.welling@uva.nl

### ABSTRACT

The effectiveness of convolutional neural networks stems in large part from their ability to exploit the translation invariance that is inherent in many learning problems. Recently, it was shown that CNNs can exploit other sources of invariance, such as rotation invariance, by using *group convolutions* instead of planar convolutions. However, for reasons of performance and ease of implementation, it has been necessary to limit the group convolution to transformations that can be applied to the filters without interpolation. Thus, for images with square pixels, only integer translations, rotations by multiples of 90 degrees, and reflections are admissible.

Whereas the square tiling provides a 4-fold rotational symmetry, a hexagonal tiling of the plane has a 6-fold rotational symmetry. In this paper we show how one can efficiently implement planar convolution and group convolution over hexagonal lattices, by reusing existing highly optimized convolution routines. We find that, due to the reduced anisotropy of hexagonal filters, planar HexaConv provides better accuracy than planar convolution with square filters, given a fixed parameter budget. Furthermore, we find that the increased degree of symmetry of the hexagonal grid increases the effectiveness of group convolutions, by allowing for more parameter sharing. We show that our method significantly outperforms conventional CNNs on the AID aerial scene classification dataset, even outperforming ImageNet pretrained models.

### 1 INTRODUCTION

For sensory perception tasks, neural networks have mostly replaced handcrafted features. Instead of defining features by hand using domain knowledge, it is now possible to learn them, resulting in improved accuracy and saving a considerable amount of work. However, successful generalization is still critically dependent on the inductive bias encoded in the network architecture, whether this bias is understood by the network architect or not.

The canonical example of a successful network architecture is the Convolutional Neural Network (CNN, ConvNet). Through convolutional weight sharing, these networks exploit the fact that a given visual pattern may appear in different locations in the image with approximately equal likelihood. Furthermore, this translation symmetry is preserved throughout the network because a translation of the input image leads to a translation of the feature maps at each layer: convolution is translation equivariant.

Very often, the true label function (the mapping from image to label that we wish to learn) is invariant to more transformations than just translations. Rotations are an obvious example, but standard translational convolutions cannot exploit this symmetry, because they are not rotation equivariant. As it turns out, a convolution operation can be defined for almost any group of transformation — not just translations. By simply replacing convolutions with group convolutions (wherein filters are not

\*Equal contribution





Figure 4: Hexagonal convolution filters (left) represented in 2D memory (right) for filters of size three (blue) and five (blue and green). Standard 2D convolution using both feature map and filter stored according to the coordinate system is equivalent to convolution on the hexagonal lattice. Note that for the offset coordinate system two separate planar convolution are required — one for even and one for odd rows.





Figure 6: CIFAR-10 (top) and AID (bottom) examples sampled from Cartesian to hexagonal axial coordinates. Zero padding enlarges the images in axial systems.

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>




## Cesa-Lang-Weiler 2022 $\mathbb{R}^d \rtimes H$ with $H$ compact

<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>



# Cesa-Lang-Weiler 2022

## $\mathbb{R}^d \times H$ with $H$ compact

<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

en-Cai-OSU/awesome-gcnn

## Roto-Translation Covariant Convolutional Networks for Medical Image Analysis

Erik J Bekkers<sup>a,1</sup>, Maxime W Lafarge<sup>a,2</sup>, Mitko Veta<sup>a</sup>, Koen AJ Eppenhof<sup>b</sup>, Josien PW Pluim<sup>a</sup> and Remco Duits<sup>a</sup>

Eindhoven University of Technology, <sup>a</sup>Department of Mathematics and Computer Science and <sup>b</sup>Department of Biomedical Engineering, Eindhoven, The Netherlands

<sup>a</sup>Joint main authors - e.j.bekkers@tue.nl, m.w.lafarge@tue.nl

**Abstract.** We propose a framework for rotation and translation covariant deep learning using  $SE(2)$  group convolutions. The group product of the special Euclidean motion group  $SE(2)$  describes how a concatenation of two roto-translations results in a net roto-translation. We encode this geometric structure into convolutional neural networks (CNNs) via  $SE(2)$  group convolutional layers, which fit into the standard 2D CNN framework, and which allow to generically deal with rotated input samples without the need for data augmentation.

We introduce three layers: a *lifting layer* which lifts a 2D (vector valued) image to an  $SE(2)$ -image, i.e., 3D (vector valued) data whose domain is  $SE(2)$ ; a *group convolution layer* from and to an  $SE(2)$ -image; and a *projection layer* from an  $SE(2)$ -image to a 2D image. The lifting and group convolution layers are  $SE(2)$  covariant (the output roto-translates with the input). The final projection layer, a maximum intensity projection over rotations, makes the full CNN rotation invariant.

We show with three different problems in histopathology, retinal imaging, and electron microscopy that with the proposed group CNNs, state-of-the-art performance can be achieved, without the need for data augmentation by rotation and with increased performance compared to standard CNNs that do rely on augmentation.

**Keywords:** Group convolutional network, roto-translation group, mitosis detection, vessel segmentation, cell boundary segmentation

### 1 Introduction

In this work we generalize  $\mathbb{R}^2$  convolutional neural networks (CNNs) to  $SE(2)$  group CNNs (G-CNNs) in which the data lives on position orientation space, and in which the convolution layers are defined in terms of representations of the special Euclidean motion group  $SE(2)$ . In essence this means that we replace the convolutions (with translations of a kernel) by  $SE(2)$  group convolutions (with roto-translations of a kernel). The advantage of the proposed approach compared to standard  $\mathbb{R}^2$  CNNs is that rotation covariance is encoded in the network design and does not have to be learned by the convolution kernels. E.g., a feature that may appear in the data under several orientations does not have

Yanruo Zhou<sup>1</sup>, Qixing Ye<sup>1</sup>, Qianqian Guo<sup>1</sup> and Huikai Huo<sup>1</sup>  
<sup>1</sup>University of Chinese Academy of Sciences, Beijing, China  
On the Generalization of Equivariance and Convolution in Neural Networks to the Action of Compact Groups

Carlo Belotti<sup>1</sup>, Christos Alou Blandeau<sup>1</sup>, Ameri GRASP Laboratory, University of Pennsylvania, Philadelphia, PA, USA  
A Roto-Translation Covariant Convolutional Network for Medical Image Analysis

Ricardo Krebs<sup>1</sup>, Shubhabrata Ghosh<sup>1</sup>  
Roto-Translation Covariant Convolutional Network for Medical Image Analysis

MEDICAL IMAGE ANALYSIS  
Journal of Medical Imaging and Biometrics  
Journal of Medical Imaging and Biometrics

Medical Image Analysis  
journal homepage: www.elsevier.com/locate/media

## Roto-translation equivariant convolutional networks: Application to histopathology image analysis

Maxime W. Lafarge<sup>a,\*</sup>, Erik J. Bekkers<sup>b</sup>, Josien P.W. Pluim<sup>a</sup>, Remco Duits<sup>b</sup>, Mitko Veta<sup>a</sup>

<sup>a</sup>Department of Biomedical Engineering, Eindhoven University of Technology, Eindhoven, the Netherlands

<sup>b</sup>Department of Mathematics and Computer Science, Eindhoven University of Technology, Eindhoven, the Netherlands

### ARTICLE INFO

#### Article history:

Received 4 December 2019

Revised 3 June 2020

Accepted 14 August 2020

Available online 31 October 2020

#### Keywords:

Group convolutional neural network  
Roto-translation equivariance  
Computational pathology  
Mitosis detection  
Tumor detection  
Nuclei segmentation

### ABSTRACT

Rotation-invariance is a desired property of machine-learning models for medical image analysis and in particular for computational pathology applications. We propose a framework to encode the geometric structure of the special Euclidean motion group  $SE(2)$  in convolutional networks to yield translation and rotation equivariance via the introduction of  $SE(2)$ -group convolution layers. This structure enables models to learn feature representations with a discretized orientation dimension that guarantees that their outputs are invariant under a discrete set of rotations.


Conventional approaches for rotation invariance rely mostly on data augmentation, but this does not guarantee the robustness of the output when the input is rotated. At that, trained conventional CNNs may require test-time rotation augmentation to reach their full capability.

This study is focused on histopathology image analysis applications for which it is desirable that the arbitrary global orientation information of the imaged tissues is not captured by the machine learning models. The proposed framework is evaluated on three different histopathology image analysis tasks (mitosis detection, tumor detection, nuclei segmentation) and provides a comparative analysis for each problem when using the proposed framework.

**Author(s).** Published by Elsevier B.V. This article under the CC BY-NC-ND license [creativecommons.org/licenses/by-nc-nd/4.0/](https://creativecommons.org/licenses/by-nc-nd/4.0/)

experiments that we used to analyze and validate them. In the construction of the G-CNNs we adhere to the following principle of group equivariant architecture design.

**G-CNN design principle** A sequence of layers starting with a lifting layer (Eq. (7)) and followed by one or more group convolution layers (Eq. (9)), possibly intertwined with point-wise non-linearities, results in the encoding of roto-translation equivariant feature maps. If such a block is followed by a projection layer (Eq. (10)) then the entire block results in an encoding of features that is guaranteed to be rotationally invariant. Our implementation of the G-CNN layers is available at <https://github.com/tueimage/se2cnn>.




**Fig. 2.** Illustration of the process generating a rotated set of effective kernels from a trainable vector of base weights via the introduction of fixed interpolation matrix in the computational pipeline.

<https://quva-lab.github.io/esCNN/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>




## Cesa-Lang-Weiler 2022 $\mathbb{R}^d \rtimes H$ with $H$ compact

<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>



Cesa-Lang-Weiler 2022  
 $G = \mathbb{R}^d \rtimes H$  with  $H$  compact

<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

CVF

This CVPR paper is the Open Access version, provided by the Computer Vision Foundation.  
 Except for this watermark, it is identical to the version available on IEEE Xplore.

# Harmonic Networks: Deep Translation and Rotation Equivariance

Daniel E. Worrall, Stephan J. Garbin, Daniyar Tumukhambetov and Gabriel J. Brostow

{d.worrall, s.garbin, d.tumukhambetov, g.brostow}@cs.ucl.ac.uk

University College London\*

## Abstract

Translating or rotating an input image should not affect the results of many computer vision tasks. Convolutional neural networks (CNNs) are already translation equivariant: input image translations produce proportionate feature map translations. This is not the case for rotations. Global rotation equivariance is typically sought through data augmentation, but patch-wise equivariance is more difficult. We present Harmonic Networks or H-Nets, a CNN exhibiting equivariance to patch-wise translation and 360°-rotation. We achieve this by replacing regular CNN filters with circular harmonics, returning a maximal response and orientation for every receptive field patch.

H-Nets use a rich, parameter-efficient and fixed computational complexity representation, and we show that deep feature maps within the network encode complicated rotational invariants. We demonstrate that our layers are general enough to be used in conjunction with the latest architectures and techniques, such as deep supervision and batch normalization. We also achieve state-of-the-art classification on rotated-MNIST, and competitive results on other benchmark challenges.

## 1. Introduction

representing 360°-rotations (CNNs) [19]. Currently, by design to map an image to rotated versions of the image map the same feature vector [Figure 1]. However, until now, when the feature vectors do not have a clear or easy to predict manner relating input transformations, is called *equivariance*.

*Equivariance*, where feature maps are invariant to global transformations of the input, is a key constraint for a model, such as a CNN, to restrict all intermediate representations to be invariant. For example,

Figure 1. Patch-wise translation equivariance in CNNs arises from translational weight tying, so that a translation  $\pi$  of the input image  $I$ , leads to a corresponding translation  $\psi$  of the feature maps  $f(I)$ , where  $\pi \neq \psi$  in general, due to pooling effects. However, for rotations, CNNs do not yet have a feature space transformation  $\psi$  ‘hard-baked’ into their structure, and it is complicated to discover what  $\psi$  may be, if it exists at all. Harmonic Networks have a hard-baked representation, which allows for easier interpretation of feature maps—see Figure 3.

consider detecting a deformable object, such as a butterfly. The pose of the wings is limited in range, and so there are only certain poses our detector should normally see. A transformation invariant detector, good at detecting wings, would detect them whether they were bigger, further apart, rotated, etc., and it would encode all these cases with the same representation. It would fail to notice nonsensical situations, however, such as a butterfly with wings rotated past the usual range, because it has thrown that extra pose information away. An equivariant detector, on the other hand, does not dispose of local pose information, and so it hands on a richer and more useful representation to downstream processes. Equivariance conveys more information about an input to downstream processes, it also constrains the space of possible learned models to those that are valid under the rules of natural image formation [30]. This makes learning more reliable and helps with generalization. For instance, consider CNNs. The key insight is that the statistics of natural images, embodied in the correlations between pixels, are a) invariant to translation, and b) highly localized. Thus features at every layer in a CNN are computed on local receptive fields, where weights are shared

Figure 2. Real and imaginary parts of the complex Gaussian filter  $W_m(r, \phi) e^{-r^2/2} = e^{-r^2} e^{im\phi}$ , for some rotation orders. As a simple example, we have set  $R(r) = e^{-r^2}$  and  $\beta = 0$ , but in general we learn these quantities. Cross-correlation of a feature map of rotation order  $n$  with one of these filters of rotation order  $m$ , results in a feature map of rotation order  $m+n$ . Note the negative rotation order filters have flipped imaginary parts compared to the positive orders.

feature maps, which live in a discrete domain. We shall instead use continuous spaces, because the analysis is easier. Later on in Section 4.2 we shall demonstrate how to convert back to the discrete domain for practical implementation, but for now we work entirely in continuous Euclidean space.

### 3.1. Equivariance

Equivariance is a useful property to have because transforma-

Figure 3. DOWN: Cross-correlation of the input patch with  $W_m$  yields a scalar complex-valued response. ACROSS-THEN-DOWN: Cross-correlation with the  $\theta$ -rotated image yields another complex-valued response. BOTTOM: We transform from the unrotated response to the rotated response, through multiplication by  $e^{im\theta}$ .

Here  $r, \phi$  are the spatial coordinates of image/feature maps, expressed in polar form,  $m \in \mathbb{Z}$  is known as the *rotation order*.

5028

**Antipov et al. 2020** Translation-Equivariant Convolutional Networks for Application to Histopathology

**Rotation-equivariant CNNs (manifolds)**

**Cohen et al. 2019** Icosahedron

**de Haan et al. 2020** Meshes

**Bekkers 2019** Lie groups (rotation, scale)

**Finzi et al. 2020** Lie groups

**Chakraborty et al. 2018** Riemannian Hom. spaces

**Weiler-Cesa 2019** SE(2)

**Weiler et al. 2018** SE(3)

**Weiler et al. 2018** SE(3)

**2018** *or-Trivedi 2018*

**2018**

**Continuous G-CNNs (Lie groups)**

**Weiler et al. 2021** Coordinate Independent CNNs

**Weiler et al. 2020** Scale-translation

**Sosnovik et al. 2020** Scale-translation

**Antipov et al. 2020** Translation-Equivariant Convolutional Networks for Application to Histopathology

**Group-Equivariant CNNs**

**Marina W. Unger<sup>1</sup>, Erik J. Bekkers<sup>2</sup>, Justin P.W. Pieters<sup>3</sup>**

<sup>1</sup> Department of Pathology  
Donders Institute for Medical  
Research and Technology  
University of Maastricht  
The Netherlands

<sup>2</sup> Department of Pathology  
Donders Institute for Medical  
Research and Technology  
University of Maastricht  
The Netherlands

<sup>3</sup> Department of Pathology  
Donders Institute for Medical  
Research and Technology  
University of Maastricht  
The Netherlands

**Abstract**

We introduce Group-equivariant Convolutional Kernel Networks (G-CNNs), a natural generalization of convolutional neural networks that reduces model complexity by exploiting symmetries. G-CNNs use G-convolutions, a new “spectral” layer that enjoys a substantially higher degree of weight sharing than regular convolution layers. G-convolutions merge the expressive capacity of G-networks with the low cost of learning the order of operations. Group-equivariant layers are easy to construct and can be implemented with orthogonal convolutional kernels. G-equivariant group actions are generated by translations, reflections, and rotations. G-CNNs achieve state-of-the-art results on CIFAR-10 and CIFAR-100.

**1. Introduction**

Learning to invariantly detect properties of manifolds during training is a challenging task, due to its generality and inherent difficulties in dealing with global symmetries. For example, the problem of identifying previously unseen objects in a pre-positioned sequence of images is often addressed by learning local features that are invariant to global orientation and scale. This is done by extracting features in every location of the input image. However, this approach is not able to handle more complex transformations such as rotation and global scaling. It has been shown that the model performance is significantly improved by learning features that are invariant to global scaling. Although data augmentation provides a good solution to this problem, we propose

**2. Motivation**

Convolutional neural networks (CNNs) have proven to be very powerful models of sensory data such as images, video, and audio. Although using theories of neural networks design is currently lacking, a large amount of empirical evidence supports the notion that both overparameterized learning and depth (among other factors) are important for or predictive performance.

**3. Motivation**


Convolutional weight sharing is effective because there is a considerable symmetry in most practical tasks. The location and data distribution are both approximately invariant to translation. By using the same weights to analyze different parts of the image, a convolution layer takes advantage of this invariance, while performing fewer computations compared to non-equi-

© 2021 International Conference on Machine Learning (ICML). 978-1-4503-8869-6/21 © 2021 PMLR. All rights reserved. CC-BY 4.0 International license (http://creativecommons.org/licenses/by/4.0/).

*Cesa-Lang-Weiler 2022*  
 $G = \mathbb{R}^d \rtimes H$  with  $H$  compact

<https://quva-lab.github.io/escnn/>

# A brief history of China



<https://quva-lab.github.io/escnn/>

Published as a conference paper at ICLR 2022

## A PROGRAM TO BUILD $E(n)$ -EQUIVARIANT STEERABLE CNNs

**Gabriele Cesa**  
Qualcomm AI Research\*  
University of Amsterdam  
gcesa@qti.qualcomm.com

**Leon Lang**  
University of Amsterdam  
l.lang@uva.nl

**Maurice Weiler**  
University of Amsterdam  
m.weiler.m@gmail.com

### ABSTRACT

Equivariance is becoming an increasingly popular design choice to build data efficient neural networks by exploiting prior knowledge about the symmetries of the problem at hand. Euclidean steerable CNNs are one of the most common classes of equivariant networks. While the constraints these architectures need to satisfy are understood, existing approaches are tailored to specific (classes of) groups. No generally applicable method that is *practical* for implementation has been described so far. In this work, we generalize the Wigner-Eckart theorem proposed in Lang & Weiler (2020), which characterizes general  $G$ -steerable kernel spaces for compact groups  $G$  over their homogeneous spaces, to arbitrary  $G$ -spaces. This enables us to directly parameterize filters in terms of a band-limited basis on the whole space rather than on  $G$ 's orbits, but also to easily implement steerable CNNs equivariant to a large number of groups. To demonstrate its generality, we instantiate our method on a variety of isometry groups acting on the Euclidean space  $\mathbb{R}^3$ . Our framework allows us to build  $E(3)$  and  $SE(3)$ -steerable CNNs like previous works, but also CNNs with arbitrary  $G \leq O(3)$ -steerable kernels. For example, we build 3D CNNs equivariant to the symmetries of platonic solids or choose  $G = SO(2)$  when working with 3D data having only azimuthal symmetries. We compare these models on 3D shapes and molecular datasets, observing improved performance by matching the model's symmetries to the ones of the data.

### 1 INTRODUCTION

In machine learning, it is common for learning tasks to present a number of *symmetries*. A symmetry in the data occurs, for example, when some property (e.g., the label) does not change if a set of transformations is applied to the data itself, e.g., translations or rotations of images. Symmetries are algebraically described by *groups*. If prior knowledge about the symmetries of a task is available, it is usually beneficial to encode them in the models used (Shawe-Taylor 1989; Cohen & Welling 2016a). The property of such models is referred to as *equivariance* and is obtained by introducing some *equivariance constraints* in the architecture (see Eq. 2). A classical example are convolutional neural networks (CNNs), which achieve translation equivariance by constraining linear layers to be convolution operators. A wider class of equivariant models are Euclidean steerable CNNs (Cohen & Welling 2016b; Weiler et al. 2018a; Weiler & Cesa, 2019; Jenner & Weiler, 2022), which guarantee equivariance to isometries  $\mathbb{R}^n \times G$  of a Euclidean space  $\mathbb{R}^n$ , i.e., to translations and a group  $G$  of origin-preserving transformations, such as rotations and reflections. As proven in Weiler et al. (2018a; 2021); Jenner & Weiler (2022), this requires convolutions with *G-steerable* (equivariant) kernels.

Our goal is developing a program to parameterize with minimal requirements arbitrary  $G$ -steerable kernel spaces, with compact  $G$ , which are required to implement  $\mathbb{R}^n \times G$  equivariant CNNs. Lang & Weiler (2020) provides a first step in this direction by generalizing the *Wigner-Eckart theorem* from quantum mechanics to obtain a general technique to parametrize  $G$ -steerable kernel spaces over *orbits* of a compact  $G$ . The theorem reduces the task of building steerable kernel bases to that of finding some pure representation theoretic ingredients. Since the equivariance constraint only relates points  $g \cdot x \in \mathbb{R}^n$  in the same *orbit*  $G \cdot x \in \mathbb{R}^n$ , a kernel can take independent values on different orbits. Fig. 1 shows

\*Qualcomm AI Research is an initiative of Qualcomm Technologies, Inc.

## README.md

### General $E(2)$ -Equivariant Steerable CNNs

[Documentation](#) | [Experiments](#) | [Paper](#) | [Thesis](#)

e2cnn is a PyTorch extension for equivariant deep learning.

*Equivariant neural networks* guarantee a specified transformation behavior of their feature spaces under transformations of their input. For instance, classical convolutional neural networks (CNNs) are by design equivariant to translations of their input. This means that a translation of an image leads to a corresponding translation of the network's feature maps. This package provides implementations of neural network modules which are equivariant under all *isometries*  $E(2)$  of the image plane  $\mathbb{R}^2$ , that is, under *translations*, *rotations* and *reflections*. In contrast to conventional CNNs,  $E(2)$ -equivariant models are guaranteed to generalize over such transformations, and are therefore more data efficient.

The feature spaces of  $E(2)$ -Equivariant Steerable CNNs are defined as spaces of feature fields, being characterized by gray-

### Getting Started

e2cnn is easy to use since it provides a high level user interface which abstracts most intricacies of group and representation theory away. The following code snippet shows how to perform an equivariant convolution from an RGB-image to 10 *regular* feature fields (corresponding to a [group convolution](#)).

```
from e2cnn import gspaces
from e2cnn import nn
import torch

r2_act = gspaces.Rot2dOnR2(N=8)
feat_type_in = nn.FieldType(r2_act, 3*[r2_act.trivial_repr])
feat_type_out = nn.FieldType(r2_act, 10*[r2_act.regular_repr])

conv = nn.R2Conv(feat_type_in, feat_type_out, kernel_size=5)
relu = nn.ReLU(feat_type_out)

x = torch.randn(16, 3, 32, 32)
x = nn.GeometricTensor(x, feat_type_in)


y = relu(conv(x))
```

Line 5 specifies the symmetry group action on the image plane  $\mathbb{R}^2$  under which the network should be equivariant. We choose the *cyclic group*  $C_8$ , which describes discrete rotations by multiples of  $2\pi/8$ . Line 6 specifies the input feature field types. The three color channels of an RGB image are thereby to be identified as three independent scalar fields, which transform under the *trivial representation* of  $C_8$ . Similarly, the output feature space is in line 7

<https://quva-lab.github.io/eschn/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>




# Cesa-Lang-Weiler 2022

## $\mathbb{R}^d \times H$ with $H$ compact

<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>



# Cesa-Lang-Weiler 2022


## $\mathbb{R}^d \times H$ with $H$ compact

<https://quva-lab.github.io/escnn/>



# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>




## Cesa-Lang-Weiler 2022 $\mathbb{R}^d \rtimes H$ with $H$ compact

<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>



Cesa-Lang-Weiler 2022  
 $G = \mathbb{R}^d \rtimes H$  with  $H$  compact

<https://quva-lab.github.io/escnn/>





Figure 1: Many modalities of spatial data are null spaces important symmetries. We propose a more general framework that can support a great variety of symmetry groups.

In addition, we can support more modalities. A group action on a manifold, used in group convolutional networks (Cohen and Welling, 2016), is a continuous action of our group on the manifold. In this case, the convolution operation is implemented as a Lie group, and Hamiltonian dynamics. For Hamiltonian systems, the evolution of our modality is especially intricate, leading to exact conservation of linear and angular momentum.

**1. Introduction**  
 Symmetry pervades the natural world. The same law of gravitation governs a grain of sand, the orbit of an planet, and the formation of galaxies. It is precisely because of the role of the universe that we can hope to understand it. Once we started to understand the symmetries inherent in physical laws, we could predict behavior in galaxies billions of light years away by studying our own local region, a star and galaxy. One of the first tools to achieve these full potential, it is essential to incorporate our knowledge of naturally existing symmetries into the design of algorithms and architectures. An example of this principle is the convolution operation of convolutional neural networks (LeCun et al., 1989). When an input (e.g. an image) is processed, the output of a convolutional layer is localized in the same way.

Group theory provides mathematical language for symmetry and equivariance. Convolutional layers are equivariant under the modality we use for convolutional layers and transformation groups and data types.




Figure 1: Many realizations of spatial coordinate systems will possess important symmetries. A group equivariant convolutional network can be used to automatically learn invariant features of our data. Our architecture is local, and handles both Euclidean and non-Euclidean systems. It is especially well suited for learning

in dimensions, and on a greater level in time. A group equivariant convolutional network is a generalization of a group equivariant convolutional network (Cohen and Welling, 2016). In this paper, we develop a general framework for convolutional networks on arbitrary continuous data. A group equivariant convolutional network is a generalization of a group equivariant convolutional network (Cohen and Welling, 2016).

In this paper, we develop a general framework for convolutional networks on arbitrary continuous data. A group equivariant convolutional network is a generalization of a group equivariant convolutional network (Cohen and Welling, 2016). In this paper, we develop a general framework for convolutional networks on arbitrary continuous data. A group equivariant convolutional network is a generalization of a group equivariant convolutional network (Cohen and Welling, 2016).

The same line of thought can also be applied to other types of continuous data, such as signals in polar coordinates. As  $G$ -structures are required to be continuous, we removed the origin 0 where polar coordinates are singular. One can once again define an  $\mathcal{R}$ -structure by adding reflected frames as shown in Fig. 5f. These  $G$ -structures model convolutions on  $\mathbb{R}^2 \setminus \{0\}$  which are  $SO(2)$  and  $O(2)$  equivariant but not translation equivariant. Fig. 5h shows the usual  $SO(2)$ -structure on the embedded 2-sphere  $S^2$ , which is underlying  $SO(3)$ -equivariant spherical CNNs. Another popular choice is the  $\{\epsilon\}$ -structure in Fig. 5k, which is induced by spherical coordinates. Note that this  $\{\epsilon\}$ -structure would be singular at the poles, which are therefore cut out. Continuous (i.e. non-singular) reductions of the structure group beyond  $SO(2)$  are on the sphere topologically obstructed.  $G$  steerable kernels with  $G \geq SO(2)$  are therefore strictly necessary for continuous convolutions on topological spheres like the mesh in Fig. 5i. Fig. 5j shows an  $\mathcal{R}$ -structure on the Möbius strip. As the Möbius strip is non-orientable, it does not admit a continuous reduction of the structure group beyond the reflection group  $G = \mathcal{R}$ .

continuous symmetries are equivariant.

# History of Convolutional Networks

## /Chen-Cai-OSU/

### COORDINATE INDEPENDENT CONVOLUTIONAL NETWORKS ISOMETRY AND GAUGE EQUIVARIANT CONVOLUTIONS ON RIEMANNIAN MANIFOLDS

Maurice Weiler  
University of Amsterdam  
m.weiler.ml@gmail.com


Patrick Forré  
University of Amsterdam  
p.d.forre@uva.nl

Erik Verlinde  
University of Amsterdam  
e.p.verlinde@uva.nl

Max Welling  
University of Amsterdam,  
Qualcomm AI Research  
m.welling@uva.nl

#### ABSTRACT

Motivated by the vast success of deep convolutional networks, there is a great interest in generalizing convolutions to non-Euclidean manifolds. A major complication in comparison to flat spaces is that it is unclear in which alignment a convolution kernel should be applied on a manifold. The underlying reason for this ambiguity is that general manifolds do not come with a canonical choice of reference frames (gauge). Kernels and features therefore have to be expressed relative to *arbitrary coordinates*. We argue that the particular choice of coordinatization should not affect a network's inference – it should be *coordinate independent*. A simultaneous demand for coordinate independence and weight sharing is shown to result in a requirement on the network to be *equivariant under local gauge transformations* (changes of local reference frames). The ambiguity of reference frames depends thereby on the *G-structure* of the manifold, such that the necessary level of gauge equivariance is prescribed by the corresponding *structure group*  $G$ . Coordinate independent convolutions are proven to be equivariant w.r.t. those *isometries* that are symmetries of the  $G$ -structure. The resulting theory is formulated in a coordinate free fashion in terms of fiber bundles. To exemplify the design of coordinate independent convolutions, we implement a convolutional network on the Möbius strip. The generality of our differential geometric formulation of convolutional networks is demonstrated by an extensive literature review which explains a large number of Euclidean CNNs, spherical CNNs and CNNs on general surfaces as specific instances of coordinate independent convolutions.



Chen 2018  
roto-transl.

Hoogeboom et al. 2018  
 $SE(2,6)$

Worrall-Brostow 2018  
3D discrete roto-t

#### G-CNNs

Cohen et al. 2017  
 $SE(2)$

Esteves 2017  
 $SO(3)$

Kondor 2018  
 $SE(3)$

Thomas et al. 2017  
 $SE(3)$

Continuous  
(spherical)

<https://qava.csail.mit.edu/conv/>










(a)  $M = \mathbb{R}^2$ ,  $G = \{e\}$ (b)  $M = \mathbb{R}^2$ ,  $G = \{e\}$ (c)  $M = \mathbb{R}^2 \setminus \{0\}$ ,  $G = \{e\}$ (d)  $M = \mathbb{R}^2$ ,  $G = \mathcal{R}$ (e)  $M = \mathbb{R}^2$ ,  $G = \mathcal{R}$ (f)  $M = \mathbb{R}^2$ ,  $G = \mathcal{S}$ (g)  $M = \mathbb{R}^2$ ,  $G = SO(2)$ (h)  $M = S^2$ ,  $G = SO(2)$ (i)  $M = \text{"Suzanne"}$ ,  $G = SO(2)$ (j)  $M = S^2 \setminus \text{poles}$ ,  $G = \{e\}$ (k)  $M = \text{M\"obius}$ ,  $G = \mathcal{R}$ (l)  $M = \text{M\"obius}$ ,  $G = \mathcal{R}$ 

Figure 5: Exemplary  $G$ -structures  $GM$  for different structure groups  $G$  and on different manifolds  $M$ . The structure group  $G$  signals which values gauge transformations can take, and therefore how “big” the subset of distinguished frames at each point  $p$  is. Fig. 5a shows the canonical  $\{e\}$ -structure (frame field) of  $\mathbb{R}^2$ , which corresponds to conventional Euclidean CNNs. The  $G$ -structures in Figs. 5d, 5g and 5j are constructed by adding reflected ( $G = \mathcal{R}$ ), rotated ( $G = SO(2)$ ) and scaled ( $G = \mathcal{S}$ ) frames, respectively. The corresponding  $GM$ -convolutions are not only translation equivariant but equivariant under the action of affine groups  $Aff(G)$ .  $G$ -structures are usually not unique. Figs. 5b and 5e show alternative  $G$  structures on  $\mathbb{R}^2$  (corresponding to an alternative metric relative to which their frames are orthonormal). They might not be practically relevant but demonstrate the flexibility of our framework. The  $\{e\}$ -structure in Fig. 5c corresponds to polar coordinates. As  $G$ -structures are required to be continuous, we removed the origin 0 where polar coordinates are singular. One can once again define an  $\mathcal{R}$ -structure by adding reflected frames as shown in Fig. 5f. These  $G$ -structures model convolutions on  $\mathbb{R}^2 \setminus \{0\}$  which are  $SO(2)$  and  $O(2)$  equivariant but not translation equivariant. Fig. 5h shows the usual  $SO(2)$ -structure on the embedded 2-sphere  $S^2$ , which is underlying  $SO(3)$ -equivariant spherical CNNs. Another popular choice is the  $\{e\}$ -structure in Fig. 5k, which is induced by spherical coordinates. Note that this  $\{e\}$ -structure would be singular at the poles, which are therefore cut out. Continuous (i.e. non-singular) reductions of the structure group beyond  $SO(2)$  are on the sphere topologically obstructed.  $G$  steerable kernels with  $G \geq SO(2)$  are therefore strictly necessary for continuous convolutions on topological spheres like the mesh in Fig. 5i. Fig. 5j shows an  $\mathcal{R}$ -structure on the Möbius strip. As the Möbius strip is non-orientable, it does not admit a continuous reduction of the structure group beyond the reflection group  $G = \mathcal{R}$ .









# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>




Cesa-Lang-Weiler 2022  
 $G = \mathbb{R}^d \rtimes H$  with  $H$  compact

<https://quva-lab.github.io/escnn/>



# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>




G-CNNs are NNs under  
continuous rotation G-CNNs  
(singular discretizations)

<https://quva-lab.github.io/escnn/>

# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>



G-CNNs are NNs under  
equivariance constraints

Continuous rotation G-CNNs  
(singular discriminative)

Continuous rotation G-CNNs  
(singular discriminative)

Continuous G-CNNs  
(Lie groups)

**"Inductive bias"**

Reduce the search space for NNs  
to only the sensible ones!

<https://equvalab.github.io/escnn/>




# A brief history of G-CNNs

<https://github.com/Chen-Cai-OSU/awesome-equivariant-network>

**Up next:**

A linear map is equivariant  
if and only if it is a group convolution



G-CNNs are NNs under  
equivariance constraints

**"Inductive bias"**

Reduce the search space for NNs  
to only the sensible ones!

<https://equivariant.github.io/escnn/>