

# Lecture 9: Explicit Generative Models

## Efstratios Gavves

# Lecture overview

---


- Gentle intro to Bayesian Modelling and Variational Inference
- Restricted Boltzmann Machines
- Deep Boltzmann Machines
- Deep Belief Network
- Contrastive Divergence
- Variational Autoencoders
- Normalizing Flows

# Explicit density models

- Plug in the model density function to likelihood
- Then maximize the likelihood


## Problems

- Design complex enough model that meets data complexity
- At the same time, make sure model is computationally tractable
- More details in the next lecture




Bayesian Modelling  
Variational Inference

**Deep Belief  
Network**



**Deep Boltzmann  
Machine**




# How to define a generative model?

- We can define an explicit density function over all possible relations  $\psi_c$  between the input variables  $x_c$

$$p(x) = \prod_c \psi_c (x_c)$$


$$\psi_c(x_1, x_2) = \psi_1 x_1 x_2$$

- Quite inefficient → think of all possible relations (not just pairwise) between  $256 \times 256 = 65K$  input variables
- Solution: Define an energy function to model the relations between the inputs variables



# Restricted Boltzmann Machines

0 1



- Boltzmann (or Gibbs) distribution defined over a free energy function  $E(x)$

$$p(x) = \frac{1}{Z} \exp(-E(x))$$

- $Z$  is the normalization factor that makes sure  $\int_x p(x) dx = 1$ 
  - Very expensive to compute → if  $x = \{0, 1\}$  computing  $Z$  requires  $2^d$  computations
- Better restrict the model further to a bottleneck

$$E(x) = -x^T W h - b^T x - c^T h$$

Pixel input  $x$   $x^T W x$  Feature maps

# Why Boltzmann?

---

- In statistical mechanics and mathematics, a Boltzmann distribution (also called Gibbs distribution) is a probability distribution, probability measure, or frequency distribution of particles in a system over various possible states. The distribution is expressed in the form

$$F(\text{state}) \propto \exp\left(-\frac{E}{kT}\right)$$

- $E$  is the state energy,  $k$  is the Boltzmann constant,  $T$  is the thermodynamic temperature

[https://en.wikipedia.org/wiki/Boltzmann\\_distribution](https://en.wikipedia.org/wiki/Boltzmann_distribution)

# Restricted Boltzmann Machines


---

- $E(x) = -x^T Wh - b^T x - c^T h$
- The  $x^T Wh$  models correlations between  $x$  and the latent activations via the parameter matrix  $W$
- The  $b^T x, c^T h$  model the priors
- Restricted Boltzmann Machines (RBM) assume  $\underline{x}, \underline{h}$  to be binary

# Restricted Boltzmann Machines

- $E(x) = -x^T Wh - b^T x - c^T h, \quad \theta = \{W, b, c\}$

- The free energy function  $F(x) = -\log \sum_h \exp(-E(x, h))$  defines a bipartite graph with undirected connections
  - Information flows forward and backward




# Restricted Boltzmann Machines

- The hidden units  $h_j$  are independent to each other conditioned on the visible units

$$p(\underline{h} | \underline{x}) = \prod_j p(h_j | x, \theta)$$

- The hidden units  $x_i$  are independent to each other conditioned on the visible units

$$p(\underline{x} | \underline{h}) = \prod_i p(x_i | h, \theta)$$



# Training RBMs

- The conditional probabilities are defined as sigmoids

$$p(h_j|x, \theta) = \sigma(W_{\cdot j}x + b_j)$$
$$p(x_i|h, \theta) = \sigma(W_{\cdot i}h + c_i)$$




- Maximize log-likelihood

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_n \log p(x_n|\theta)$$

$$\theta = \{w, b\}$$

- Let's take the gradients

$$\frac{\partial \log p(x_n|\theta)}{\partial \theta} = -\frac{\partial F(x_n)}{\partial \theta} - \frac{\partial \log Z}{\partial \theta}$$
$$= -\sum_h p(h|x_n, \theta) \frac{\partial E(x_n|h, \theta)}{\partial \theta} + \sum_{\tilde{x}, h} p(\tilde{x}, h|\theta) \frac{\partial E(\tilde{x}, h|\theta)}{\partial \theta}$$



Hidden unit (features)

# Training RBMs

- Let's take the gradients

$$\begin{aligned}\frac{\partial \log p(x_n|\theta)}{\partial \theta} &= -\frac{\partial F(x_n)}{\partial \theta} - \frac{\partial \log Z}{\partial \theta} \\ &= -\sum_h p(h|x_n, \theta) \frac{\partial E(x_n|h, \theta)}{\partial \theta} + \sum_{\tilde{x}, h} p(\tilde{x}, h|\theta) \frac{\partial E(\tilde{x}, h|\theta)}{\partial \theta}\end{aligned}$$

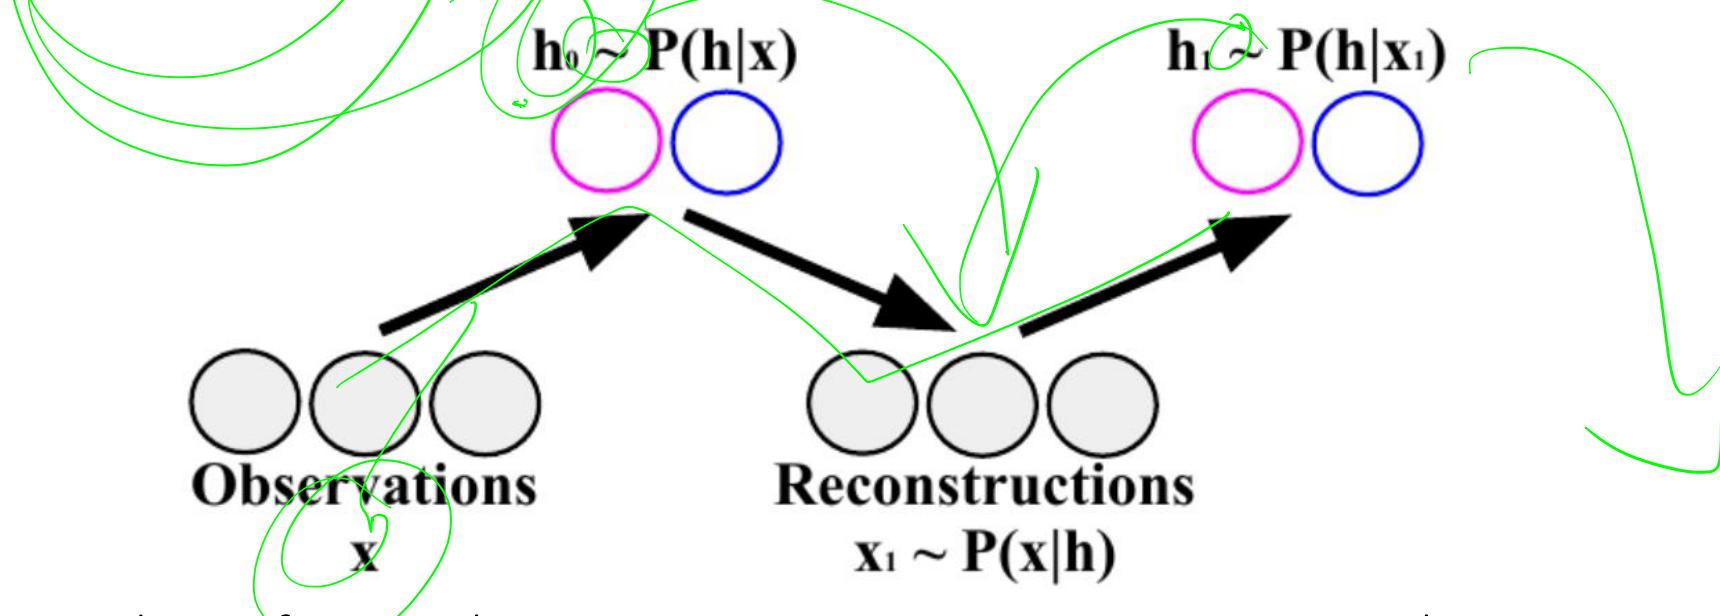
hard

- Easy because we just substitute in the definitions the  $x_n$  and sum over  $h$
- Hard because you need to sum over both  $\tilde{x}, h$  which can be huge
  - It requires approximate inference, e.g., MCMC

# Training RBMs with Contrastive Divergence

- Approximate the gradient with Contrastive Divergence
- Specifically, apply Gibbs sampler for  $k$  steps and approximate the gradient


$$\frac{\partial \log p(x_n|\theta)}{\partial \theta} = -\frac{\partial E(x_n, h_0|\theta)}{\partial \theta} - \frac{\partial E(x_k, h_k|\theta)}{\partial \theta}$$



Hinton, *Training Products of Experts by Minimizing Contrastive Divergence*, Neural Computation, 2002

# Deep Belief Network

- RBMs are just one layer
- Use RBM as a building block
- Stack multiple RBMs one on top of the other  
$$p(x, h_1, h_2) = p(x|h_1) \cdot p(h_1|h_2)$$
- Deep Belief Networks (DBN) are directed models
  - The layers are densely connected and have a single forward flow
  - This is because the RBN is directional,  $p(x_i|h, \theta) = \sigma(W_i^T x + c_i)$ , namely the input argument has only variable only from below




# Deep Boltzmann Machines

- Stacking layers again, but now with connection from the **above** and from the **below** layers
- Since it's a Boltzmann machine, we need an energy function

$$E(x, h_1, h_2 | \theta) = x^T W_1 h_1 + h_1^T W_2 h_2 + h_2^T W_3 h_3$$


$$p(h_2^k | h_1, h_3) = \sigma\left(\sum_j W_1^{jk} h_1^j + \sum_l W_3^{kl} h_3^l\right)$$



# Deep Boltzmann Machines


- Schematically similar to Deep Belief Networks
- But, Deep Boltzmann Machines (DBM) are undirected models
  - Belong to the Markov Random Field family
- So, two types of relationships: bottom-up and up-bottom

$$p(h_2^k | h_1, h_3) = \sigma \left( \sum_j W_1^{jk} h_1^j + \sum_l W_3^{kl} h_3^l \right)$$




# Training Deep Boltzmann Machines

- Computing gradients is intractable
- Instead, variational methods (mean-field) or sampling methods are used




# Bayesian Modelling Variational Inference



# Bayesian Terminology

- Observed variables  $x$
- Latent variables  $\theta$ 
  - Both unobservable model parameters  $w$  and unobservable model activations  $z$
  - $\theta = \{w, z\}$
- Joint probability density function (pdf):  $p(\underline{x}, \theta)$
- Marginal pdf:  $p(x) = \int_{\theta} p(x, \theta) d\theta$
- Prior pdf → marginal over input:  $p(\theta) = \int_x p(x, \theta) dx$ 
  - Usually a user defined pdf
- Posterior pdf:  $p(\theta|x)$
- Likelihood pdf:  $p(x|\theta)$



# Bayesian Terminology

- Posterior pdf

$$\begin{aligned} p(\theta|x) &= \frac{p(x|\theta)p(\theta)}{\int_{\theta'} p(x|\theta') d\theta'} \\ &= \frac{p(x|\theta)p(\theta)}{p(x)} \end{aligned}$$

← Conditional probability

← Bayes Rule

← Marginal probability

←  $p(x)$  is constant

*Bayes Rule*


$$p(x|\theta) = p(x|\theta)p(\theta)$$

- Posterior Predictive pdf

$$p(x_{new}|\bar{x}) = \int_{\theta} p(x_{new}|\theta)p(\theta|\bar{x}) d\theta$$

# Bayesian Terminology

- Conjugate priors
  - when posterior and prior belong to the same family, so the joint pdf is easy to compute
- Point estimate approximations of latent variables
  - instead of computing a distribution over all possible values for the variable, compute one point only, e.g. the most likely (maximum likelihood or max a posteriori estimate)
$$\theta^* = \arg_{\theta} \max p(x|\theta)p(\theta) \text{ (MAP)}$$
$$\theta^* = \arg_{\theta} \max p(x|\theta) \quad (\text{MLE})$$
  - Quite good when the posterior distribution is peaky (low variance)




# Bayesian Modelling

---


- Estimate the posterior density  $p(\theta|x)$  for your training data  $x$
- To do so, need to define the prior  $p(\theta)$  and likelihood  $p(x|\theta)$  distributions
- Once the  $p(\theta|x)$  density is estimated, Bayesian Inference is possible
  - $p(\theta|x)$  is a (density) function, not just a single number (point estimate)
- But how to estimate the posterior density?
  - Markov Chain Monte Carlo (MCMC) → Simulation-like estimation
  - Variational Inference → Turn estimation to optimization

# Variational Inference


- Estimating the true posterior  $p(\theta|x)$  is not always possible
  - especially for complicated models like neural networks
- Variational Inference assumes another function  $q(\theta|\varphi)$  with which we want to approximate the true posterior  $p(\theta|x)$ 
  - $q(\theta|\varphi)$  is the approximate posterior
  - Note that the approximate posterior does not depend on the observable variables  $x$
- We approximate by minimizing the **reverse KL-divergence** w.r.t.  $\varphi$ 
$$\varphi^* = \arg \min_{\varphi} KL(q(\theta|\varphi) || p(\theta|x))$$
- Turn inference into optimization




# Variational Inference (graphically)




# Variational Inference (graphically)





# Variational Inference (graphically)




How to overestimate variance?





Underestimating variance. Why?




# Variational Inference (graphically)



How to overestimate variance? Forward KL



Underestimating variance. Why?



# Mean-Field Approximation and CAVI Optimization

---

- To make the optimization of the VI easier, one can assume the latent variables are independent of each other

$$q(\theta|\varphi) = \prod_j q_j(\theta_j|\varphi_j)$$

- The optimization is often done with CAVI
  - Coordinate-Ascent Variational Inference
  - Initially set  $\varphi$  randomly
  - For each  $j$  in turn you set  $q_j(\theta_j|\varphi_j) = \mathbb{E}_{g_{-j}}[\log p(\theta|x)]$

# Variational Inference - Evidence Lower Bound (ELBO)

- Given latent variables  $\theta$  and the approximate posterior

$$q_\varphi(\theta) = q(\theta|\varphi)$$

- The log marginal is

$$\begin{aligned}\log p(x) &= \log \int p(x, \theta) d\theta \\ &= \log \int_{\theta} p(x, \theta) \frac{q_\varphi(\theta)}{q_\varphi(\theta)} d\theta \\ &= \log \mathbb{E}_{q_\varphi(\theta)} \left[ \frac{p(x, \theta)}{q_\varphi(\theta)} \right] \\ &\leq \mathbb{E}_{q_\varphi(\theta)} \left[ \log \frac{p(x, \theta)}{q_\varphi(\theta)} \right]\end{aligned}$$

$$\begin{aligned}&= \mathbb{E}_{q_\varphi(\theta)} [\log p(x, \theta)] - \mathbb{E}_{q_\varphi(\theta)} [\log q_\varphi(\theta)] \\ &= \mathbb{E}_{q_\varphi(\theta)} [\log p(x, \theta)] + H(\theta) \\ &= \text{ELBO}_{\theta, \varphi}(x)\end{aligned}$$

or

$$\begin{aligned}&= \mathbb{E}_{q_\varphi(\theta)} [\log p(x|\theta)] - \mathbb{E}_{q_\varphi(\theta)} [\log p(\theta)] \\ &\quad + \mathbb{E}_{q_\varphi(\theta)} [\log q_\varphi(\theta)] \\ &= \mathbb{E}_{q_\varphi(\theta)} [\log p(x|\theta)] - \text{KL}(q_\varphi(\theta)||p(\theta)) \\ &= \text{ELBO}_{\theta, \varphi}(x)\end{aligned}$$

# Variational Inference - Evidence Lower Bound (ELBO)

- Given latent variables  $\theta$  and the approximate posterior

$$q_\phi(\theta) = q(\theta|\phi)$$

- The log marginal is

$$\begin{aligned} \log p(x) &= \log \int_{\theta} p(x, \theta) d\theta \\ &\xrightarrow{\text{multiply and divide by } q(\theta|\phi)} \log \int_{\theta} \frac{p(x, \theta)}{q(\theta|\phi)} q(\theta|\phi) d\theta \\ &\xrightarrow{\text{by definition of } q_\phi(\theta)} \log \left[ \frac{\int_{\theta} p(x, \theta) q(\theta|\phi)}{q(\theta|\phi)} \right] \\ &\xrightarrow{\text{Tendances}} \leq E_{q(\theta|\phi)} \left[ \log \frac{p(x, \theta)}{q(\theta|\phi)} \right] \end{aligned}$$

$$= E_{q(\theta|\phi)} [\log p(x, \theta)] - E_{q(\theta|\phi)} [\log q(\theta|\phi)]$$

$$= \bar{E}_{q(\theta|\phi)} [\log p(x, \theta)] - H(q)$$

$\text{ELBO}(\theta|\phi)$

$$= E_{q(\theta|\phi)} [\log p(x|\theta)] - E_{q(\theta|\phi)} [\log q(\theta|\phi)]$$

$$+ E_{q(\theta|\phi)} [\log q_\phi(\theta)]$$

$$= E_{q(\theta|\phi)} [\log p(x|\theta)] - \mathcal{H}(q(\theta|\phi))$$

$\text{ELBO}$

# ELBO and the marginal

- It is easy to see that the ELBO is directly related to the marginal

$$\begin{aligned} \text{ELBO}_{\theta, \varphi}(x) &= \\ &= \mathbb{E}_{q_{\varphi}(\theta)}[\log p(x, \theta)] - \mathbb{E}_{q_{\varphi}(\theta)}[\log q_{\varphi}(\theta)] \\ &= \mathbb{E}_{q_{\varphi}(\theta)}[\log p(\theta|x)] + \mathbb{E}_{q_{\varphi}(\theta)}[\log p(x)] - \mathbb{E}_{q_{\varphi}(\theta)}[\log q_{\varphi}(\theta)] \\ &= \mathbb{E}_{q_{\varphi}(\theta)}[\log p(x)] - KL(q_{\varphi}(\theta)||p(\theta|x)) \\ &= \log p(x) - KL(q_{\varphi}(\theta)||p(\theta|x)) && \leftarrow \log p(x) \text{ does not depend on } q_{\varphi}(\theta) \\ &\Rightarrow \log p(x) = \text{ELBO}_{\theta, \varphi}(x) + KL(q_{\varphi}(\theta)||p(\theta|x)) && \leftarrow \mathbb{E}_{q_{\varphi}(\theta)}[1]=1 \end{aligned}$$

log p(x) = ELBO<sub>θ, φ</sub>(x) + KL(q<sub>φ</sub>(θ)||p(θ|x))

- You can also see  $\text{ELBO}_{\theta, \varphi}(x)$  as Variational Free Energy

# ELBO and the marginal

---

- It is easy to see that the ELBO is directly related to the marginal

$$\text{ELBO}_{\theta, \varphi}(x) =$$

# ELBO interpretations

$$\log p(x) = \text{ELBO}_{\theta, \varphi}(x) + KL(q_{\varphi}(\theta) || p(\theta|x))$$

- The log-likelihood is constant, as it does not depends on any parameter
- Also, both  $\text{ELBO}_{\theta, \varphi}(x) > 0$  and  $KL(q_{\varphi}(\theta) || p(\theta|x)) > 0$

1. The higher the Variational Lower Bound  $\text{ELBO}_{\theta, \varphi}(x)$ , the smaller the difference between the approximate posterior  $q_{\varphi}(\theta)$  and the true posterior  $p(\theta|x)$  → better latent representation
2. The Variational Lower Bound  $\text{ELBO}_{\theta, \varphi}(x)$  approaches the log-likelihood → better density model

# Amortized Inference

---


- The variational distribution  $\underline{q(\theta|\varphi)}$  does not depend directly on data
  - Only indirectly, via minimizing its ~~distance~~ to the true posterior  $KL(q(\theta|\varphi)||p(\theta|x))$
- So, with  $q(\theta|\varphi)$  we have a major optimization problem, as the approximate posterior must approximate the whole dataset  $x = [x_1, x_2, \dots, x_N]$  jointly
- As this is obviously quite complex, one can amortize the optimization on individual data points by setting

$$q(\theta|\varphi) = q_\varphi(\theta|x)$$

- Predict model parameters  $\theta$  using a  $\varphi$ -parameterized model of the input  $x$
- Use it for parameters that depend on data, such as the latent activations

# Amortized Inference (Intuitively)

- Originally, Variational Inference assumed that  $q(\theta|\varphi)$  describes the approximate posterior of the dataset as a whole
  - Think of  $\theta$  not as the latent activations  $z$ , but only the latent model variables  $w$



# Variational Autoencoders


- Let's rewrite the ELBO a bit more explicitly

$$\begin{aligned}\text{ELBO}_{\theta, \varphi}(x) &= \mathbb{E}_{q_\varphi(\theta)}[\log p(x|\theta)] - \text{KL}(q_\varphi(\theta)||p(\theta)) \\ &= \mathbb{E}_{q_\varphi(z|x)}[\log p_\theta(x|z)] - \text{KL}(q_\varphi(z|x)||p_\lambda(z))\end{aligned}$$

- Instead of  $p(x|\theta)$  we have  $p_\theta(x|z)$  to indicate that the model for the ~~posterior~~ density has weights parameterized by  $\theta$  and latent model activations parameterized by  $z$
- Instead of  $p(\theta)$  we have  $p_\lambda(z)$ , namely we put a  $\lambda$ -parameterized prior only on the latent activations  $z$  and not the model weights
- Instead of  $q(\theta|\varphi)$  we have  $q_\varphi(z|x)$  to indicate that the model approximates the posterior density of the latent activations, and the model weights are parameterized by  $\varphi$

# Variational Autoencoders

- So, we have  $\text{ELBO}_{\theta, \varphi}(x) = \mathbb{E}_{q_\varphi(z|x)}[\log p_\theta(x|z)] - \text{KL}(q_\varphi(z|x) || p_\lambda(z))$
- What if we model the densities  $p_\theta(x|z)$  and  $q_\varphi(z|x)$  as neural networks?
- The approximate posterior looks like a standard CovnNet (or MLP), which receives an image input  $x$  and returns a feature map/latent variable  $z$ 
  - Also known as encoder or inference network
- The likelihood term  $p_\theta(x|z)$  looks like an inverted ConvNet (deconvolutions), which given a latent feature map  $z$  reconstructs the input  $x$ 
  - Also known as decoder or generator network, because it recognizes the input given the latent variable
- A difference from a standard autoencoder is we now have an opinion of what the distribution of the latents  $z$  should look like, with  $p_\lambda(z)$



# Training Variational Autoencoders

---


- Maximize the Evidence Lower Bound (ELBO)

- Or minimize the negative ELBO

$$\mathcal{L}(\theta, \varphi) = \mathbb{E}_{q_\varphi(z|x)}[\log p_\theta(x|z)] - \text{KL}(q_\varphi(z|x)||p_\lambda(z))$$

- How to we optimize the ELBO?

# Training Variational Autoencoders



- Maximize the Evidence Lower Bound (ELBO)

- Or minimize the negative ELBO

$$\begin{aligned}\mathcal{L}(\theta, \varphi) &= \mathbb{E}_{q_\varphi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\varphi(Z|x) || p_\lambda(Z)) \\ &= \int_Z q_\varphi(z|x) \log p_\theta(x|z) dz - \int_Z q_\varphi(z|x) \log \frac{q_\varphi(z|x)}{p_\lambda(z)} dz\end{aligned}$$


- Forward propagation → compute the two terms
- The first term is an integral (expectation) that we cannot solve analytically.  
So, we need to sample from the pdf instead
  - When  $p_\theta(x|z)$  contains even a few nonlinearities, like in a neural network, the integral is hard to compute analytically

# Complex integrals

---

# Training Variational Autoencoders


- Maximize the Evidence Lower Bound (ELBO)
  - Or minimize the negative ELBO

$$\begin{aligned}\mathcal{L}(\theta, \varphi) &= \mathbb{E}_{q_\varphi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\varphi(z|x) || p_\lambda(z)) \\ &= \int_Z q_\varphi(z|x) \log p_\theta(x|z) dz - \int_Z q_\varphi(z|x) \log \frac{q_\varphi(z|x)}{p_\lambda(z)} dz\end{aligned}$$





- Forward propagation → compute the two terms
- The **first term** is an integral (expectation) that we cannot solve analytically. So, we need to sample from the pdf instead
  - When  $p_\theta(x|z)$  contains even a few nonlinearities, like in a neural network, the integral is hard to compute analytically
- The **second term** is the KL divergence between two distributions that we know

# Typical VAE

- We set the prior  $p_\lambda(Z)$  to be the unit Gaussian  
$$p(Z) \sim N(0, 1)$$
- We set the likelihood to be a Bernoulli for binary data  
$$p(X|Z) \sim \text{Bernoulli}(\pi)$$
- We set  $q_\varphi(Z|x)$  to be a neural network (MLP, ConvNet), which maps an input  $x$  to the Gaussian distribution, specifically it's mean and variance
  - $\mu_z, \sigma_z \sim q_\varphi(Z|x)$
  - The neural network has two outputs, one is the mean  $\mu_x$  and the other the  $\sigma_x$ , which corresponds to the covariance of the Gaussian
- We set  $p_\theta(X|Z)$  to be an inverse neural network, which maps  $Z$  to the Bernoulli distribution if our outputs binary (e.g. Binary MNIST)



# VAE: Interpolation in the latent space



# Forward propagation in VAE

---

- Sample  $z$  from the approximate posterior density  $z \sim q_\varphi(Z|x)$ 
  - As  $q_\varphi$  is a neural network that outputs values from a specific and known parametric pdf, e.g. a Gaussian, sampling from it is rather easy
  - Often even a single draw is enough
- Second, compute the  $\log p_\theta(x|Z)$ 
  - As  $p_\theta$  is a neural network that outputs values from a specific and known parametric pdf, e.g. a Bernoulli for white/black pixels, computing the log-prob is easy
- Computing the ELBO is rather straightforward in the standard case
- How should we optimize the ELBO?

# Forward propagation in VAE

---

- Sample  $z$  from the approximate posterior density  $z \sim q_\varphi(Z|x)$ 
  - As  $q_\varphi$  is a neural network that outputs values from a specific and known parametric pdf, e.g. a Gaussian, sampling from it is rather easy
  - Often even a single draw is enough
- Second, compute the  $\log p_\theta(x|Z)$ 
  - As  $p_\theta$  is a neural network that outputs values from a specific and known parametric pdf, e.g. a Bernoulli for white/black pixels, computing the log-prob is easy
- Computing the ELBO is rather straightforward in the standard case
- How should we optimize the ELBO? Backpropagation?

# Backward propagation in VAE

- Backpropagation → compute the gradients of
$$\mathcal{L}(\theta, \varphi) = \mathbb{E}_{z \sim q_\varphi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\varphi(z|x) || p_\lambda(z))$$

- $\nabla_\theta \mathcal{L} = \mathbb{E}_{z \sim q_\varphi(z|x)} [\nabla_\theta \log p_\theta(x|z)]$

- The expectation and sampling in  $\mathbb{E}_{z \sim q_\varphi(z|x)}$  does not depend on  $\theta$ , so no problem!
- Also, the KL does not depend on  $\theta$ , so no gradient from over there!

- $\nabla_\varphi \mathcal{L} = \nabla_\varphi \left[ \mathbb{E}_{z \sim q_\varphi(z|x)} [\log p_\theta(x|z)] \right] - \nabla_\varphi [\text{KL}(q_\varphi(z|x) || p_\lambda(z))]$

# Backward propagation in VAE

---

- Backpropagation → compute the gradients of

$$\mathcal{L}(\theta, \varphi) = \mathbb{E}_{z \sim q_\varphi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\varphi(Z|x) || p_\lambda(Z))$$

# Backward propagation in VAE

---

- Backpropagation → compute the gradients of

$$\mathcal{L}(\theta, \varphi) = \mathbb{E}_{z \sim q_\varphi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\varphi(Z|x) || p_\lambda(Z))$$

- $\nabla_\theta \mathcal{L} = \mathbb{E}_{z \sim q_\varphi(z|x)} [\nabla_\theta \log p_\theta(x|z)]$

- The expectation and sampling in  $\mathbb{E}_{z \sim q_\varphi(z|x)}$  does not depend on  $\theta$ , so no problem!
  - Also, the KL does not depend on  $\theta$ , so no gradient from over there!

- $\nabla_\varphi \mathcal{L} = \nabla_\varphi \left[ \mathbb{E}_{z \sim q_\varphi(z|x)} [\log p_\theta(x|z)] \right] - \nabla_\varphi [\text{KL}(q_\varphi(Z|x) || p_\lambda(Z))]$

- Problem?

# Backward propagation in VAE

---

- Backpropagation → compute the gradients of

$$\mathcal{L}(\theta, \varphi) = \mathbb{E}_{z \sim q_\varphi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\varphi(Z|x) || p_\lambda(Z))$$

- $\nabla_\theta \mathcal{L} = \mathbb{E}_{z \sim q_\varphi(z|x)} [\nabla_\theta \log p_\theta(x|z)]$

- The expectation and sampling in  $\mathbb{E}_{z \sim q_\varphi(z|x)}$  does not depend on  $\theta$ , so no problem!
  - Also, the KL does not depend on  $\theta$ , so no gradient from over there!

- $\nabla_\varphi \mathcal{L} = \nabla_\varphi \left[ \mathbb{E}_{z \sim q_\varphi(z|x)} [\log p_\theta(x|z)] \right] - \nabla_\varphi [\text{KL}(q_\varphi(Z|x) || p_\lambda(Z))]$

- **Problem?** Sampling  $z \sim q_\varphi(Z|x)$  is not differentiable → no gradients

- No gradients → No backprop → No training! → Solution?

# Solution: REINFORCE?

---

- So, our latent variable  $Z$  is a Gaussian (in the standard VAE) represented by the mean and variance  $\mu_Z, \sigma_Z$ , which are the output of a neural net
- So, we can train by sampling randomly from that Gaussian

$$z \sim N(\mu_Z, \sigma_Z)$$

- Once we have that  $z$ , however, it's a fixed value (not a function), so we cannot backprop
- We could use, however, the REINFORCE algorithm to compute an approximation to the gradient
  - High-variance gradients → slow and not very effective learning




# Solution: Reparameterization trick

- Remember, we have a Gaussian output  $z \sim N(\mu_z, \sigma_z)$
- For certain pdfs, including the Gaussian, we can rewrite their random variable  $z$  as deterministic transformations of a simpler random variable  $\varepsilon$
- For the Gaussian specifically, the following two formulations are equivalent  
$$z \sim N(\mu_z, \sigma_z) \Leftrightarrow z = \mu_z + \varepsilon \cdot \sigma_z,$$
where  $\varepsilon \sim N(0, 1)$  and  $\mu_z, \sigma_z$  are deterministic values from the NN function

# Solution: Reparameterization trick

- Instead of sampling from  $z \sim N(\mu_z, \sigma_z)$ , we sample from  $\varepsilon \sim N(0, 1)$  and then we compute  $z$
- Sampling directly from  $z \sim N(\mu_z, \sigma_z)$  leads to high-variance estimates
- Sampling directly from  $\varepsilon \sim N(0, 1)$  leads to low-variance estimates
  - Why low variance? Exercise for the interested reader
- Remember: since we are sampling for  $z$ , we are also sampling gradients
- More distributions beyond Gaussian possible: Laplace, Student-t, Logistic, Cauchy, Rayleigh, Pareto




# Check what is random

---



- Again, the latent variable is  $z = \mu_z + \varepsilon \cdot \sigma_z$
- $\mu_z$  and  $\sigma_z$  are deterministic functions (via the neural network encoder)
- $\varepsilon$  is a random variable, which comes externally
- The  $z$  as a result is itself a random variable, because of  $\varepsilon$
- However, now the randomness is not associated with the neural network and its parameters that we have to learn
  - The randomness instead comes from the external  $\varepsilon$
  - The gradients flow through  $\mu_z$  and  $\sigma_z$

# Reparameterization Trick (graphically)

Original form



Reparameterised form



: Deterministic node

: Random node

[Kingma, 2013]  
[Bengio, 2013]  
[Kingma and Welling 2014]  
[Rezende et al 2014]

# VAE Training Pseudocode

---

**Data:**

$\mathcal{D}$ : Dataset

$q_{\phi}(\mathbf{z}|\mathbf{x})$ : Inference model

$p_{\theta}(\mathbf{x}, \mathbf{z})$ : Generative model

**Result:**

$\theta, \phi$ : Learned parameters

$(\theta, \phi) \leftarrow$  Initialize parameters

**while** SGD not converged **do**

$\mathcal{M} \sim \mathcal{D}$  (Random minibatch of data)

$\epsilon \sim p(\epsilon)$  (Random noise for every datapoint in  $\mathcal{M}$ )

Compute  $\tilde{\mathcal{L}}_{\theta, \phi}(\mathcal{M}, \epsilon)$  and its gradients  $\nabla_{\theta, \phi} \tilde{\mathcal{L}}_{\theta, \phi}(\mathcal{M}, \epsilon)$

Update  $\theta$  and  $\phi$  using SGD optimizer

**end**

The ELBO's gradients

“ **i want to talk to you .** ”  
“ *i want to be with you .* ”  
“ *i do n’t want to be with you .* ”  
*i do n’t want to be with you .*  
**she did n’t want to be with him .**


---

**he was silent for a long moment .**  
*he was silent for a moment .*  
*it was quiet for a moment .*  
*it was dark and cold .*  
*there was a pause .*  
**it was my turn .**

---

Figure 2.D.2: An application of VAEs to interpolation between pairs of sentences, from [Bowman et al., 2015]. The intermediate sentences are grammatically correct, and the topic and syntactic structure are typically locally consistent.

# VAE for Image Resynthesis



*Smile vector:*  
mean smiling faces –  
mean no-smile faces

**Latent space arithmetic**

Figure 2.D.3: VAEs can be used for image re-synthesis. In this example by White [2016], an original image (left) is modified in a latent space in the direction of a *smile vector*, producing a range of versions of the original, from smiling to sadness. Notice how changing the image along a single vector in latent space, modifies the image in many subtle and less-subtle ways in pixel space.

# VAE for designing chemical compounds




Figure 2.D.1: Example application of a VAE in [Gómez-Bombarelli et al., 2016]: design of new molecules with desired chemical properties. (a) A latent continuous representation  $\mathbf{z}$  of molecules is learned on a large dataset of molecules. (b) This continuous representation enables gradient-based search of new molecules that maximizes some chosen desired chemical property given by objective function  $f(\mathbf{z})$ .


# Normalizing Flows

<https://www.shakirm.com/slides/DeepGenModelsTutorial.pdf>

<https://blog.evjang.com/2018/01/hf1.html>

<https://arxiv.org/pdf/1505.05770.pdf>

- Using simple pdfs, like a Gaussian, for the approximate posterior limits the expressivity of the model
- Better make sure the approximate posterior comes from a class of models that can even contain the true posterior
- Use a series of  $K$  invertible transformations to construct the approximate posterior
  - $z_k = f_k \circ f_{k-1} \circ \dots \circ f_1(z_0)$
  - Rule of change for variables




Changing from the  $x$  variable to  $y$  using the transformation  $y = f(x) = 2x + 1$

# Normalizing Flows


<https://www.shakirm.com/slides/DeepGenModelsTutorial.pdf>

<https://blog.evjang.com/2018/01/hf1.html>

<https://arxiv.org/pdf/1505.05770.pdf>



# Normalizing Flows



<https://www.shakirm.com/slides/DeepGenModelsTutorial.pdf>

# Normalizing Flows on Non-Euclidean Manifolds











Figure 1: Left: Construction of a complex density on  $S^n$  by first projecting the manifold to  $R^n$ , transforming the density and projecting it back to  $S^n$ . Right: Illustration of transformed ( $S^2 \rightarrow R^2$ ) densities corresponding to an uniform density on the sphere. Blue: empirical density (obtained by Monte Carlo); Red: Analytical density from equation (4); Green: Density computed ignoring the intrinsic dimensionality of  $S^n$ .

$$\log q_K(\mathbf{z}_K) = \log q_0(\mathbf{z}_0) - \frac{1}{2} \sum_{k=1}^K \log \det |\mathbf{J}_\phi^\top \mathbf{J}_\phi|$$

Gemici et al., 2016

<https://www.shakirm.com/slides/DeepGenModelsTutorial.pdf>

# Normalizing Flows on Non-Euclidean Manifolds



## Summary

- Gentle intro to Bayesian Modelling and Variational Inference
- Restricted Boltzmann Machines
- Deep Boltzmann Machines
- Deep Belief Network
- Contrastive Divergence
- Variational Autoencoders
- Normalizing Flows